

**Fundamentals of Database Systems**  
**Laboratory Manual<sup>1</sup>**

**Rajshekhar Sunderraman**  
**Georgia State University**

**June 2006**

---

<sup>1</sup> To accompany Elmasri and Navathe, Fundamentals of Database Systems, 5<sup>th</sup> Edition, Addison-Wesley, 2007.

## Preface

This laboratory manual accompanies the popular database textbook *Elmasri and Navathe, Fundamentals of Database Systems, 5<sup>th</sup> Edition, Addison-Wesley, 2007*. It provides supplemental materials to enhance the practical coverage of concepts in an introductory database systems course. The material presented in this laboratory manual complements the first 11 chapters of the Elmasri/Navathe text typically covered in most introductory database systems courses.

### Chapter Mappings

The laboratory manual consists of 6 chapters and the following table shows the mapping to the chapters in the Elmasri/Navathe textbook:

Laboratory Manual Chapter	Elmasri/Navathe 5th Edition Chapter(s)
Chapter 1	Chapters 3, 4, and 7
Chapter 2	Chapter 6
Chapter 3	Chapter 9
Chapter 4	Chapter 26
Chapter 5	Chapters 10 and 11
Chapter 6	Chapters 9 and 26

Chapter 1 presents ERWin, a popular data modeling software that allows database designers to represent Entity-Relationship diagrams and automatically generate relational SQL code to create the database in one of several commercial relational database management systems such as Oracle or Microsoft SQLServer. The material presented in this chapter is tutorial in nature and covers the COMPANY database design of the Elmasri/Navathe text in detail.

Chapter 2 presents two interpreters that can be used to execute queries in Relational Algebra and Domain Relational Calculus. These interpreters are part of a Java package that includes a rudimentary database engine capable of storing relations and able to perform basic relational algebraic operations on these relations. It is hoped that these interpreters will allow the student to get a better understanding of abstract query languages such as relational algebra and relational calculus.

Chapter 3 presents interacting and programming with Oracle database management system. A popular data loading tool for Oracle databases called SQL Loader is introduced and the COMPANY database of the Elmasri/Navathe text is extended with additional data to make it more interesting to program with. Programming applications that access Oracle databases is then introduced in Java using the JDBC interface. Several non-trivial example programs are discussed.

Chapter 4 covers MySQL database management system, a popular open source database system that is increasing used by small and medium sized organizations. Programming Web applications in PHP that accesses MySQL databases is introduced with a complete database browser application for the COMPANY database.

Chapter 5 introduces a Prolog-based toolkit for relational database design. The toolkit, called Database Designer (DBD), allows the student to work with numerous concepts and algorithms that deal with functional dependency theory and data normalization. The student may use DBD to verify answers to many questions related to functional dependency theory and normalization algorithms.

Chapter 6 presents several semester long projects for students in introductory database courses to complete. These projects may be implemented in Java, PHP or any other favorite programming language and may access Oracle, MySQL or any other relational database management system.

### **Code**

The laboratory manual comes with all the code and data presented in the different chapters. The software for the relational query interpreters as well as the database designer (DBD) also accompanies the laboratory manual.

### **Software**

The software systems discussed and used in the laboratory manual are ERWin from Computer Associates, Oracle DBMS from Oracle, and MySQL, PHP, and SWI-Prolog from open source. Both Computer Associates and Oracle have educational pricing for their software and we expect the individual universities and colleges that use this laboratory manual to provide the software for use by their students.

### **Acknowledgements**

The author wishes to acknowledge the contributions of many of his students. In particular, the author appreciates the diligent effort put forth by Ms. Manisha Jain in building and testing the DRC interpreter of Chapter 2.

Rajshekhar Sunderraman  
Atlanta, Georgia  
June 2006

# Contents

<b>1. ER MODELING TOOLS .....</b>	<b>6</b>
1.1 STARTING WITH ERWIN.....	6
1.2 ADDING ENTITY TYPES.....	7
1.3 ADDING RELATIONSHIPS .....	10
1.4 FORWARD ENGINEERING .....	12
1.5 SUPERTYPE/SUBTYPE EXAMPLE.....	15
EXERCISES.....	17
<b>2. INTERPRETERS FOR ABSTRACT QUERY LANGUAGES.....</b>	<b>21</b>
2.1 CREATING THE DATABASE .....	21
2.2 RELATIONAL ALGEBRA INTERPRETER .....	23
2.2.1 <i>Relational Algebra Syntax</i> .....	23
2.2.2 <i>Naming of Intermediate Relations and Attributes</i> .....	25
2.2.3 <i>Relational Algebraic Operators Supported by the RA Interpreter</i> .....	26
2.2.4 <i>Examples</i> .....	27
2.3 DOMAIN RELATIONAL CALCULUS INTERPRETER.....	30
2.3.1 <i>Domain Relational Calculus Syntax</i> .....	31
2.3.2 <i>Safe DRC Queries</i> .....	32
2.3.3 <i>DRC Query Examples</i> .....	34
EXERCISES.....	35
<b>3. RELATIONAL DATABASE MANAGEMENT SYSTEMS: ORACLE.....</b>	<b>38</b>
3.1 COMPANY DATABASE.....	38
3.2 <i>SQL*PLUS</i> UTILITY .....	42
3.3 <i>SQL*LOADER</i> UTILITY .....	43
3.4 PROGRAMMING WITH ORACLE USING THE JDBC API .....	46
EXERCISES.....	56
<b>4. RELATIONAL DATABASE MANAGEMENT SYSTEMS: MYSQL .....</b>	<b>62</b>
4.1 COMPANY DATABASE.....	62
4.2 <i>MYSQL</i> UTILITY .....	66
4.3 <i>MYSQL</i> AND <i>PHP</i> PROGRAMMING .....	68
EXERCISES.....	80
<b>5. DATABASE DESIGN (DBD) TOOLKIT .....</b>	<b>83</b>
5.1 CODING RELATIONAL SCHEMAS AND FUNCTIONAL DEPENDENCIES .....	83
5.2 INVOKING THE SWI-PROLOG INTERPRETER.....	83
5.3 DBD SYSTEM PREDICATES .....	85
5.3.1 <i>xplus(R,F,X,Xplus)</i> .....	85
5.3.2 <i>finfplus(R,F,[X,Y])</i> .....	86
5.3.3 <i>fplus(R,F,Fplus)</i> .....	86
5.3.4 <i>implies(R,F1,F2)</i> and <i>equiv(R,F1,F2)</i> .....	87
5.3.5 <i>superkey(R,F,K)</i> and <i>candkey(R,F,K)</i> .....	88
5.3.6 <i>mincover(R,F,FC)</i> .....	89
5.3.7 <i>ljd(R,F,R1,R2)</i> , <i>ljd(R,F,D)</i> , and <i>fpd(R,F,D)</i> .....	90
5.3.8 <i>is3NF(R,F)</i> and <i>threenf(R,F,D)</i> .....	93
5.3.9 <i>isBCNF(R,F)</i> and <i>bcnf(R,F,D)</i> .....	93
EXERCISES.....	94

<b>6. PROJECTS .....</b>	<b>99</b>
<b>6.1 STUDENT REGISTRATION SYSTEM (GoLUNAR) .....</b>	<b>99</b>
<b>6.2 ONLINE BOOK STORE DATABASE SYSTEM .....</b>	<b>108</b>
<b>6.3 ONLINE SHOPPING SYSTEM .....</b>	<b>117</b>
<b>6.4 ONLINE BULLETIN BOARD SYSTEM .....</b>	<b>123</b>
<b>BIBLIOGRAPHY .....</b>	<b>127</b>

# CHAPTER 1

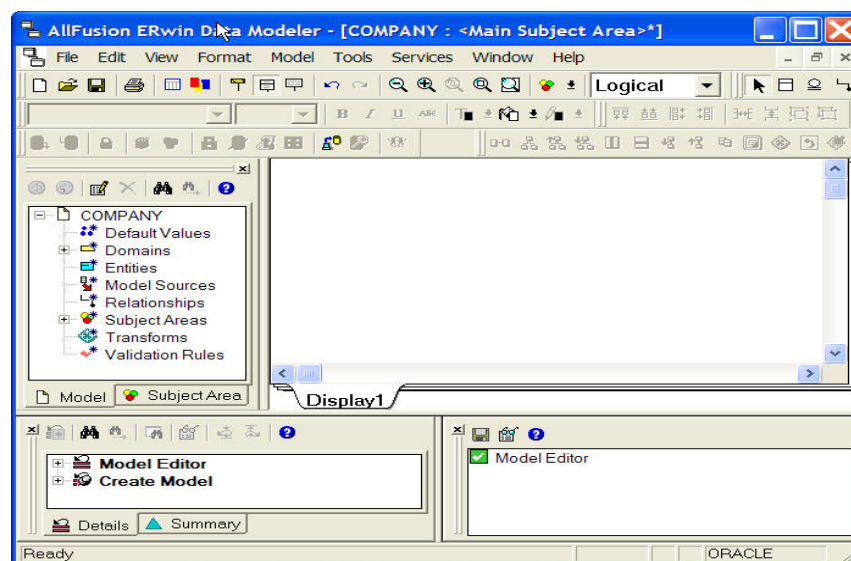
## ER Modeling Tools

This chapter introduces ERWin, a popular data modeling tool used in the industry. ERWin is a powerful tool which allows database designers to enter their Entity Relationship (ER) diagrams in a graphical form and produce physical database designs for popular relational database management systems such as Oracle and Microsoft SQLServer.

The use of ERWin is illustrated in this chapter using the ER schema diagram for the COMPANY database shown in Figure 3.2 of the Elmasri/Navathe text.

### 1.1 Starting with ERWin

The ERWin Data Modeler workspace is shown in Figure 1.1.



**Figure 1.1: ERWin Data Modeler Workspace**

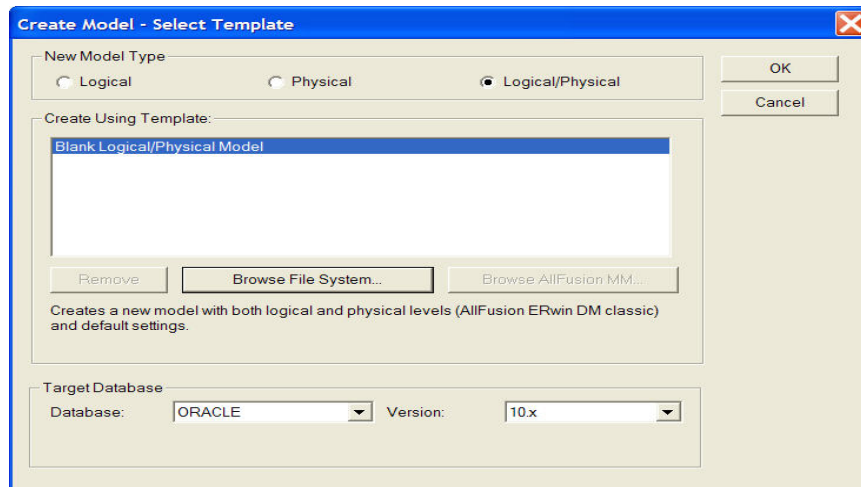
The top part of the workspace consists of Menu and Toolbars. The middle part of the workspace consists of two panes: the *model explorer pane* on the left providing a text based view of the data model and the *diagram window pane* on the right providing a graphical view of the data model. The lower part of the workspace consists of two panes: the *action log pane* on the left that displays a log of all changes made to the data model under design and the *advisories pane* that displays messages associated with the actions performed on the data model under design.

ERWin supports three model types for use by the database designer:

1. Logical: A conceptual model that includes entities, relationships, and attributes. This model type is essentially at the ER modeling level.

2. Physical: A database specific model that contains relational tables, columns and associated data types.
3. Logical/Physical: A single model that includes both the conceptual level objects as well as physical level tables. In this chapter we will use this model type.

To create a model in ERWin, one should launch the program and then choose the New option from the File menu. The Create Model dialog appears as shown in Figure 1.2.



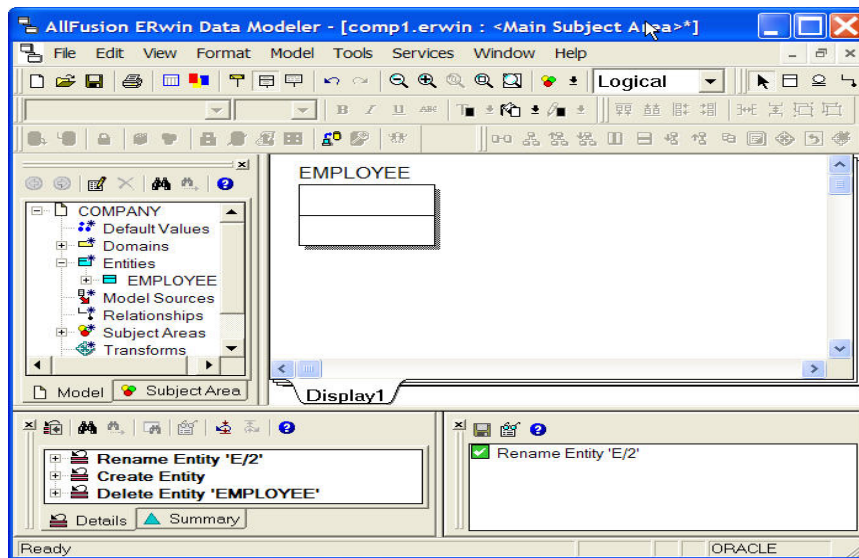
**Figure 1.2: Create Model dialog window**

In this dialog window, the user should choose the type of model. Typically the Logical/Physical model type should be chosen if the final goal is to produce a relational design for the database. The target database may also be chosen. In this case, Oracle 10.x version is chosen as the target database. In a future step, we will illustrate how ERWin can be used to generate SQL code to create the database objects in Oracle 10.x database.

The workspace for the new model will be populated by the system with a default name of Model\_n. This name may be changed in the model explorer pane by right clicking the model name and choosing the Properties option. This brings up a new window in which the name and other properties of the model may be changed. Besides changing the model name, the “Transform” options should be checked. This would allow for many-to-many relationships to be transformed correctly into separate relational tables in the physical model. In addition any sub-type/super-type relationships will also be transformed correctly in the physical model.

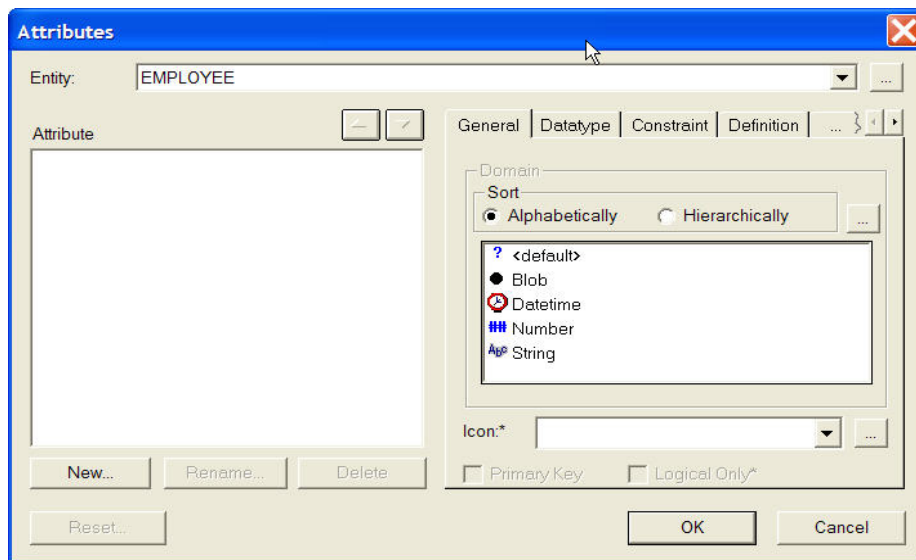
## 1.2 Adding Entity Types

To add an entity type to the database design, the user may either right click the “Entities” entry in the model explorer pane and choose “New” or choose the “Entity” icon in the Menus and Toolbars section of the workspace and click in the diagram window pane. An entity box shows up in the diagram window pane with a default entity name (E/n) which can be changed either in the diagram window pane or in the model explorer pane. Figure 1.3 shows the addition of the EMPLOYEE entity type in the COMPANY database.



**Figure 1.3: Add EMPLOYEE entity to the COMPANY database**

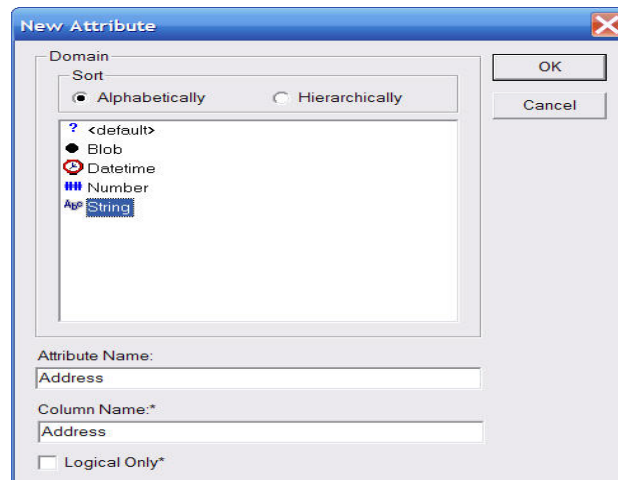
To add attributes to the EMPLOYEE entity type, the user may right click within the EMPLOYEE entity box in the diagram window pane and choose “Attributes”. This brings up a separate window using which new attributes may be added. The attribute window is shown in Figure 1.4.



**Figure 1.4: Attribute Window**

The user may now add attributes one at a time by clicking the “New” button. A separate window pops up as shown in Figure 1.5.

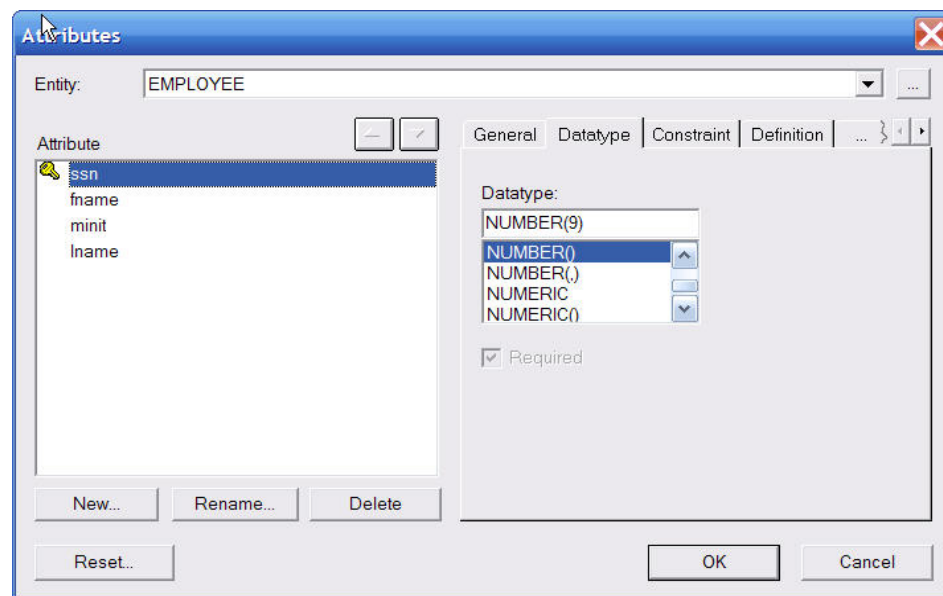




**Figure 1.5: New Attribute Window**

The user may choose an appropriate Domain (data type) and enter the Attribute Name and click OK. The data type may be further refined in the Attribute Window by choosing the Datatype tab and entering a precise data type. The user may also choose to designate this attribute as a primary key by selecting this option in the Attribute window.

After adding a few attributes to the EMPLOYEE entity type the Attributes window is shown in Figure 1.6.



**Figure 1.6: Attribute Window with four attributes**

In this way, we can create each of entity types: EMPLOYEE, DEPARTMENT, PROJECT, and DEPENDENT for the COMPANY database.

## Weak Entity Sets

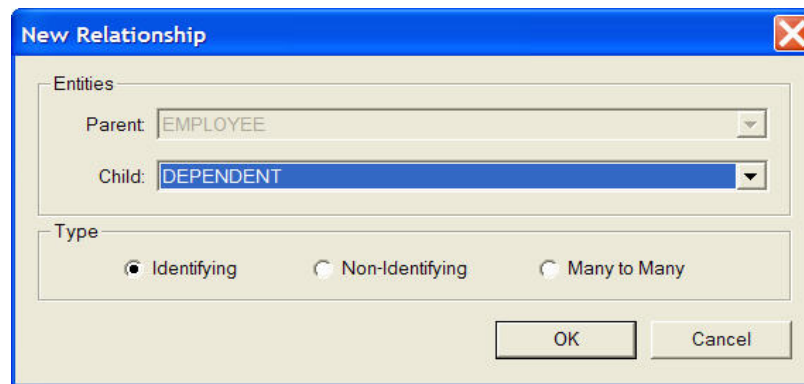
By default any entity type created as discussed so far is classified as an independent entity type. ERWin will classify an entity type as “weak” as soon as it participates in an identifying relationship. For example, the entity type DEPENDENT will be classified as “weak” in a subsequent step when we add the identifying relationship from EMPLOYEE to DEPENDENT in the next section. Weak entity types are denoted by rounded rectangles in the diagram window pane.

## Multi-Valued Attributes

Multi-valued attributes such as the locations attribute for the DEPARTMENT entity type cannot be modeled easily with ERWin. To handle such attributes, a separate entity type LOCATIONS is created and a many-to-many relationship between DEPARTMENT and LOCATIONS will be established in the next section.

## 1.3 Adding Relationships

Three types of relationships are supported in ERWin: identifying, non-identifying, and many-to-many. ERWin classifies the child entity type in an identifying relationship as “weak”. To add a relationship, the user may simply right click the Relationships entry in the model explorer pane and choose “New”. This pops up a new relationship window as shown in Figure 1.7.



**Figure 1.7: New Relationship Window**

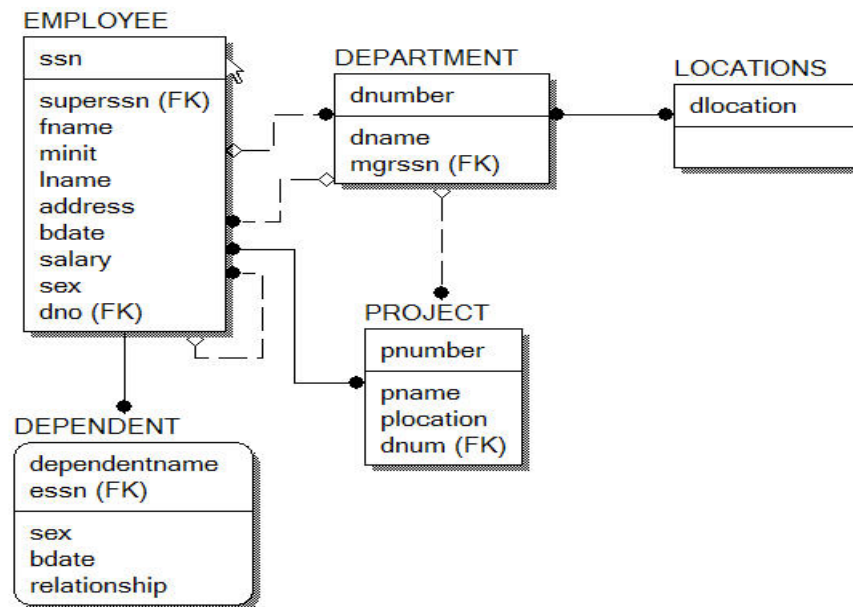
After choosing the parent and child entity types and the type of relationship and clicking OK, the new relationship is added and is reflected by a line connecting the two entity types in the diagram window pane. Many-to-many relationships are denoted by a solid connecting line with two black dots at the two ends. Non-identifying relationships are denoted by a dashed connecting line with a black dot at many-end and a square-shaped symbol at the one-end. Identifying relationships are denoted by a solid connecting line with a black dot at the many-end and nothing special at the one-end.

After creating a new relationship, the user may add verb phrases and other properties of the relationship by right clicking the connecting line in the diagram and choosing properties.

In the case of the COMPANY database, we create the following relationships:

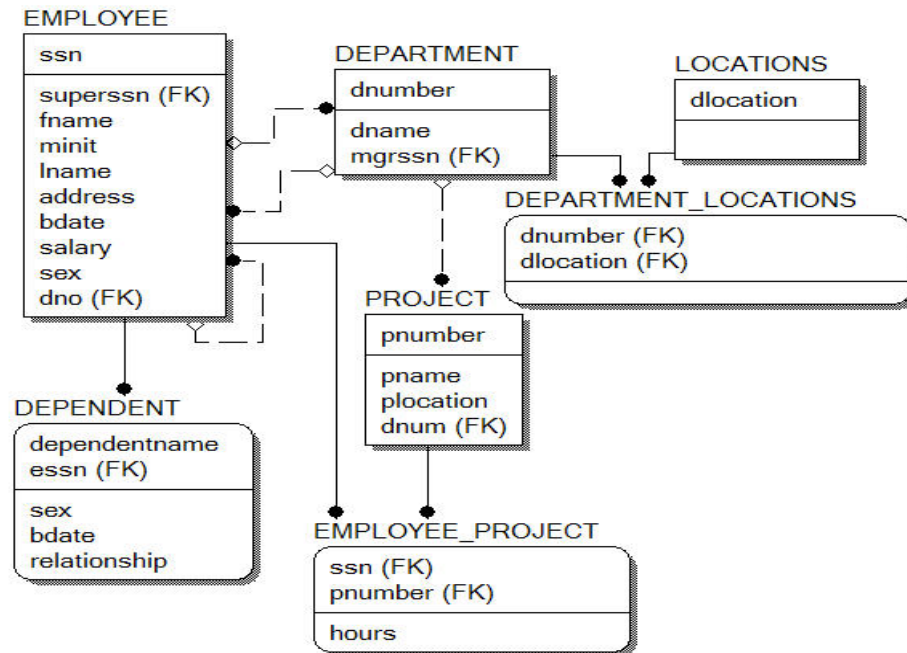
- One identifying relationship from EMPLOYEE to DEPENDENT.
- Two many-to-many relationships, one from EMPLOYEE to PROJECT and the other from DEPARTMENTS to LOCATIONS, and
- Four non-identifying relationships: from EMPLOYEE to DEPARTMENT (one-to-one for manages), from DEPARTMENT to EMPLOYEE (one-to-many for works for relationship), from EMPLOYEE to EMPLOYEE (one-to-many for supervisor/supervisee relationship), and from DEPARTMENT to PROJECT (one-to-many for the controls relationship).

The final logical ER diagram from the diagram window pane is shown in Figure 1.8.



**Figure 1.8: Final Logical ER Diagram**

Notice that the two many-to-many relationships do not have the transforms applied yet. The transforms are shown in the physical ER diagram (obtained by switching from Logical to Physical in the Menu and Toolbar section) in Figure 1.9. Notice the introduction of the two new “entity types” for the two many-to-many relationships. These are introduced by ERWin as a result of defining the transforms at the model level.



**Figure 1.9: Final Physical ER Diagram**

## 1.4 Forward Engineering

ERWin provides a powerful feature called forward engineering which allows the database designer to convert the ER design into a schema generation SQL script for one or more target relational databases. The following SQL script is obtained for the COMPANY database by choosing Tools→Forward Engineering→Schema-Generation option in the Menus and Toolbars section and clicking the “Preview” button.

```
CREATE TABLE DEPARTMENT
(
    dname VARCHAR2(20) NOT NULL ,
    dnumber INTEGER NOT NULL ,
    mgrssn NUMBER(9) NULL
);

ALTER TABLE DEPARTMENT
    ADD PRIMARY KEY (dnumber);

CREATE TABLE DEPARTMENT_LOCATIONS
(
    dnumber INTEGER NOT NULL ,
    dlocation VARCHAR2(20) NOT NULL
);

ALTER TABLE DEPARTMENT_LOCATIONS
    ADD PRIMARY KEY (dnumber,dlocation);
```

```
CREATE TABLE DEPENDENT
(
    dependentname VARCHAR2(20) NOT NULL ,
    sex CHAR NULL ,
    bdate DATE NULL ,
    relationship VARCHAR2(20) NULL ,
    essn NUMBER(9) NOT NULL
);
```

```
ALTER TABLE DEPENDENT
    ADD PRIMARY KEY (dependentname,essn);
```

```
CREATE TABLE EMPLOYEE
(
    ssn NUMBER(9) NOT NULL ,
    superssn NUMBER(9) NULL ,
    fname VARCHAR2(20) NULL ,
    minit CHAR NULL ,
    lname VARCHAR2(20) NOT NULL ,
    address VARCHAR2(50) NULL ,
    bdate DATE NULL ,
    salary NUMBER(8) NULL ,
    sex CHAR NULL ,
    dno INTEGER NULL
);
```

```
ALTER TABLE EMPLOYEE
    ADD PRIMARY KEY (ssn);
```

```
CREATE TABLE EMPLOYEE_PROJECT
(
    ssn NUMBER(9) NOT NULL ,
    pnumber INTEGER NOT NULL ,
    hours NUMBER(3) NULL
);
```

```
ALTER TABLE EMPLOYEE_PROJECT
    ADD PRIMARY KEY (ssn,pnumber);
```

```
CREATE TABLE LOCATIONS
(
    dlocation VARCHAR2(20) NOT NULL
);
```

```
ALTER TABLE LOCATIONS
    ADD PRIMARY KEY (dlocation);
```

```

CREATE TABLE PROJECT
(
    pnumber INTEGER NOT NULL ,
    pname VARCHAR2(20) NULL ,
    plocation VARCHAR2(20) NULL ,
    dnum INTEGER NULL
);

ALTER TABLE PROJECT
    ADD PRIMARY KEY (pnumber);

ALTER TABLE DEPARTMENT
    ADD ( FOREIGN KEY (mgrssn) REFERENCES EMPLOYEE(ssn) ON DELETE SET NULL);

ALTER TABLE DEPARTMENT_LOCATIONS
    ADD ( FOREIGN KEY (dnumber) REFERENCES DEPARTMENT(dnumber));

ALTER TABLE DEPARTMENT_LOCATIONS
    ADD ( FOREIGN KEY (dlocation) REFERENCES LOCATIONS(dlocation));

ALTER TABLE DEPENDENT
    ADD ( FOREIGN KEY (essn) REFERENCES EMPLOYEE(ssn));

ALTER TABLE EMPLOYEE
    ADD ( FOREIGN KEY (superssn) REFERENCES EMPLOYEE(ssn) ON DELETE SET
NULL);

ALTER TABLE EMPLOYEE
    ADD ( FOREIGN KEY (dno) REFERENCES DEPARTMENT(dnumber) ON DELETE SET
NULL);

ALTER TABLE EMPLOYEE_PROJECT
    ADD ( FOREIGN KEY (ssn) REFERENCES EMPLOYEE(ssn));

ALTER TABLE EMPLOYEE_PROJECT
    ADD ( FOREIGN KEY (pnumber) REFERENCES PROJECT(pnumber));

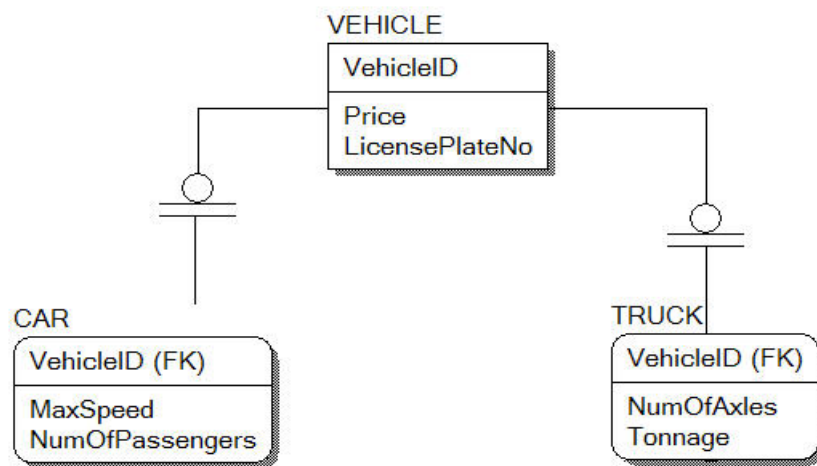
ALTER TABLE PROJECT
    ADD ( FOREIGN KEY (dnum) REFERENCES DEPARTMENT(dnumber) ON DELETE SET
NULL);

```

The above SQL script contains table definitions and basic primary and foreign key constraints definitions. ERWin does provide a number of options to generate views, triggers, indices etc. and these can be set in the forward engineering schema generation window.

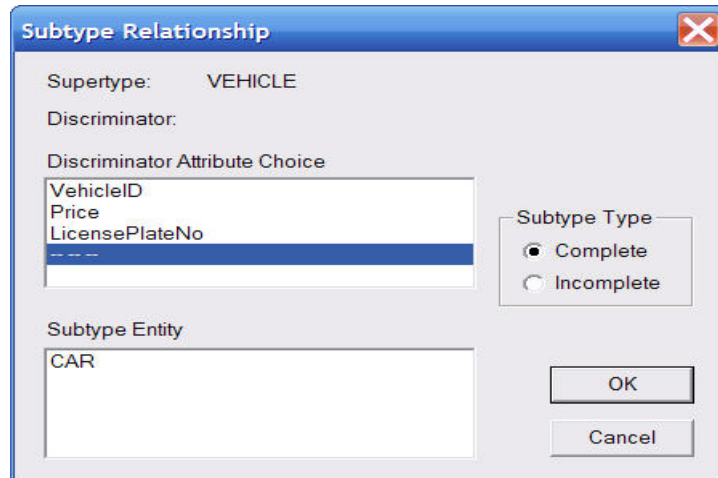
## 1.5 Supertype/Subtype Example

ERWin supports the creation of sub-type/super-type relationships between entity types. Consider the example in Figure 4.3 of the Elmasri/Navathe text in which a super-type entity VEHICLE consists of two sub-types CAR and TRUCK. To create this design in ERWin, the three entity types are created first. Then, the user may click the sub-type button (a circle with two parallel lines below the circle) in the Menus and Toolbars section, followed by clicking the super-type entity (VEHICLES) in the diagram window pane, followed by clicking the sub-type entity (CAR) in the diagram window pane. This process may be repeated for adding other sub-types (TRUCK in this example). The logical model for this example is shown in Figure 1.10.



**Figure 1.10: Sub-type/Super-type Logical ER Diagram**

To customize the properties of the sub-type/super-type relationship, the user may right click the relationship symbol (circle with two parallel lines) and choose Subtype Relationship. This brings up a window shown in Figure 1.11. The user may choose “Complete” subtype (when all categories are known) or “Incomplete” subtype (when all categories may not be known). The user may also add verb phrases etc by right-clicking the relationship line and choosing properties as was done for ordinary relationships. ERWin also allows the user to choose a “discriminator” attribute for the sub-types (an attribute in the super-type whose values determine the sub-type object). If no discriminator attribute is defined, the user may choose “-- -- --”.



The dialog box titled "Subtype Relationship" has a blue title bar with a close button. It contains the following fields and controls:

- Supertype:** VEHICLE
- Discriminator:** (empty)
- Discriminator Attribute Choice:** A list box containing "VehicleID", "Price", "LicensePlateNo", and "----". "VehicleID" is selected.
- Subtype Type:** Two radio buttons: "Complete" (selected) and "Incomplete".
- Subtype Entity:** A text box containing "CAR".
- Buttons:** "OK" and "Cancel" at the bottom right.

**Figure 1.11: Subtype Relationship Properties**

The following SQL script is produced using the forward engineering feature of ERWin for the Vehicles example:

```
CREATE TABLE CAR
(
    MaxSpeed INTEGER NULL ,
    NumOfPassengers INTEGER NULL ,
    VehicleID INTEGER NOT NULL
);

ALTER TABLE CAR
    ADD PRIMARY KEY (VehicleID);

CREATE TABLE TRUCK
(
    NumOfAxles INTEGER NULL ,
    Tonnage INTEGER NULL ,
    VehicleID INTEGER NOT NULL
);

ALTER TABLE TRUCK
    ADD PRIMARY KEY (VehicleID);

CREATE TABLE VEHICLE
(
    VehicleID INTEGER NOT NULL ,
    Price NUMBER(8,2) NULL ,
    LicensePlateNo VARCHAR2(20) NULL
);

ALTER TABLE VEHICLE
    ADD PRIMARY KEY (VehicleID);
```



```
ALTER TABLE CAR
  ADD ( FOREIGN KEY (VehicleID) REFERENCES VEHICLE(VehicleID));
```

```
ALTER TABLE TRUCK
  ADD ( FOREIGN KEY (VehicleID) REFERENCES VEHICLE(VehicleID));
```

## Exercises

### ER Modeling Problems

1. Consider the *university* database described in Exercise 3.16 of the Elmasri/Navathe text. Enter the ER schema for this database using a data modeling tool such as ERWin.
2. Consider a *mail order* database in which employees take orders for parts from customers. The data requirements are summarized as follows:
  - The mail order company has employees identified by a unique employee number, their first and last names, and a zip code where they are located.
  - The customers of the company are identified by a unique customer number, their first and last names, and a zip code where they are located.
  - The parts being sold by the company are identified by a unique part number, a part name, their price, and quantity in stock.
  - Orders placed by customers are taken by employees and are given a unique order number. Each order may contain certain quantities of one or more parts and their received date as well as a shipped date is recorded.

Design an Entity-Relationship diagram for the mail order database and enter the design using a data modeling tool such as ERWin.

3. Consider a *movie* database in which data is recorded about the movie industry. The data requirements are summarized as follows:
  - Movies are identified by their title and year of release. They have a length in minutes. They also have a studio that produces the movie and are classified under one or more genres (such as horror, action, drama etc). Movies are directed by one or more directors and have one or more actors acting in them. The movie also has a plot outline. Each movie also has zero or more quotable quotes which are spoken by a particular actor acting in the movie.
  - Actors are identified by their names and date of birth and act in one or more movies. Each actor has a role in the movie.
  - Directors are also identified by their names and date of birth and direct one or more movies. It is possible for a director to act in a movie (not necessarily in a movie they direct).

- Studios are identified by their names and have an address. They produce one or more movies.

Design an Entity-Relationship diagram for the movie order database and enter the design using a data modeling tool such as ERWin.

4. Consider a *conference review* system database in which researchers submit their research papers for consideration. The database system also caters to reviewers of papers who make recommendations on whether to accept or reject the paper. The data requirements are summarized as follows:

Authors of papers are uniquely identified by their email id. Their first and last names are also recorded.

- Papers are assigned unique identifiers by the system and are described by a title, an abstract, and a file name containing the actual paper.
- Papers may have multiple authors, but one of the authors is designated as the contact author.
- Reviewers of papers are uniquely identified by their email id. Their first and last names are also recorded.
- Each paper is assigned between two and four reviewers. The reviewers rate the paper assigned to them on a scale of 1 to 10.
- Each review contains two types of written comments: one to be seen by the review committee only and the other by the author(s) as well.

Design an Entity-Relationship diagram for the conference review database and enter the design using a data modeling tool such as ERWin.

5. Consider the ER diagram for the AIRLINE database shown in Figure 3.17 of the Elmasri/Navathe text. Enter this design using a data modeling tool such as ERWin.

### Enhanced ER Modeling Problems

6. Consider a *grade book* database in which instructors within an academic department maintain scores/points obtained by individual students in their classes. The data requirements are summarized as follows:
  - Students are identified by a unique student id, their first and last names, and an email address.
  - The instructor teaches certain courses each term. The courses are identified by a course number, a section number, and the term in which they are taught. The instructor also assigns grade cutoffs (example 90, 80, 70, and 60) for letter grades A, B, C, D, and F for each course he or she teaches.
  - Students are enrolled in courses taught by the instructor.
  - Each course being taught by the instructor has a number of grading components (such as mid-term, final exam, project, etc.). Each grading component has a maximum number of points (such as 100 or 50) and a weight (such as 20% or 10%). The weights of all the grading components of a course usually add up to 100.

- Finally, the instructor records the points earned by each student in each of the grading components in each of the courses. For example, student with id=1234 earns 84 points for the grading component mid-term for the course CSc 2310 section 2 in the fall 2005 term. The mid-term grading component may have been defined to have a maximum of 100 points and a weight of 20% of the course grade.

Design an enhanced Entity-Relationship diagram for the grade book database and enter the design using a data modeling tool such as ERWin.

7. Consider an *online auction* database system in which members (buyers and sellers) participate in the sale of items. The data requirements for this system are summarized as follows:

- The online site has members who are identified by a unique member id and are described by an email address, their name, a password, their home address, and a phone number.
- A member may be a buyer or a seller. A buyer has a shipping address recorded in the database. A seller has a bank account number and routing number recorded in the database.
- Items are placed by a seller for sale and are identified by a unique item number assigned by the system. Items are also described by an item title, an item description, a starting bid price, bidding increment, the start date of the auction, and the end date of the auction.
- Items are also categorized based on a fixed classification hierarchy (for example a modem may be classified as /COMPUTER/HARDWARE/MODEM).
- Buyers make bids for items they are interested in. A bidding price and time of bid placement is recorded. The person at the end of the auction with the highest bid price is declared the winner and a transaction between the buyer and the seller may proceed soon after.
- Buyers and sellers may place feedback ratings on the purchase or sale of an item. The feedback contains a rating between 1 and 10 and a comment. Note that the ratings are placed on a completed transaction by the buyer or seller of the item in the transaction.

Design an Entity-Relationship diagram for the auction database and enter the design using a data modeling tool such as ERWin.

8. Consider a database system for a baseball organization such as the major leagues. The data requirements are summarized as follows:

- The personnel involved in the league include players, coaches, managers, and umpires. Each is identified by a unique personnel id. They are also described by their first and last names along with the date and place of birth.
- Players are further described by other attributes such as their batting orientation (left, right, or switch) and have a lifetime batting average (BA).

- Within the players group is a subset of players called pitchers. Pitchers have a life time ERA (earned run average) associated with them.
- Teams are uniquely identified by their names. Teams are also described by the city in which they are located and the division and league in which they play (such as Central division of the American league).
- Teams have one manager, a number of coaches, and a number of players.
- Games are played between two teams with one designated as the home team and the other the visiting team on a particular date. The score (runs, hits, and errors) are recorded for each team. The team with more number of runs is declared the winner of the game.
- With each finished game, a winning pitcher and a losing pitcher are recorded. In case there is a save awarded, the save pitcher is also recorded.
- With each finished game, the number of hits (singles, doubles, triples, and home runs) obtained by each player is also recorded.

Design an Entity-Relationship diagram for the baseball database and enter the design using a data modeling tool such as ERWin.

9. Consider the ER diagram for the university database shown in Figure 4.9 of the Elmasri/Navathe text. Enter this design using a data modeling tool such as ERWin.
10. Consider the ER diagram for the small airport database shown in Figure 4.12 of the Elmasri/Navathe text. Enter this design using a data modeling tool such as ERWin.

## CHAPTER 2

### Interpreters for Abstract Query Languages

This chapter introduces Java-based interpreters for two abstract query languages: Relational Algebra (RA) and Domain Relational Calculus (DRC). The interpreters have been implemented using the parser generator tools JCup and JFlex. In order to use these interpreters, one needs to only download two jar files: `dbengine.jar` and `aql.jar` and include them in the Java CLASSPATH. The JCup libraries are included as part of the jar files and hence the only other software that is required to use the interpreters is a standard Java environment.

The system is simple to use and comes with a database engine that implements a set of basic relational algebraic operators. The interpreter reads a query from the terminal and performs the following three steps:

- (1) **Syntax Check:** The query is checked for any syntax errors. If there are any syntactic errors, the interpreter reports these to the terminal and waits to read another query; otherwise the interpreter proceeds to the second step.
- (2) **Semantics Check:** The syntactically correct query is checked for semantic errors including type mismatches, invalid column references, and invalid relation names. Safety checks are also made by the DRC interpreter. If there are any semantic errors or if the DRC query is unsafe, the interpreter reports these to the terminal and waits to read another query; otherwise the interpreter proceeds to the third step.
- (3) **Query Evaluation:** The query is evaluated using the primitives provided by the database engine and the results are displayed.

### 2.1 Creating the Database

Before the user can start using the interpreters, they must create a database against which they will submit queries. The database consists of several text files all stored within a directory. The directory is named after the database name. For example, to create a database identified with the name `db1` and containing two tables:

```
student(sid:integer,sname:varchar,phone:varchar,gpa:decimal)
skills(sid:integer,language:varchar)
```

a directory called `db1` should be created along with the following three files (one for the catalog description and the remaining two for the data for the two tables):

```
catalog.dat
STUDENT.dat
```

SKILLS.dat

The file names are case sensitive and should strictly follow the convention used, i.e. `catalog.dat` should be all lower case and the data files should be named after their relation name in upper case followed by the file suffix, `.dat`, in lower case.

The `catalog.dat` file contains the number of relations in the first line followed by the descriptions of each relation. The description of each relation begins with the name of the relation in a separate line followed by the number of attributes in a separate line followed by attribute descriptions. Each attribute description includes the name of the attribute in a separate line followed by the data type (`VARCHAR`, `INTEGER`, or `DECIMAL`) in a separate line. All names and data types are in upper case. There should be no leading or trailing white space in any of the lines. The `catalog.dat` file for database `db1` is shown below:

```
2
STUDENT
4
SID
INTEGER
SNAME
VARCHAR
PHONE
VARCHAR
GPA
DECIMAL
SKILLS
2
SID
INTEGER
LANGUAGE
VARCHAR
```

The `db1` directory must include one data file for each relation. In the case of `db1`, they should be named `STUDENT.dat` and `SKILLS.dat`. The data file for relations contains the number of tuples in the first line followed by the description of each tuple. Tuples are described by the values under each column with each value in a separate line. For example, let the `SKILLS` relation have three tuples:

```
(111, Java)
(111, C++)
(222, Java)
```

These tuples will be represented in the `SKILLS.dat` data file as follows:

```

3
111
Java
111
C++
222
Java

```

Again, there should be no leading or trailing white spaces in any of the lines. Some pre-defined databases are available along with this laboratory manual. New data may be added to existing databases as well as new databases may be created when needed.

## 2.2 Relational Algebra Interpreter

The RA interpreter is invoked using the following terminal command:

```
$ java edu.gsu.cs.ra.RA company
```

Here `$` is the command prompt and `company` is the name of the database (as well as the name of the directory where the database files are stored). This command assumes that the `company` directory is present in the same directory where this command is issued. Of course, one can issue this command in a different directory by providing the full path to the database directory.

The interpreter responds with the following prompt:

```
RA>
```

At this prompt the user may enter a Relational Algebra query or type the `exit` command. Each query must be terminated by a semi-colon. Even the `exit` command must end with a semi-colon. Queries may span more than one line; upon typing the `ENTER` key the interpreter prints the `RA>` prompt and waits for further input unless the `ENTER` key is typed after a semi-colon, in which case the query is processed by the interpreter.

### 2.2.1 Relational Algebra Syntax

A subset of Relational Algebra that includes the union, minus, intersect, Cartesian product, natural join, select, project, and rename operators is implemented in the interpreter. The context-free grammar for this subset is shown below:

```

<Query> ::= <Expr> SEMI;
<Expr>  ::= <ProjExpr>   | <RenameExpr>       | <UnionExpr> |

```

```

        <MinusExpr> | <IntersectExpr> | <JoinExpr> |
        <TimesExpr> | <SelectExpr>      | RELATION
<ProjExpr>    ::= PROJECT [<AttrList>] (<Expr>)
<RenameExpr> ::= RENAME [<AttrList>] (<Expr>)
<AttrList>   ::= ATTRIBUTE | <AttrList> , ATTRIBUTE
<UnionExpr>  ::= (<Expr> UNION <Expr>)
<MinusExpr>  ::= (<Expr> MINUS <Expr>)
<IntersectExpr> ::= (<Expr> INTERSECT <Expr>)
<JoinExpr>   ::= (<Expr> JOIN <Expr>)
<TimesExpr>  ::= (<Expr> TIMES <Expr>)
<SelectExpr> ::= SELECT [<Condition>] (<Expr>)
<Condition>  ::= <SimpleCondition> |
                <SimpleCondition> AND <Condition>
<SimpleCondition> ::= <Operand> <Comparison> <Operand>
<Operand>        ::= ATTRIBUTE | STRING-CONST | NUMBER-CONST
<Comparison>     ::= < | <= | = | <> | > | >=

```

The terminal strings in the grammar include

- Keywords for the relational algebraic operators: PROJECT, RENAME, UNION, MINUS, INTERSECT, JOIN, TIMES, and SELECT. These keywords are case-insensitive.
- Logical keyword AND (case-insensitive).
- Miscellaneous syntactic character strings such as (, ), <, <=, =, <>, >, >=, ,, and comma (,).
- Name strings: RELATION and ATTRIBUTE (case-insensitive names of relations and their attributes).
- Constant strings: STRING-CONST (a string enclosed within single quotes; e.g. 'Thomas') and NUMBER-CONST (integer as well as decimal numbers; e.g. 232 and -36.1).

An example of a well-formed syntactically correct query for the company database of the Elmasri/Navathe text is:

```

( project[ssn] (select[lname='Jones'] (employee))
  union
  project[superssn] (select[dno=5] (employee))
);

```

All relational algebra queries must be terminated by a semi-colon.

A relational algebra query in the simplest form is a “relation name”. For example the following terminal session with the interpreter illustrates the execution of this simple query form:

```

$ java edu.gsu.cs.ra.RA company
RA> departments;

```



```
SEMANTIC ERROR in RA Query: Relation DEPARTMENTS does not exist
RA> department;
DEPARTMENT (DNAME:VARCHAR,DNUMBER:INTEGER,MGRSSN:VARCHAR,MGRSTARTDATE:VARCHAR)
```

```
Number of tuples = 6
Research:5:333445555:22-MAY-1978:
Administration:4:987654321:01-JAN-1985:
Headquarters:1:888665555:19-JUN-1971:
Software:6:111111100:15-MAY-1999:
Hardware:7:444444400:15-MAY-1998:
Sales:8:555555500:01-JAN-1997:
```

```
RA> exit;
$
```

In response to a query, the interpreter displays the schema of the result followed by the answer to the query. Individual values within a tuple are terminated by a colon. The simplest query form is useful to display the database contents.

More complicated relational algebra queries involve one or more applications of one or more of the several operators such as select, project, times, join, union, etc. For example, consider the query “*Retrieve the names of all employees working for Dept. No. 5*”. This would be expressed by the query execution in the following RA session:

```
RA> project [fname,lname] (select [dno=5] (employee));
temp1 (FNAME:VARCHAR,LNAME:VARCHAR)
```

```
Number of tuples = 4
Franklin:Wong:
John:Smith:
Ramesh:Narayan:
Joyce:English:
```

```
RA>
```

### 2.2.2 Naming of Intermediate Relations and Attributes

The RA interpreter assigns temporary relation names such as temp0, temp1, etc. to each intermediate relation encountered in the execution of the entire query. The RA interpreter also employs the following rules as far as naming of attributes/columns of intermediate relations:

1. Union, Minus, and Intersect: The attribute/column names from the left operand are used to name the attributes of the output relation.
2. Times (Cartesian Product): Attribute/Column names from both operands are used to name the attributes of the output relation. Attribute/Column names that are common to both operands are prefixed by relation name (tempN).

3. **Select:** The attribute names of the output relation are the same as the attribute/column names of the operand.
4. **Project, Rename:** Attribute/Column names present in the attribute list parameter of the operator are used to name the attributes of the output relation. Duplicate attribute/column names are not allowed in the attribute list.
5. **Join (Natural Join):** Attribute/Column names from both operands are used to name the attributes of the output relation. Common attribute/column names appear only once.

As another example, consider the query “*Retrieve the social security numbers of employees who either work in department 5 or directly supervise an employee who works in department 5*”. The query is illustrated in the following RA session:

```
RA> (project[ssn] (select[dno=5] (employee)))
RA> union project[superssn] (select[dno=5] (employee))) ;
temp4 (SSN:VARCHAR)
```

Number of tuples = 5

```
333445555:
123456789:
666884444:
453453453:
888665555:
```

RA>

### 2.2.3 Relational Algebraic Operators Supported by the RA Interpreter

**Select:** As can be noted from the grammar, the select operator supported by the interpreter has the following syntax:

```
select[condition] (expression)
```

where `condition` is a conjunction of one or more simple conditions involving comparisons of attributes or constants with other attributes or constants. The attributes used in the condition must be present in the attributes of the relation corresponding to `expression`.

**Project:** The project operator supported by the interpreter has the following syntax:

```
project[attribute-list] (expression)
```

where `attribute-list` is a comma separated list of attributes, each of which is present in the attributes of the relation corresponding to `expression`.

**Rename:** The syntax for the rename operator is

```
rename[attribute-list] (expression)
```

where `attribute-list` is a comma separated list of attribute names. The number of attributes mentioned in the list must be equal to the number of attributes of the relation corresponding to `expression`.

**Join:** The syntax for the join operator is

```
(expression1 join expression2)
```

There is no restriction on the schemas of the two expressions.

**Times:** The syntax for the times operator is

```
(expression1 times expression2)
```

There is no restriction on the schemas of the two expressions.

**Union:** The syntax for the union operator is

```
(expression1 union expression2)
```

The schemas of the two expressions must be compatible (same number of attributes and same data types; the names of the attributes may be different).

**Minus:** The syntax for the minus operator is

```
(expression1 minus expression2)
```

The schemas of the two expressions must be compatible (same number of attributes and same data types; the names of the attributes may be different).

**Intersect:** The syntax for the intersect operator is

```
(expression1 intersect expression2)
```

The schemas of the two expressions must be compatible (same number of attributes and same data types; the names of the attributes may be different).

## 2.2.4 Examples

The queries from Section 6.5 of the Elmasri/Navathe text modified to work with the RA interpreter are shown below:

**Query 1:** Retrieve the name and address of employees who work for the "Research" department.

```

project[fname,lname,address] (
  (rename[dname,dno,mgrssn,mgrstartdate] (
    select[dname='Research'] (department))
  join
  employee
)
);

```

**Query 2:** For every project located in "Stafford", list the project number, the controlling department number, and the department manager's last name, address, and birth date.

```

project[pnumber,dnum,lname,address,bdate] (
  (
    (select[plocation='Stafford'] (projects)
    join
    rename[dname,dnum,ssn,mgrstartdate] (department)
  )
  join employee
)
);

```

**Query 3:** Find the names of employees who work on all the projects controlled by department number 5.

```

project[lname,fname] (
  (employee
  join
  (project[ssn] (employee)
  minus
  project[ssn] (
    (
      (project[ssn] (employee)
      times
      project[pnumber] (select[dnum=5] (projects))
    )
    minus
    rename[ssn,pnumber] (project[essn,pno] (works_on))
  )
  )
  )
);

```

**Query 4:** Make a list of project numbers for projects that involve an employee whose last name is "Smith", either as a worker or as a manager of the department that controls the project.

```

( project[pno] (
  (rename[essn] (project[ssn] (select[lname='Smith'] (employee)))
  join
  works_on
  )
  )
union
project[pnumber] (
  ( rename[dnum] (project[dnumber] (select[lname='Smith'] (
    (employee
    join
    rename[dname,dnumber,ssn,mgrstartdate] (department)
    )
    )
    )
  )
  join
  projects
  )
)
);

```

**Query 5:** List the names of all employees with two or more dependents.

```

project[lname,fname] (
  (rename[ssn] (
    project[essn1] (
      select[essn1=essn2 and dname1<>dname2] (
        (rename[essn1,dname1] (project[essn,dependent_name] (dependent))
        times
        rename[essn2,dname2] (project[essn,dependent_name] (dependent)))
      )
    )
  )
  join
  employee)
);

```

**Query 6:** Retrieve the names of employees who have no dependents.

```

project[lname,fname] (
  ( ( project[ssn] (employee)
    minus project[essn] (dependent)
    )
  join
  employee
  )
);

```

```
);
```

**Query 7:** List the names of managers who have at least one dependent.

```
project[lname, fname] (
  ((rename[ssn] (project[mgrssn] (department))
    join
    rename[ssn] (project[essn] (dependent))
  )
  join
  employee
)
);
```

**Important Tip:** Since many of the queries shown above are long and span multiple lines, the best way to use the interpreter is to create a text file in which the queries are typed. These queries are then cut and pasted into the interpreter prompt. Any errors in syntax or semantics should be corrected in the text file and then the process of cut and paste should be repeated until a correct solution is reached.

## 2.3 Domain Relational Calculus Interpreter

The DRC interpreter is invoked using the following terminal command:

```
$ java edu.gsu.cs.drc.DRC company
```

Here \$ is the command prompt and `company` is the name of the database (as well as the name of the directory where the database files are stored). This command assumes that the `company` directory is present in the same directory where this command is issued. Of course, one can issue this command in a different directory by providing the full path to the database directory.

The interpreter responds with the following prompt:

```
DRC>
```

At this prompt the user may enter a Domain Relational Calculus query or type the `exit` command. Each DRC query is expressed in a set-notation using a pair of curly brackets as follows:

```
{ variable-list | P(variable-list) }
```

where `variable-list` is a comma-separated list of variables which must all be present in the body predicate of the query `P(variable-list)` as free-variables.

The `exit` command must end with a semi-colon. Queries may span more than one line; upon typing the ENTER key the interpreter prints the `DRC>` prompt and waits for further input unless

the ENTER key is typed after a right curly bracket (}), in which case the query is processed by the interpreter.

### 2.3.1 Domain Relational Calculus Syntax

The context-free grammar for DRC queries implemented within the DRC interpreter is shown below:

```
Query ::= LBRACE VarList BAR Formula RBRACE;
VarList ::= NAME | VarList COMMA NAME;
Formula ::= AtomicFormula |
           Formula AND Formula |
           Formula OR Formula |
           NOT LPAREN Formula RPAREN |
           LPAREN EXISTS VarList RPAREN LPAREN Formula RPAREN |
           LPAREN FORALL VarList RPAREN LPAREN Formula RPAREN;
AtomicFormula ::=
           NAME LPAREN ArgList RPAREN | Arg Comparison Arg;
ArgList ::= Arg | ArgList COMMA Arg;
Arg ::= NAME | STRING | NUMBER;
Comparison ::= < | <= | = | <> | > | >=
```

The terminal strings in the grammar include

- Keywords for the logical operators: AND, OR, and NOT. These keywords are case-insensitive.
- Quantifier keywords EXISTS and FORALL (case-insensitive).
- Miscellaneous syntactic character strings such as (, ), <, <=, =, <>, >, >=, and comma (,).
- NAME strings: used for named relations and variables (case-insensitive).
- Constant strings: STRING (a string enclosed within single quotes; e.g. 'Thomas') and NUMBER (integer as well as decimal numbers; e.g. 232 and -36.1).

An example of a well-formed syntactically correct query on the company database of the Elmasri/Navathe text is:

```
{ x | (exists a1,a2,a3,a4,a5,a6,a7,a8) (
        employee(a1,a2,'Jones',x,a3,a4,a5,a6,a7,a8)) or
      (exists a1,a2,a3,a4,a5,a6,a7,a8) (
        employee(a1,a2,a3,x,a4,a5,a6,a7,a8,5)) }
```

All DRC queries must be enclosed within a pair of matching curly brackets.

A DRC query in the simplest form displays the contents of a relation. For example the following terminal session with the interpreter illustrates the execution of this simple query form that displays the contents of the DEPARTMENT relation:

```
$ java edu.gsu.cs.drc.DRC company
DRC> { a,b,c,d | department(a,b,c,d) }
ANSWER (A:VARCHAR,B:INTEGER,C:VARCHAR,D:VARCHAR)
```

```
Number of tuples = 6
Research:5:333445555:22-MAY-1978:
Administration:4:987654321:01-JAN-1985:
Headquarters:1:888665555:19-JUN-1971:
Software:6:111111100:15-MAY-1999:
Hardware:7:444444400:15-MAY-1998:
Sales:8:555555500:01-JAN-1997:
```

```
DRC>
```

In response to a query, the interpreter displays the schema of the result followed by the answer to the query. Individual values are terminated by a colon.

### 2.3.2 Safe DRC Queries

The DRC interpreter checks for the “safety” of queries and evaluates only those that are determined to be safe. An error message is generated for unsafe queries.

For the discussion of safe DRC queries, we will assume that the formula defining the query does not contain the `forall` quantifier. If the `forall` quantifier does appear in the formula, the user can convert such a formula to an equivalent one without the `forall` quantifier using the logical equivalence:

$$(\text{forall } X) (F) \equiv \text{NOT } ((\text{exists } X) (\text{NOT } (F)))$$

It is almost always the case that the `F` in the `forall` quantified formula above is of the form

$$\text{NOT } (P) \text{ or } Q$$

In case the user does not eliminate the `forall` quantifier, the DRC interpreter would automatically convert all `forall` quantified formulas into equivalent `exists` quantified formulas using the above equivalence. In addition, the interpreter would also apply the DeMorgan’s law:

$$\text{NOT } (P \text{ or } Q) \equiv \text{NOT } (P) \text{ and } \text{NOT } (Q)$$

to push the `NOT` further inside the formula.

As an example of this automatic transformation, consider the following query provided by the user:



```
{a,b | (exists c) (r(a,b,c) and
              (forall d,e) (not(s(a,d,e)) or (exists f) (t(d,f)))) }
```

The DRC interpreter would convert the above query to:

```
{a,b | (exists c) (r(a,b,c) and
              not(exists d,e) (s(a,d,e) and not(exists f) (t(d,f)))) }
```

**Definition:** A DRC query (without forall quantifiers) is defined to be *safe* if it satisfies the following three conditions:

- (a) for every subformula in the query connected with an “or”, the two operand formulas have the same set of free variables, i.e. the “or” formula is of the form:

$$F(X_1, \dots, X_n) \text{ or } G(X_1, \dots, X_n)$$

- (b) All free variables appearing in “maximal sub-conjuncts”,  $F_1$  and ... and  $F_n$ , must be “limited” in that they either appear in (i) a positive sub-formula  $F_i$  or (ii) as  $X$  in an sub-formula of the form  $X=a$  or  $a=X$  or (iii) as  $X$  in a sub-formula of the form  $X=Y$  where  $Y$  is determined to be “limited”.
- (c) The NOT operator may be applied only to a term in a maximal sub-conjunct of type discussed in (b), i.e. all free variables in the NOT term must be shown to be “limited” in the positive terms of the maximal sub-conjunct.

Some examples follow. The following query would be considered safe as it satisfies condition (a).

```
{a,b | (exists c) (r(a,b,c)) or s(a,b) }
```

But the following would not be safe:

```
{a,b | (exists b,c) (r(a,b,c)) or s(a,b) }
```

because the free variables on the left operand of the “or” formula consists of only one variable,  $a$ , and the free variables on the right operand consists of two variables,  $a$  and  $b$ .

The query formula from an earlier query:

```
(exists c) (r(a,b,c) and
            not(exists d,e) (s(a,d,e) and not(exists f) (t(d,f))))
```

is safe. The formula has the following two maximal sub-conjuncts (ignoring atomic formulas which are maximal sub-conjuncts of size 1):

- (1)  $s(a, d, e)$  and  $\text{not}(\text{exists } f)(t(d, f))$   
 all three free variables  $a$ ,  $d$ , and  $e$  are limited as they appear in  $s(a, d, e)$
- (2)  $r(a, b, c)$  and  
 $\text{not}(\text{exists } d, e)(s(a, d, e) \text{ and } \text{not}(\text{exists } f)(t(d, f)))$   
 all three free variables  $a$ ,  $b$ , and  $c$  are limited as they appear in  $r(a, b, c)$ .

The free variables in each of the maximal sub-conjuncts are shown to be “limited” and hence the overall query is safe.

The following query formula is unsafe:

$p(a, b) \text{ and } \text{not}((\text{exists } c)(q(b, c, d)))$

since the free variable  $d$  is not “limited” as it is not grounded in a positive term in the maximal sub-conjunct.

### 2.3.3 DRC Query Examples

The queries from Section 6.7 of the Elmasri/Navathe text modified to work with DRC interpreter are shown below:

**Query 0:** Retrieve the birthdate and address of the employees whose name is "John B. Smith".

```
{ u, v | (exists t, w, x, y, z) (
    employee('John', 'B', 'Smith', t, u, v, w, x, y, z) ) }
```

**Query 1:** Retrieve the name and address of all employees who work for the "Research" department.

```
{ q, s, v | (exists r, t, u, w, x, y, z, n, o) (
    employee(q, r, s, t, u, v, w, x, y, z) and
    department('Research', z, n, o) ) }
```

**Query 2:** For every project located in "Stafford", list the project number, the controlling department number, and the department manager's last name, birth date, and address.

```
{ i, k, s, u, v | (exists h, q, r, t, w, x, y, z, l, o) (
    projects(h, i, 'Stafford', k) and
    employee(q, r, s, t, u, v, w, x, y, z) and
    department(l, k, t, o) ) }
```

**Query 6:** List the names of employees who have no dependents.

```
{ q, s | (exists r, t, u, v, w, x, y, z) (
```

```
employee(q,r,s,t,u,v,w,x,y,z) and
not ((exists m,n,o,p) (dependent(t,m,n,o,p))) ) }
```

The following is not SAFE and would not work

```
{ q,s | (exists r,t,u,v,w,x,y,z) (
    employee(q,r,s,t,u,v,w,x,y,z) and
    (forall l,m,n,o,p) (not (dependent(l,m,n,o,p)) or t<>l) ) }
```

**Query 7:** List the names of managers who have at least one dependent.

```
{ s,q | (exists r,t,u,v,w,x,y,z,h,i,k,m,n,o,p) (
    employee(q,r,s,t,u,v,w,x,y,z) and
    department(h,i,t,k) and
    dependent(t,m,n,o,p) ) }
```

## Exercises

1. Specify and execute the following queries using the RA interpreter on the COMPANY database schema.
  - a. Retrieve the names of all employees in department 5 who work more than 10 hours per week on the 'ProductX' project.
  - b. List the names of all employees who have a dependent with the same first name as themselves.
  - c. Find the names of employees who are directly supervised by 'Franklin Wong'.
  - d. Retrieve the names of employees who work on every project.
  - e. Retrieve the names of employees who do not work on any project.
  - f. Retrieve the names and addresses of all employees who work on at least one project located in Houston but whose department has no location in Houston.
  - g. Retrieve the last names of all department managers who have no dependents.
2. Consider the following MAILORDER relational schema describing the data for a mail order company:

```
parts(pno,pname,qoh,price,olevel)
customers(cno,cname,street,zip,phone)
employees(eno,ename,zip,hdate)
zipcodes(zip,city)
orders(ono,cno,eno,received,shipped)
odetails(ono,pno,qty)
```

The attribute names are self-explanatory. "qoh" stands for quantity on hand. Specify and execute the following queries using the RA interpreter on the MAILORDER database schema.

- a. Retrieve the names of parts that cost less than \$20.00.
  - b. Retrieve the names and cities of employees who have taken orders for parts costing more than \$50.00
  - c. Retrieve the pairs of customer number values of customers who live in the same zip code.
  - d. Retrieve the names of customers who have ordered parts only from employees living in the city of Wichita.
  - e. Retrieve the names of customers who have ordered all parts costing less than \$20.00.
  - f. Retrieve the names of customers who have not placed a single order.
  - g. Retrieve the names of customers who have placed exactly two orders.
3. Consider the following GRADEBOOK relational schema describing the data for a grade book of a particular instructor (Note: The attributes A, B, C, and D store grade cutoffs.)

```
catalog(cno, ctitle)
students(sid, fname, lname, minit)
courses(term, secno, cno, A, B, C, D)
enrolls(sid, term, secno)
```

Specify and execute the following queries using the RA interpreter on the GRADEBOOK database schema.

- a. Retrieve the names of students enrolled in the 'Automata' class in the term of Fall 1996.
  - b. Retrieve the SID values of students who have enrolled in CSc226 as well as CSc227.
  - c. Retrieve the SID values of students who have enrolled in CSc226 or CSc227.
  - d. Retrieve the names of students who have not enrolled in any class.
  - e. Retrieve the names of students who have enrolled in all courses in the `catalog` table.
4. Consider the database consisting of the following relations:

```
supplier(sno, sname)
part(pno, pname)
project(jno, jname)
supply(sno, pno, jno)
```

The database records information about suppliers, parts, and projects and includes a ternary relationship between suppliers, parts, and projects. This relationship is a many-many-many relationship. Specify and execute the following queries using the RA interpreter.

- a. Retrieve part numbers of parts that are supplied to exactly two projects.
- b. Retrieve supplier names of suppliers who supply more than two parts to project 'JI'.
- c. Retrieve part numbers of parts that are supplied by every supplier.
- d. Retrieve project names of projects that are supplied only by suppliers 'SI'.

- e. Retrieve supplier names of suppliers who supply at least two different parts each to at least two different projects.
5. Specify and execute the following queries for the database in Exercise 5.16 of the Elmasri/Navathe text using the RA interpreter.
- a. Retrieve the names of students who have enrolled in a course that uses a textbook published by Addison Wesley.
  - b. Retrieve the course names of courses which have changed their textbook at least once.
  - c. Retrieve the names of departments that only adopt textbooks from Addison Wesley.
  - d. Retrieve the names of departments that have adopted all textbooks written by Navathe and published by Addison Wesley in their courses.
  - e. Retrieve the names of students who have never used a book (in a course) written by Navathe and published by Addison Wesley.
6. Repeat Exercises 1 through 5 in domain relational calculus (DRC) by using the DRC interpreter.

## CHAPTER 3

### Relational Database Management Systems: Oracle™

This chapter introduces the student to the basic utilities used to interact with Oracle DBMS. The chapter also introduces the student to programming applications in Java and JDBC that interact with Oracle databases. The discussion in this chapter is not specific to any version of Oracle and all examples would work with any version of Oracle higher than Oracle 8.

The company database of the Elmasri/Navathe text is used throughout this chapter. In Section 3.1 a larger data set is introduced for the company database. In Section 3.2 and 3.3, the SQL\*Plus and the SQL\*Loader utilities are introduced using which the student can create database tables and load the tables with data. Finally, in Section 3.4, programming in Java using the JDBC API and connecting to Oracle databases is introduced through three complete application programs.

#### 3.1 COMPANY Database

Consider the COMPANY database state shown in Figure 5.6 of the Elmasri/Navathe text. Let us assume that the company is expanding with 3 new departments: Software, Hardware, and Sales. The Software department has 2 locations: Atlanta and Sacramento, the Hardware department is located in Milwaukee, and the Sales department has 5 locations: Chicago, Dallas, Philadelphia, Seattle, and Miami.

The Software department has 3 new projects: OperatingSystems, DatabaseSystems, and Middleware and the Hardware department has 2 new projects: InkjetPrinters and LaserPrinters. The company has added 32 new employees in this expansion process.

The updated COMPANY database is shown in the following tables.

DEPARTMENT	DNAME	DNUMBER	MGRSSN	MGRSTARTDATE
	Research	5	333445555	22-May-78
	Administration	4	987654321	1-Jan-85
	Headquarters	1	888665555	19-Jun-71
	Software	6	111111100	15-May-99
	Hardware	7	444444400	15-May-98
	Sales	8	555555500	1-Jan-97

PROJECT	PNAME	PNUMBER	PLOCATION	DNUM
	ProductX	1	Bellaire	5
	ProductY	2	Sugarland	5
	ProductZ	3	Houston	5
	Computerization	10	Stafford	4
	Reorganization	20	Houston	1
	Newbenefits	30	Stafford	4
	OperatingSystems	61	Jacksonville	6
	DatabaseSystems	62	Birmingham	6
	Middleware	63	Jackson	6
	InkjetPrinters	91	Phoenix	7
	LaserPrinters	92	LasVegas	7

DEPT_LOCATIONS	DNUMBER	DLOCATION
	1	Houston
	4	Stafford
	5	Bellaire
	5	Sugarland
	5	Houston
	6	Atlanta
	6	Sacramento
	7	Milwaukee
	8	Chicago
	8	Dallas
	8	Philadephia
	8	Seattle

DEPENDENT	ESSN	DEPENDENT_NAME	SEX	BDATE	RELATION
	333445555	Alice	F	5-Apr-76	Daughter
	333445555	Theodore	M	25-Oct-73	Son
	333445555	Joy	F	3-May-48	Spouse
	987654321	Abner	M	29-Feb-32	Spouse
	123456789	Michael	M	1-Jan-78	Son
	123456789	Alice	F	31-Dec-78	Daughter
	123456789	Elizabeth	F	5-May-57	Spouse
	444444400	Johnny	M	4-Apr-97	Son
	444444400	Tommy	M	7-Jun-99	Son
	444444401	Chris	M	19-Apr-69	Spouse
	444444402	Sam	M	14-Feb-64	Spouse

WORKS_ON	ESSN	PNO	HOURS
	123456789	1	32.5
	123456789	2	7.5
	666884444	3	40
	453453453	1	20
	453453453	2	20
	333445555	2	10
	333445555	3	10
	333445555	10	10
	333445555	20	10
	999887777	30	30
	999887777	10	10
	987987987	10	35
	987987987	30	5
	987654321	30	20
	987654321	20	15
	888665555	20 null	
	111111100	61	40
	111111101	61	40
	111111102	61	40
	111111103	61	40
	222222200	62	40
	222222201	62	48
	222222202	62	40
	222222203	62	40
	222222204	62	40
	222222205	62	40
	333333300	63	40
	333333301	63	46
	444444400	91	40
	444444401	91	40
	444444402	91	40
	444444403	91	40
	555555500	92	40
	555555501	92	44
	666666601	91	40
	666666603	91	40
	666666604	91	40
	666666605	92	40
	666666606	91	40
	666666607	61	40
	666666608	62	40
	666666609	63	40
	666666610	61	40
	666666611	61	40
	666666612	61	40
	666666613	61	30
	666666613	62	10
	666666613	63	10



## EMPLOYEE

FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
James	E	Borg	888665555	10-Nov-27	450 Stone, Houston, TX	M	55000	null	1
Franklin	T	Wong	333445555	8-Dec-45	638 Voss, Houston, TX	M	40000	888665555	5
Jennifer	S	Wallace	987654321	20-Jun-31	291 Berry, Bellaire, TX	F	43000	888665555	4
John	B	Smith	123456789	9-Jan-55	731 Fondren, Houston, TX	M	30000	333445555	5
Alicia	J	Zelaya	999887777	19-Jul-58	3321 Castle, Spring, TX	F	25000	987654321	4
Ramesh	K	Narayan	666884444	15-Sep-52	971 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	31-Jul-62	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	29-Mar-59	980 Dallas, Houston, TX	M	25000	987654321	4
Jared	D	James	111111100	10-Oct-66	123 Peachtree, Atlanta, GA	M	85000	null	6
Alex	D	Freed	444444400	9-Oct-50	4333 Pillsbury, Milwaukee, WI	M	89000	null	7
John	C	James	555555500	30-Jun-75	7676 Bloomington, Sacramento, CA	M	81000	null	6
Jon	C	Jones	111111101	14-Nov-67	111 Allgood, Atlanta, GA	M	45000	111111100	6
Justin	null	Mark	111111102	12-Jan-66	2342 May, Atlanta, GA	M	40000	111111100	6
Brad	C	Knight	111111103	13-Feb-68	176 Main St., Atlanta, GA	M	44000	111111100	6
Evan	E	Wallis	222222200	16-Jan-58	134 Pelham, Milwaukee, WI	M	92000	null	7
Josh	U	Zell	222222201	22-May-54	266 McGrady, Milwaukee, WI	M	56000	222222200	7
Andy	C	Vile	222222202	21-Jun-44	1967 Jordan, Milwaukee, WI	M	53000	222222200	7
Tom	G	Brand	222222203	16-Dec-66	112 Third St, Milwaukee, WI	M	62500	222222200	7
Jenny	F	Vos	222222204	11-Nov-67	263 Mayberry, Milwaukee, WI	F	61000	222222201	7
Chris	A	Carter	222222205	21-Mar-60	565 Jordan, Milwaukee, WI	F	43000	222222201	7
Kim	C	Grace	333333300	23-Oct-70	6677 Mills Ave, Sacramento, CA	F	79000	null	6
Jeff	H	Chase	333333301	7-Jan-70	145 Bradbury, Sacramento, CA	M	44000	333333300	6
Bonnie	S	Bays	444444401	19-Jun-56	111 Hollow, Milwaukee, WI	F	70000	444444400	7
Alec	C	Best	444444402	18-Jun-66	233 Solid, Milwaukee, WI	M	60000	444444400	7
Sam	S	Snedden	444444403	31-Jul-77	987 Windy St, Milwaukee, WI	M	48000	444444400	7
Nandita	K	Ball	555555501	16-Apr-69	222 Howard, Sacramento, CA	M	62000	555555500	6
Bob	B	Bender	666666600	17-Apr-68	8794 Garfield, Chicago, IL	M	96000	null	8
Jill	J	Jarvis	666666601	14-Jan-66	6234 Lincoln, Chicago, IL	F	36000	666666600	8
Kate	W	King	666666602	16-Apr-66	1976 Boone Trace, Chicago, IL	F	44000	666666600	8
Lyle	G	Leslie	666666603	9-Jun-63	417 Hancock Ave, Chicago, IL	M	41000	666666601	8
Billie	J	King	666666604	1-Jan-60	556 Washington, Chicago, IL	F	38000	666666603	8
Jon	A	Kramer	666666605	22-Aug-64	1988 Windy Creek, Seattle, WA	M	41500	666666603	8
Ray	H	King	666666606	16-Aug-49	213 Delk Road, Seattle, WA	M	44500	666666604	8
Gerald	D	Small	666666607	15-May-62	122 Ball Street, Dallas, TX	M	29000	666666602	8
Arnold	A	Head	666666608	19-May-67	233 Spring St, Dallas, TX	M	33000	666666602	8
Helga	C	Pataki	666666609	11-Mar-69	101 Holyoke St, Dallas, TX	F	32000	666666602	8
Naveen	B	Drew	666666610	23-May-70	198 Elm St, Philadelphia, PA	M	34000	666666607	8
Carl	E	Reedy	666666611	21-Jun-77	213 Ball St, Philadelphia, PA	M	32000	666666610	8
Sammy	G	Hall	666666612	11-Jan-70	433 Main Street, Miami, FL	M	37000	666666611	8

### 3.2 SQL\*Plus Utility

Oracle provides a command line utility, called SQL\*Plus, that allows the user to execute SQL statements interactively. The user may enter the statements directly on the SQL\*Plus prompt or may save the statements in a file and have them executed by entering the “start” command. We shall now look at how one can use this utility to create the database tables for the COMPANY database.

#### Creating Database Tables

Let us assume that all the SQL statements to create the COMPANY database tables are present in a file called `company_schema.sql`. The part that defines the DEPARTMENT table is shown below:

```
CREATE TABLE department (
    dname          varchar(25),
    dnumber        number(4),
    mgrssn         char(9),
    mgrstartdate   date,
    primary key (dnumber),
    foreign key (mgrssn) references employee
);
```

The following SQL\*Plus session illustrates how these statements would be executed resulting in the creation of the COMPANY database tables:

```
$ sqlplus
SQL>start company_schema

Table created

SQL>exit
$
```

The \$ symbol used above is the command line prompt on a Unix system or a Windows command line interface.

In general, when several tables are defined in a single file, these definitions should be arranged in a particular order of foreign key references. For example, the definition of the WORKS\_ON table should follow the definition of the EMPLOYEE table and the PROJECT table since it has foreign keys referencing the EMPLOYEE and the PROJECT tables.

Sometimes, it is possible for two tables to have a circular reference between them. For example, the foreign key constraints defined for the EMPLOYEE and the DEPARTMENT tables indicate a mutual dependency between the two tables. The EMPLOYEE table contains an attribute DNO which

refers to the DNUMBER attribute of the DEPARTMENT table and the MGRSSN attribute of the DEPARTMENT table refers to the SSN attribute of the EMPLOYEE table. This mutual dependency causes some trouble while creating the tables as well as loading the data. To avoid these problems, it may be simplest to omit the forward references, i.e. omit the DNO foreign key in the EMPLOYEE table.

A more appropriate way to handle the circular reference problem is as follows: Create the tables by omitting the forward references. After the tables are created, use the ALTER TABLE statement to add the omitted reference. In the case of the EMPLOYEE table, use the following ALTER command to add the forward reference after the DEPARTMENT table is created:

```
ALTER TABLE employee ADD (
    foreign key (dno) references department(dnumber)
);
```

The SQL\*Plus utility is very versatile program and allows the user to submit any SQL statement for execution. In addition, it allows users to format query results in different ways. Please consult the SQL\*Plus documentation that comes with any Oracle installation for further details on its use.

### 3.3 SQL\*Loader Utility

Oracle provides a data loading utility called SQL\*Loader<sup>2</sup> (sqlldr) that allows the user to populate the tables with data. To use sqlldr, we need a control file that indicates how that data is to be loaded and a data file that contains the data to be loaded. A typical control file is shown below:

```
LOAD DATA
INFILE <dataFile>
APPEND INTO TABLE <tableName>
FIELDS TERMINATED BY '<separator>'
(<attribute-list>)
```

Here, <datafile> is the name of the file containing the data, <tableName> is the database table into which the data is to be loaded, <separator> is a string that does not appear as part of the data, and <attribute-list> is the list of attributes or columns under which data is to be loaded. The attribute list need not be the entire list of columns and the attributes need not appear in the order in which they were defined in the create table statement,

As a concrete example, the following control file (department.ctl) will load data into the department table:

---

<sup>2</sup> For more details on this utility, please refer to Oracle's online documentation.

```

LOAD DATA
INFILE department.csv
APPEND INTO TABLE department
FIELDS TERMINATED BY ',' OPTIONALLY ENCLOSED BY '"'
TRAILING NULLCOLS
(dname,dnumber,mgrssn,mgrstartdate DATE 'yyyy-mm-dd')

```

The `OPTIONALLY ENCLOSED` phrase is useful in situations when one of the data fields contained the separator symbol. In this case, the entire data value would be enclosed with a pair of double quotes. The `TRAILING NULLCOLS` phrase is useful in case there is a null value for the last data field. Null values are indicated by an empty string followed by the separator symbol. The corresponding data file, `department.csv` is:

```

Research,5,333445555,1978-05-22
Administration,4,987654321,1985-01-01
Headquarters,1,888665555,1971-06-19
Software,6,111111100,1999-05-15
Hardware,7,444444400,1998-05-15
Sales,8,555555500,1997-01-01

```

The following command will load the data into the department table:

```
$ sqlldr OracleId/OraclePassword control=department.ctl
```

Here, `OracleId/OraclePassword` is the Oracle user id/password. Upon execution of this command, a log file under the name `department.log` is created by `sqlldr`. This log file contains information such as the number of rows successfully loaded into the table, and any errors encountered while loading the data.

The loading of the data can be verified by the following `sqlplus` session:

```
$ sqlplus OracleId
```

```
SQL*Plus: Release 9.2.0.1.0 - Production on Wed Apr 6 20:35:45
2005
```

```
Copyright (c) 1982, 2002, Oracle Corporation. All rights
reserved.
```

```
Enter password:
```

```

Connected to:
Oracle9i Enterprise Edition Release 9.2.0.1.0 - 64bit Production
With the Partitioning, OLAP and Oracle Data Mining options
JServer Release 9.2.0.1.0 - Production

```

```
SQL> select * from department;
```

DNAME	DNUMBER	MGRSSN	MGRSTARTD
Research	5	333445555	22-MAY-78
Administration	4	987654321	01-JAN-85
Headquarters	1	888665555	19-JUN-71
Software	6	111111100	15-MAY-99
Hardware	7	444444400	15-MAY-98
Sales	8	555555500	01-JAN-97

```
6 rows selected.
```

```
SQL> exit
```

```
Disconnected from Oracle9i Enterprise Edition Release 9.2.0.1.0 -
64bit Production
With the Partitioning, OLAP and Oracle Data Mining options
JServer Release 9.2.0.1.0 - Production
$
```

It is possible to load the data without a separate data file by using the following control file:

```
LOAD DATA
INFILE *
APPEND INTO TABLE department
FIELDS TERMINATED BY ','
(dname,dnumber,mgrssn,mgrstartdate DATE 'yyyy-mm-dd')
BEGINDATA
Research,5,333445555,1978-05-22
Administration,4,987654321,1985-01-01
Headquarters,1,888665555,1971-06-19
Software,6,111111100,1999-05-15
Hardware,7,444444400,1998-05-15
Sales,8,555555500,1997-01-01
```

Note that the filename has been replaced by a “\*” in the INFILE clause of the control file and the data is separated from the control information by the line containing the string BEGINDATA.

In the case of foreign key dependencies between tables, data for the tables that are being referenced should be loaded/created first before the data for the tables that refer to values in other tables. For example, the EMPLOYEE and PROJECT table data should be loaded before the data for the WORKS\_ON table is loaded.

The circular references are trickier to handle as far as bulk loading is concerned. To successfully load such data, we may use NULL values for forward references and once the referenced data is loaded, we may use the UPDATE statement to add values for the forward references.

### 3.4 Programming with Oracle using the JDBC API

This section presents Oracle JDBC programming through three complete applications. The first example illustrates basic query processing via `PreparedStatement` object in Java. The query result has at most one answer. The second example illustrates basic query processing using `Statement` object. In this example, the query may have more than one answer. The final example is a more involved one in which recursive query as well as aggregate query processing is discussed.

To be able to program with Oracle databases using Java/JDBC, one must have access to an Oracle database installation with a listener program for JDBC connections. One must also have access to the JDBC drivers (`classes12.zip` or `classes12.jar`) that comes with any standard Oracle distribution. The driver archive file must be included in the Java `CLASSPATH` variable. Any standard Java environment is sufficient to compile and run these programs.

The `COMPANY` database of the Elmasri/Navathe text is used in each of the three examples that follow.

**Example 1:** A simple Java/JDBC program that prints the last name and salary of an employee given his or her social security number is shown below:

```
import java.sql.*;
import java.io.*;

class getEmpInfo {
    public static void main (String args [])
        throws SQLException, IOException {

        // Load Oracle's JDBC Driver
        try {
            Class.forName ("oracle.jdbc.driver.OracleDriver");
        } catch (ClassNotFoundException e) {
            System.out.println ("Could not load the driver");
        }

        //Connect to the database
        String user, pass;
        user = readEntry("userid  : ");
        pass = readEntry("password: ");
        Connection conn =
            DriverManager.getConnection
                ("jdbc:oracle:thin:@tinman.cs.gsu.edu:1521:sid9ir2",
                user,pass);

        //Perform query using PreparedStatement object
        //by providing SSN at run time
```

```

String query =
    "select LNAME,SALARY from EMPLOYEE where SSN = ?";
PreparedStatement p = conn.prepareStatement (query);
String ssn = readEntry("Enter a Social Security Number: ");
p.clearParameters();
p.setString(1,ssn);
ResultSet r = p.executeQuery();

//Process the ResultSet
if (r.next ()) {
    String lname = r.getString(1);
    double salary = r.getDouble(2);
    System.out.println(lname + " " + salary);
}

//Close objects
p.close();
conn.close();
}

//readEntry function -- to read input string
static String readEntry(String prompt) {
    try {
        StringBuffer buffer = new StringBuffer();
        System.out.print(prompt);
        System.out.flush();
        int c = System.in.read();
        while(c != '\n' && c != -1) {
            buffer.append((char)c);
            c = System.in.read();
        }
        return buffer.toString().trim();
    } catch (IOException e) {
        return "";
    }
}
}

```

In the above program, the user is prompted for a database user id and password. The program uses this information to connect<sup>3</sup> to the database. The user is then prompted for the social security number of an employee. The program supplies this social security number as a parameter to an SQL query used by a PreparedStatement object. The query is executed and the resulting

---

<sup>3</sup> The connect string will be different in different installations of Oracle. Please consult your instructor/administrator for the exact connect string used in getConnection method.

ResultSet is then processed and the last name and salary of the employee is printed to the console.

The JDBC API is a very simple API to learn as there are a few important classes and methods that are required to manipulate relational data. After the driver is loaded and a connection is made (quite standard code to do this), queries and updates are submitted to the database via one of three objects: Statement, PreparedStatement, and CallableStatement.

The PreparedStatement method is illustrated in Example 1 where the SQL statement is sent to the database server with possible place holders for pluggable values for compilation using the prepareStatement method of the Connection object. After this is done, the same SQL statement may be executed a number of times with values provided for the place holder variables using the setString or similar methods. In Example 1, the prepared statement is executed only once.

The Statement method is more commonly used and is illustrated in the following example.

**Example 2:** Given a department number, the following Java/JDBC program prints the last name and salary of all employees working for the department.

```
import java.sql.*;
import java.io.*;

class printDepartmentEmps {
    public static void main (String args [])
        throws SQLException, IOException {

        //Load Oracle's JDBC Driver
        try {
            Class.forName ("oracle.jdbc.driver.OracleDriver");
        } catch (ClassNotFoundException e) {
            System.out.println ("Could not load the driver");
        }

        //Connect to the database
        String user, pass;
        user = readEntry("userid : ");
        pass = readEntry("password: ");
        Connection conn =
            DriverManager.getConnection
                ("jdbc:oracle:thin:@tinman.cs.gsu.edu:1521:sid9ir2",
                user,pass);

        //Perform query using Statement object
        String dno = readEntry("Enter a Department Number: ");
        String query =
```



```

        "select LNAME,SALARY from EMPLOYEE where DNO = " + dno;
Statement s = conn.createStatement();
ResultSet r = s.executeQuery(query);

//Process ResultSet
while (r.next ()) {
    String lname = r.getString(1);
    double salary = r.getDouble(2);
    System.out.println(lname + "    " + salary);
}

//Close objects
s.close();
conn.close();
}

```

The query is executed via a `Statement` object by including the department number as part of the query string at run time (using string concatenation). This is in contrast to the `PreparedStatement` method used in Example 1. The main difference between the two methods is that in the `PreparedStatement` method, the query string is first sent to the database server for syntax checking and then for execution subsequently. This method may be useful when the same query string is executed a number of times in a program with only a different parameter each time. On the other hand, in the `Statement` method, syntax checking and execution of the query happens at the same time.

The third method for executing SQL statements is to use the `CallableStatement` object. This is useful only in situations where the Java program needs to call a stored procedure or function. To learn about this method, the reader is referred to any standard JDBC textbook or the Oracle 9i Programming: A Primer (2004) published by Addison Wesley.

**Example 3:** In this next program, the user is presented with a menu of 3 choices:

- (a) **Find supervisees at all levels:** In this option, the user is prompted for the last name of an employee. If there are several employees with the same last name, the user is presented with a list of social security numbers of employees with the same last name and asked to choose one. The program then proceeds to list all the supervisees of the employee and all levels below him or her in the employee hierarchy.
- (b) **Find the top 5 highest paid employees:** In this option, the program finds five employees who rank in the top 5 in salary and lists them.
- (c) **Find the top 5 highest worked employees:** In this option, the program finds five employees who rank in the top 5 in number of hours worked and lists them.

A sample interaction with the user is shown below:

```

$ java example3
Enter userid: book

```

Enter password: book

QUERY OPTIONS

- (a) Find Supervisees at all levels.
- (b) Find Highest paid workers.
- (c) Find the most worked workers.
- (q) Quit.

Type in your option: a

Enter last name of employee : King

King, Kate 666666602

King, Billie 666666604

King, Ray 666666606

Select ssn from list : 666666602

SUPERVISEES

FNAME	LNAME	SSN
-----		
Gerald	Small	666666607
Arnold	Head	666666608
Helga	Pataki	666666609
Naveen	Drew	666666610
Carl	Reedy	666666611
Sammy	Hall	666666612
Red	Bacher	666666613

QUERY OPTIONS

- (a) Find Supervisees at all levels.
- (b) Find Highest paid workers.
- (c) Find the most worked workers.
- (q) Quit.

Type in your option: b

HIGHEST PAID WORKERS

-----			
666666600	Bob	Bender	96000.0
222222200	Evan	Wallis	92000.0
444444400	Alex	Freed	89000.0
111111100	Jared	James	85000.0
555555500	John	James	81000.0

QUERY OPTIONS

- (a) Find Supervisees at all levels.
- (b) Find Highest paid workers.

- (c) Find the most worked workers.
- (q) Quit.

Type in your option: c

#### MOST WORKED WORKERS

```
-----
666666613 Red          Bacher          50.0
222222201 Josh         Zell             48.0
333333301 Jeff         Chase            46.0
555555501 Nandita      Ball             44.0
111111100 Jared        James            40.0
```

#### QUERY OPTIONS

- (a) Find Supervisees at all levels.
- (b) Find Highest paid workers.
- (c) Find the most worked workers.
- (q) Quit.

Type in your option: q

\$

The program for option (a) is discussed now. Finding supervisees at all levels below an employee is a recursive query in which the data tree needs to be traversed from the employee node all the way down to the leaves in the sub-tree. One common strategy to solve this problem is to use a temporary table of social security numbers. Initially, this temporary table will store the next level supervisees. Subsequently, in an iterative manner, the supervisees at the “next” lower level will be computed with a query that involves the temporary table. These next supervisees are then added to the temporary table. This iteration is continued as long as there are new social security numbers added to the temporary table in any particular iteration. Finally, when no new social security numbers are found, the iteration is stopped and the names and social security numbers of all supervisees in the temporary table are listed.

The section of the main program that (a) creates the temporary table, (b) calls the `findSupervisees` method, and (c) drops the temporary table is shown below:

```
/* create new temporary table called tempSSN */
String sqlString = "create table tempSSN (" +
                    "ssn char(9) not null, " +
                    "primary key(ssn))";
Statement stmt1 = conn.createStatement();
try {
    stmt1.executeUpdate(sqlString);
} catch (SQLException e) {
    System.out.println("Could not create tempSSN table");
}
```

```

        stmt1.close();
        return;
    }

...
...
        case 'a': findSupervisees(conn);
...
...
    /* drop table called tmpSSN */
    sqlString = "drop table tempSSN";
    try {
        stmt1.executeUpdate(sqlString);
    } catch (SQLException e) {
    }

```

The `findSupervisees` method is divided into four stages.

**Stage 1:** The program prompts the user for input data (last name of employee) and queries the database for the social security number of employee. If there are several employees with same last name, the program lists all their social security numbers and prompts the user to choose one. At the end of this stage, the program would have the social security number of the employee whose supervisees the program needs to list. The code for this stage is shown below.

```

private static void findSupervisees(Connection conn)
    throws SQLException, IOException {

    String sqlString = null;

    Statement stmt = conn.createStatement();

    // Delete tuples from tempSSN from previous request.
    sqlString = "delete from tempSSN";
    try {
        stmt.executeUpdate(sqlString);
    } catch (SQLException e) {
        System.out.println("Could not execute Delete");
        stmt.close();
        return;
    }

    /* Get the ssn for the employee */
    sqlString = "select lname, fname, ssn " +
                "from   employee " +
                "where  lname = '";

    String lname =
        readEntry( "Enter last name of employee : ").trim();

```

```

sqlString += lname;
sqlString += "'";
ResultSet rset1;
try {
    rset1 = stmt.executeQuery(sqlString);
} catch (SQLException e) {
    System.out.println("Could not execute Query");
    stmt.close();
    return;
}
String samelName[] = new String[40];
String fName[] = new String[40];
String empssn[] = new String[40];
String ssn;
int nNames = 0;
while (rset1.next()) {
    samelName[nNames] = rset1.getString(1);
    fName[nNames] = rset1.getString(2);
    empssn[nNames] = rset1.getString(3);
    nNames++;
}
if (nNames == 0) {
    System.out.println("Name does not exist in database.");
    stmt.close();
    return;
}
else if (nNames > 1) {
    for(int i = 0; i < nNames; i++) {
        System.out.println(samelName[i] + "," +
                           fName[i] + " " + empssn[i]);
    }
    ssn = readEntry("Select ssn from list : ");
    ResultSet r = stmt.executeQuery(
        "select ssn from employee where ssn = '" + ssn + "'");
    if( !r.next()) {
        System.out.println("SSN does not exist in database.");
        stmt.close();
        return;
    }
}
else {
    ssn = empssn[0];
}

```

**Stage 2:** In the second stage, the `findSupervisees` method finds the immediate supervisees, i.e. supervisees who directly report to the employee. The social security numbers of immediate supervisees are then inserted into the temporary table.

```

/* Find immediate supervisees for that employee */
sqlString =
    "select distinct ssn from employee where superssn = '";
sqlString += ssn;
sqlString += "'";
try {
    rset1 = stmt.executeQuery(sqlString);
} catch (SQLException e) {
    System.out.println("Could not execute query");
    stmt.close();
    return;
}

/* Insert result into tempSSN table*/
Statement stmt1 = conn.createStatement();
while (rset1.next()) {
    String sqlString2 = "insert into tempSSN values ('";
    sqlString2 += rset1.getString(1);
    sqlString2 += "')";
    try {
        stmt1.executeUpdate(sqlString2);
    } catch (SQLException e) {
    }
}

```

**Stage 3:** In the third stage, the `findSupervisees` method iteratively calculates supervisees at the next lower level using the query:

```

select employee.ssn
from   employee, tempSSN
where  superssn = tempSSN.ssn;

```

The results of this query are inserted back into the `tempSSN` table to prepare for the next iteration. A boolean variable, called `newrowsadded`, is used to note if any new tuples were added to the `tempSSN` table in any particular iteration. Since the `tempSSN` table's only column (`ssn`) is also defined as a primary key, duplicate social security numbers would cause an `SQLException` which is simply ignored in this program. The code for this stage is shown below:

```

/* Recursive Querying */
ResultSet rset2;
boolean newrowsadded;
sqlString = "select employee.ssn from employee, tempSSN " +
            "where superssn = tempSSN.ssn";

```

```

do {
    newrowsadded = false;
    try {
        rset2 = stmt.executeQuery(sqlString);
    } catch (SQLException e) {
        System.out.println("Could not execute Query");
        stmt.close();
        stmt1.close();
        return;
    }
    while ( rset2.next()) {
        try {
            String sqlString2 = "insert into tempSSN values ('";
            sqlString2 += rset2.getString(1);
            sqlString2 += "')";
            stmt1.executeUpdate(sqlString2);
            newrowsadded = true;
        } catch (SQLException e) {
        }
    }
} while (newrowsadded);
stmt1.close();

```

**Stage 4:** In the final stage, the findSupervisees method prints names and social security numbers of all employees whose social security number is present in the tempSSN table. The code is shown below:

```

/* Print Results */
sqlString = "select fname, lname, e.ssn from " +
            "employee e, tempSSN t where e.ssn = t.ssn";
ResultSet rset3;
try {
    rset3 = stmt.executeQuery(sqlString);
} catch (SQLException e) {
    System.out.println("Could not execute Query");
    stmt.close();
    return;
}
System.out.println("      SUPERVISEES ");
System.out.print("FNAME");
for (int i = 0; i < 10; i++)
    System.out.print(" ");
System.out.print("LNAME");
for (int i = 0; i < 10; i++)
    System.out.print(" ");
System.out.print("SSN");

```

```

for (int i = 0; i < 6; i++)
    System.out.print(" ");
System.out.println("\n-----\n");

while(rset3.next()) {
    System.out.print(rset3.getString(1));
    for (int i = 0;
        i < (15 - rset3.getString(1).length()); i++)
        System.out.print(" ");
    System.out.print(rset3.getString(2));
    for (int i = 0;
        i < (15 - rset3.getString(2).length()); i++)
        System.out.print(" ");
    System.out.println(rset3.getString(3));
}
stmt.close();
}

```

This concludes the discussion of the `findSupervisees` method. The code to implement options (b) and (c) are quite straightforward and are omitted. The entire code for all the examples is available along with this lab manual.

## Exercises

1. Consider the SQL schema generated for the `UNIVERSITY` database in Laboratory Exercise 7.10 of the Elmasri/Navathe text.
  - a. Create the database tables in Oracle using the `SQL*Plus` interface.
  - b. Create comma separated data files containing data for at least 3 departments (CS, Math, and Biology), 20 students, and 20 courses. You may evenly distribute the students and the courses among the three departments. You may also choose to assign minors to some of the students. For a subset of the courses, create sections for the Fall 2006 and Spring 2007 terms. Make sure that you assign multiple sections for some courses. Assuming that the grades are available for Fall 2006 term, add enrollments of students in sections and assign numeric grades for these enrollments. Do not add any enrollments for the Spring 2007 sections.
  - c. Using the `SQL*Loader` utility of Oracle, load the data created in (b) into the database.
  - d. Write SQL queries for the following and execute them within an `SQL*Plus` session.
    - i. Retrieve student number, first and last names, and major department names of students who have a minor in the Biology department.
    - ii. Retrieve the student number and the first and last names of students who have never taken a class with instructor named "King".



- iii. Retrieve the student number and the first and last names of students who have taken classes only with instructor named "King".
- iv. Retrieve department names of departments along with the numbers of students with that department as their major (sorted in decreasing order of the number of students).
- v. Retrieve department names of departments along with the numbers of students with that department as their major (sorted in decreasing order of the number of students).
- vi. Retrieve the instructor names of all instructors teaching CS courses along with the sections (course number, section, semester, and year) they are teaching and the total number of students in these sections.
- vii. Retrieve the student number, first and last names, and major departments of students who do not have a grade of "A" in any of the courses they have enrolled in.
- viii. Retrieve the student number, first and last names, and major departments of all straight-A students (students who have a grade of "A" in all of the courses they have enrolled in).
- ix. For each student in the CS department, retrieve their student number, first and last names, and their GPA.

2. Consider the SQL schema generated for the MAIL\_ORDER database in Laboratory Exercise 7.11 of the Elmasri/Navathe text.

- a. Create the database tables in Oracle using the SQL\*Plus interface.
- b. Create comma separated data files containing data for at least 6 customers, 6 employees, and 20 parts. Also create data for 20 orders with an average of 3 to 4 parts in each order.
- c. Using the SQL\*Loader utility of Oracle, load the data created in (b) into the database.
- d. Write SQL queries for the following and execute them within an SQL\*Plus session.
  - i. Retrieve the names of customers who have ordered at least one part costing more than \$30.00.
  - ii. Retrieve the names of customers who have ordered all parts that cost less than \$20.00.
  - iii. Retrieve the names of customers who order parts only from employees who live in the same city as themselves.
  - iv. Retrieve a list of part number, part name, and total quantity ordered for that part. Produce a listing sorted in decreasing order of the total quantity ordered.
  - v. Retrieve the average waiting time (number of days) for all orders. The waiting time is defined as the difference between shipping date and order received date rounded up to the nearest day.

- vi. For each employee, retrieve their number, name, and total sales (in terms of dollars) in a given year, say 2005.
3. Consider the SQL schema generated for the `GRADE_BOOK` database in Laboratory Exercise 7.3 of the Elmasri/Navathe text.
- a. Create the database tables in Oracle using the `SQL*Plus` interface.
  - b. Create comma separated data files containing data for at least 20 students and 3 courses with an average of 8 students in each course. Create data for 3 to 4 grading components for each course. Then, create data to assign points to each student in every course they are enrolled in for each of the grading components.
  - c. Using the `SQL*Loader` utility of Oracle, load the data created in (b) into the database.
  - d. Write SQL queries for the following and execute them within an `SQL*Plus` session.
    - i. Retrieve the names of students who have enrolled in 'CSc 226' or 'CSc 227'.
    - ii. Retrieve the names of students who enrolled in some class during fall 2005 but no class in spring 2006.
    - iii. Retrieve the titles of courses that have had enrollments of five or fewer students.
    - iv. Retrieve the student ids, their names, and their course average (weighted average of points from all components) for the CSc 226 course offering with section number 2 during Fall 2005.

NOTE: All references to Exercises and Laboratory Exercises in the following problems refer to the numbering in the Elmasri/Navathe text.

4. Consider the `UNIVERSITY` database of Exercise 3.16 and Laboratory Exercise 3.31 for which the relational database was created and populated in Laboratory Exercise 8.26 in Oracle. Write and test a Java program that performs the functions illustrated in the following terminal session:

```
$ java p1
Student Number: 1234
Semester: Fall
Year: 2005

Main Menu
      (1) Add a class
      (2) Drop a class
      (3) See my schedule
      (4) Exit
Enter your choice: 1

Course Number: CSC 1010
Section: 2
```

Class Added

Main Menu

- (1) Add a class
- (2) Drop a class
- (3) See my schedule
- (4) Exit

Enter your choice: 1

Course Number: MATH 2010

Section: 1

Class Added

Main Menu

- (1) Add a class
- (2) Drop a class
- (3) See my schedule
- (4) Exit

Enter your choice: 3

Your current schedule is:

CSC 1010 Section 2, Introduction to Computers, Instructor: Smith  
MATH 2010 Section 1, Calculus I, Instructor: Jones

Main Menu

- (1) Add a class
- (2) Drop a class
- (3) See my schedule
- (4) Exit

Enter your choice: 2

Course Number: CSC 1010

Section: 2

Class dropped

Main Menu

- (1) Add a class
- (2) Drop a class
- (3) See my schedule
- (4) Exit

Enter your choice: 3

Your current schedule is:

MATH 2010 Section 1, Calculus I, Instructor: Jones

Main Menu

1. Add a class
2. Drop a class
3. See my schedule
4. Exit

Enter your choice: 4

\$

As the terminal session shows, the student signs into this program with an id number and the term and year of registration. The Add a class option allows the student to add a class to his/her current schedule and the Drop a class option allows the student to drop a class in his/her current schedule. The See my schedule option lists the classes in which the student is registered.

5. Consider the MAIL\_ORDER database of Laboratory Exercise 3.32 for which the relational database was created and populated in Laboratory Exercise 8.27 in Oracle. Write and test a Java program that prints an invoice for a given order number as illustrated in the following terminal session:

```
$java p2
```

```
Order Number: 1020
```

```
Customer:      Charles Smith
```

```
Customer No:   1111
```

```
Zip:           67226
```

```
Order No:      1020
```

```
Taken By:      John Jones (emp. No. 1122)
```

```
Received on:   10-DEC-1994
```

```
Shipped on:    13-DEC-1994
```

Part No.	Part Name	Quantity	Price	Sum
10506	Nut	200	1.99	398.00
10509	Bolt	100	1.55	155.00
Total:				553.00

6. Consider the GRADE\_BOOK database of Laboratory Exercise 4.28 for which the relational database was created and populated in Laboratory Exercise 8.29 in Oracle. Write and test a Java program that interacts with the Oracle database and prints a grade count report for each course. The report should list the number of A's, B's, C's, D's, and F's given in each of the courses sorted on course number. The format for the report is:

Grade Count Report

CNO	Course Title	#A's	#B's	#C's	#D's	#F's
CSc2010	Intro to CS	12	8	15	5	2
CSc2310	Java	15	2	12	8	1

## CHAPTER 4

### Relational Database Management Systems: MySQL

This chapter introduces the student to the MySQL database management system and PHP, the programming language used to program applications that access a MySQL database. The discussion in this chapter is not specific to any version of MySQL and all examples would work with MySQL 4.0 or higher version.

The COMPANY database of the Elmasri/Navathe text is used throughout this chapter. In Section 4.1, a larger data set is introduced for the COMPANY database. In Section 4.2, the `mysql` utility is introduced which allows the users to interact with the database including running commands, executing SQL statements, and running MySQL scripts. Finally, in Section 4.3 PHP programming with MySQL is introduced through a complete Web browsing application for the COMPANY database.

#### 4.1 COMPANY Database

Consider the COMPANY database state shown in Figure 5.6. Let us assume that the company is expanding with 3 new departments: Software, Hardware, and Sales. The Software department has 2 locations: Atlanta and Sacramento, the Hardware department is located in Milwaukee, and the Sales department has 5 locations: Chicago, Dallas, Philadelphia, Seattle, and Miami.

The Software department has 3 new projects: `OperatingSystems`, `DatabaseSystems`, and `Middleware` and the Hardware department has 2 new projects: `InkjetPrinters` and `LaserPrinters`. The company has added 32 new employees in this expansion process.

The updated COMPANY database is shown in the following tables.

DEPARTMENT	DNAME	DNUMBER	MGRSSN	MGRSTARTDATE
	Research	5	333445555	22-May-78
	Administration	4	987654321	1-Jan-85
	Headquarters	1	888665555	19-Jun-71
	Software	6	111111100	15-May-99
	Hardware	7	444444400	15-May-98
	Sales	8	555555500	1-Jan-97

PROJECT PNAME	PNUMBER PLOCATION	DNUM
ProductX	1 Bellaire	5
ProductY	2 Sugarland	5
ProductZ	3 Houston	5
Computerization	10 Stafford	4
Reorganization	20 Houston	1
Newbenefits	30 Stafford	4
OperatingSystems	61 Jacksonville	6
DatabaseSystems	62 Birmingham	6
Middleware	63 Jackson	6
InkjetPrinters	91 Phoenix	7
LaserPrinters	92 LasVegas	7

DEPT_LOCATIONS	DNUMBER	DLOCATION
	1	Houston
	4	Stafford
	5	Bellaire
	5	Sugarland
	5	Houston
	6	Atlanta
	6	Sacramento
	7	Milwaukee
	8	Chicago
	8	Dallas
	8	Philadephia
	8	Seattle

DEPENDENT	ESSN	DEPENDENT_NAME	SEX	BDATE	RELATION
	333445555	Alice	F	5-Apr-76	Daughter
	333445555	Theodore	M	25-Oct-73	Son
	333445555	Joy	F	3-May-48	Spouse
	987654321	Abner	M	29-Feb-32	Spouse
	123456789	Michael	M	1-Jan-78	Son
	123456789	Alice	F	31-Dec-78	Daughter
	123456789	Elizabeth	F	5-May-57	Spouse
	444444400	Johnny	M	4-Apr-97	Son
	444444400	Tommy	M	7-Jun-99	Son
	444444401	Chris	M	19-Apr-69	Spouse
	444444402	Sam	M	14-Feb-64	Spouse

WORKS_ON	ESSN	PNO	HOURS
	123456789	1	32.5
	123456789	2	7.5
	666884444	3	40
	453453453	1	20
	453453453	2	20
	333445555	2	10
	333445555	3	10
	333445555	10	10
	333445555	20	10
	999887777	30	30
	999887777	10	10
	987987987	10	35
	987987987	30	5
	987654321	30	20
	987654321	20	15
	888665555	20 null	
	111111100	61	40
	111111101	61	40
	111111102	61	40
	111111103	61	40
	222222200	62	40
	222222201	62	48
	222222202	62	40
	222222203	62	40
	222222204	62	40
	222222205	62	40
	333333300	63	40
	333333301	63	46
	444444400	91	40
	444444401	91	40
	444444402	91	40
	444444403	91	40
	555555500	92	40
	555555501	92	44
	666666601	91	40
	666666603	91	40
	666666604	91	40
	666666605	92	40
	666666606	91	40
	666666607	61	40
	666666608	62	40
	666666609	63	40
	666666610	61	40
	666666611	61	40
	666666612	61	40
	666666613	61	30
	666666613	62	10
	666666613	63	10



## EMPLOYEE

FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
James	E	Borg	888665555	10-Nov-27	450 Stone, Houston, TX	M	55000	null	1
Franklin	T	Wong	333445555	8-Dec-45	638 Voss, Houston, TX	M	40000	888665555	5
Jennifer	S	Wallace	987654321	20-Jun-31	291 Berry, Bellaire, TX	F	43000	888665555	4
John	B	Smith	123456789	9-Jan-55	731 Fondren, Houston, TX	M	30000	333445555	5
Alicia	J	Zelaya	999887777	19-Jul-58	3321 Castle, Spring, TX	F	25000	987654321	4
Ramesh	K	Narayan	666884444	15-Sep-52	971 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	31-Jul-62	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	29-Mar-59	980 Dallas, Houston, TX	M	25000	987654321	4
Jared	D	James	111111100	10-Oct-66	123 Peachtree, Atlanta, GA	M	85000	null	6
Alex	D	Freed	444444400	9-Oct-50	4333 Pillsbury, Milwaukee, WI	M	89000	null	7
John	C	James	555555500	30-Jun-75	7676 Bloomington, Sacramento, CA	M	81000	null	6
Jon	C	Jones	111111101	14-Nov-67	111 Allgood, Atlanta, GA	M	45000	111111100	6
Justin	null	Mark	111111102	12-Jan-66	2342 May, Atlanta, GA	M	40000	111111100	6
Brad	C	Knight	111111103	13-Feb-68	176 Main St., Atlanta, GA	M	44000	111111100	6
Evan	E	Wallis	222222200	16-Jan-58	134 Pelham, Milwaukee, WI	M	92000	null	7
Josh	U	Zell	222222201	22-May-54	266 McGrady, Milwaukee, WI	M	56000	222222200	7
Andy	C	Vile	222222202	21-Jun-44	1967 Jordan, Milwaukee, WI	M	53000	222222200	7
Tom	G	Brand	222222203	16-Dec-66	112 Third St, Milwaukee, WI	M	62500	222222200	7
Jenny	F	Vos	222222204	11-Nov-67	263 Mayberry, Milwaukee, WI	F	61000	222222201	7
Chris	A	Carter	222222205	21-Mar-60	565 Jordan, Milwaukee, WI	F	43000	222222201	7
Kim	C	Grace	333333300	23-Oct-70	6677 Mills Ave, Sacramento, CA	F	79000	null	6
Jeff	H	Chase	333333301	7-Jan-70	145 Bradbury, Sacramento, CA	M	44000	333333300	6
Bonnie	S	Bays	444444401	19-Jun-56	111 Hollow, Milwaukee, WI	F	70000	444444400	7
Alec	C	Best	444444402	18-Jun-66	233 Solid, Milwaukee, WI	M	60000	444444400	7
Sam	S	Snedden	444444403	31-Jul-77	987 Windy St, Milwaukee, WI	M	48000	444444400	7
Nandita	K	Ball	555555501	16-Apr-69	222 Howard, Sacramento, CA	M	62000	555555500	6
Bob	B	Bender	666666600	17-Apr-68	8794 Garfield, Chicago, IL	M	96000	null	8
Jill	J	Jarvis	666666601	14-Jan-66	6234 Lincoln, Chicago, IL	F	36000	666666600	8
Kate	W	King	666666602	16-Apr-66	1976 Boone Trace, Chicago, IL	F	44000	666666600	8
Lyle	G	Leslie	666666603	9-Jun-63	417 Hancock Ave, Chicago, IL	M	41000	666666601	8
Billie	J	King	666666604	1-Jan-60	556 Washington, Chicago, IL	F	38000	666666603	8
Jon	A	Kramer	666666605	22-Aug-64	1988 Windy Creek, Seattle, WA	M	41500	666666603	8
Ray	H	King	666666606	16-Aug-49	213 Delk Road, Seattle, WA	M	44500	666666604	8
Gerald	D	Small	666666607	15-May-62	122 Ball Street, Dallas, TX	M	29000	666666602	8
Arnold	A	Head	666666608	19-May-67	233 Spring St, Dallas, TX	M	33000	666666602	8
Helga	C	Pataki	666666609	11-Mar-69	101 Holyoke St, Dallas, TX	F	32000	666666602	8
Naveen	B	Drew	666666610	23-May-70	198 Elm St, Philadelphia, PA	M	34000	666666607	8
Carl	E	Reedy	666666611	21-Jun-77	213 Ball St, Philadelphia, PA	M	32000	666666610	8
Sammy	G	Hall	666666612	11-Jan-70	433 Main Street, Miami, FL	M	37000	666666611	8

## 4.2 mysql Utility

MySQL database system provides an interactive utility, called `mysql`, which allows the user to enter SQL commands interactively. One can also include one or more SQL statements in a file and have them executed within `mysql` using the `source` command.

In the following discussion, we will assume that your MySQL administrator has created a database called `company`<sup>4</sup>, created a MySQL user<sup>5</sup> called `book`, and has granted all rights to the `company` database to the `book` user.

The following `mysql` session creates the `department` table:

```
$ mysql -u book -p
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 1485 to server version: 4.1.9-standard

Type 'help;' or '\h' for help. Type '\c' to clear the buffer.

mysql> use company
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> source create-department.sql;
Query OK, 0 rows affected (0.04 sec)

mysql> show tables;
+-----+
| Tables_in_company |
+-----+
| department        |
| foo                |
+-----+
2 rows in set (0.00 sec)

mysql> exit;
Bye
$
```

In the above `mysql` session, the user invokes the `mysql` program with the `userid` `book` and the `-p` option to specify the password in a separate line. After connecting to the database server, the

---

<sup>4</sup> The administrator command to create a MySQL user is: `create user book identified by 'book';`  
Here the user id is `book` and the password is also `book`.

<sup>5</sup> The administrator command to create a database called `company` and to assign all rights to the `book` user is:  
`grant all on company.* to 'book'@'hostname.domain.edu';`

use company command is executed to start using the company database. Then, the source command is executed on the file create-department.sql that contains the following SQL statement:

```
CREATE TABLE department (
    dname          varchar(25) not null,
    dnumber        integer(4),
    mgrssn         char(9) not null,
    mgrstartdate   date,
    primary key (dnumber),
    key (dname)
);
```

**Note:** The data types supported in different DBMSs may have slightly different names; Please consult MySQL documentation to learn about all data types supported.

### Bulk Loading of Data

MySQL provides a “load” command that can load data stored in a text file into a table. The following mysql session illustrates the loading of the data located in department.csv file into the department table. In the following example, the load command is directly entered at the mysql prompt. It may be a good idea to create a script file with the load command and use the source command to execute the load command. This way, if there are any syntax errors, these can be corrected in the script file and the load command can be re-executed.

```
$ mysql -u raj -p
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 1493 to server version: 4.1.9-standard

Type 'help;' or '\h' for help. Type '\c' to clear the buffer.

mysql> use company;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> LOAD DATA LOCAL INFILE "department.csv"
INTO TABLE department FIELDS TERMINATED BY ","
OPTIONALLY ENCLOSED BY '"';
Query OK, 6 rows affected (0.00 sec)
Records: 6  Deleted: 0  Skipped: 0  Warnings: 0
```

```
mysql> select * from department;
+-----+-----+-----+-----+
|  dname          |  dnumber | mgrssn   | mgrstartdate |
+-----+-----+-----+-----+
| Research        |         5 | 333445555 | 1978-05-22   |
| Administration  |         4 | 987654321 | 1985-01-01   |
| Headquarters    |         1 | 888665555 | 1971-06-19   |
| Software        |         6 | 111111100 | 1999-05-15   |
| Hardware        |         7 | 444444400 | 1998-05-15   |
| Sales           |         8 | 555555500 | 1997-01-01   |
+-----+-----+-----+-----+
6 rows in set (0.00 sec)

mysql> exit;
Bye
$
```

The `LOAD DATA` command takes the name of the data file as parameter and other information such as field terminator symbols (in this case the comma) and loads the data into the table. For more details on the `load` command, please consult the MySQL documentation.

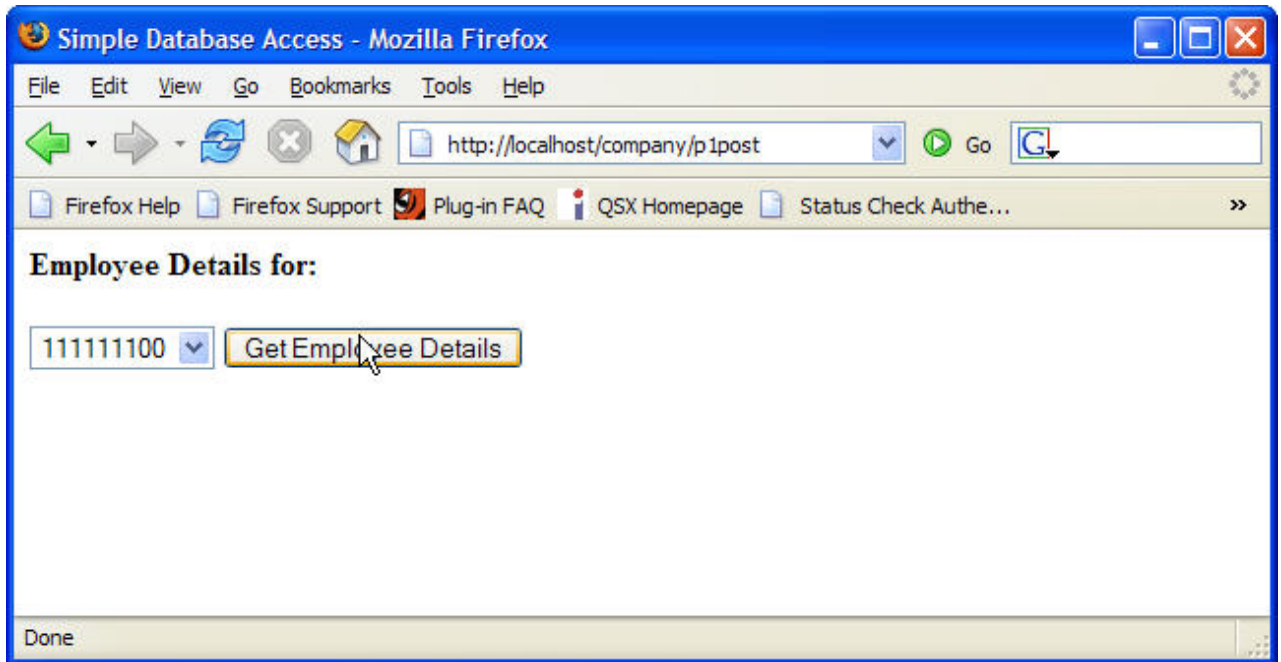
### 4.3 MySQL and PHP Programming

PHP is a very popular Web scripting language that allows the programmers to rapidly develop Web applications. In particular, PHP is most suited to develop Web applications that access a MySQL database. In this section, we illustrate the ease of programming with PHP and provide several examples of Web access to MySQL databases.

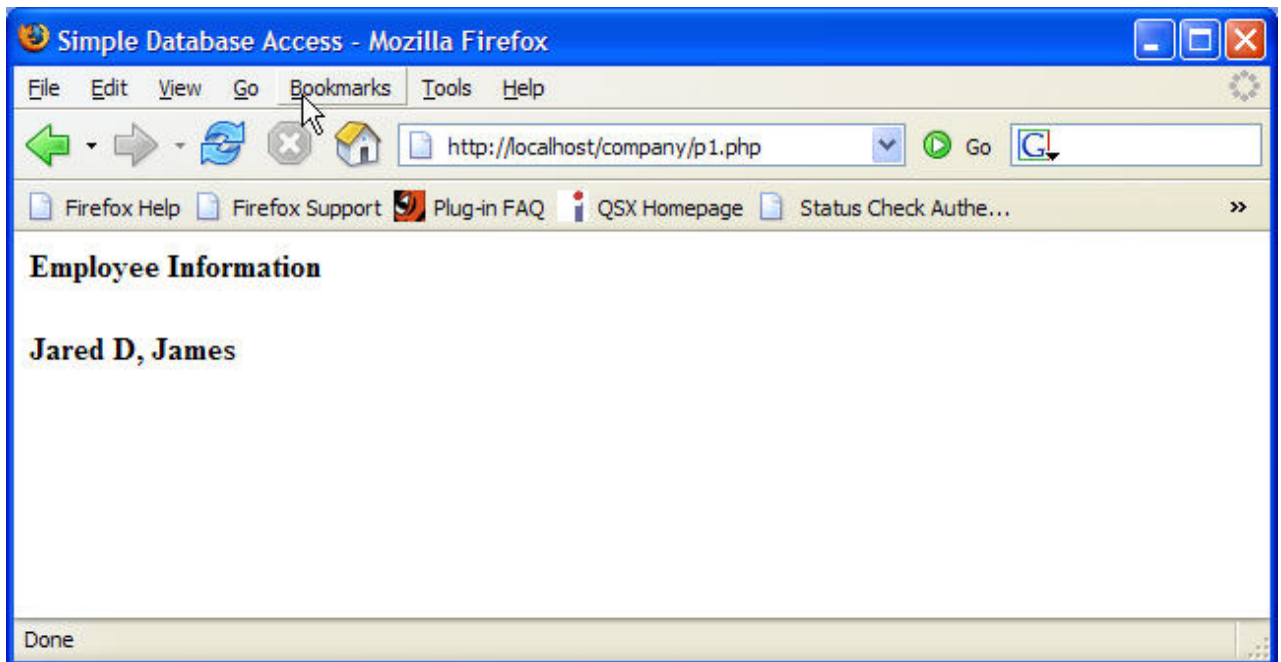
**Example 1:** Consider the problem of finding employee names given their social security number. To implement this problem as a Web application, we can design two Web pages:

1. The first Web page would contain a HTML form that contains a select list of social security numbers of employees and a submit button.
2. Upon choosing a social security number and submitting the form in the first Web page produces the second Web page that lists the name of the employee.

The two Web pages are shown in Figures 4.1 and 4.2.



**Figure 4.1: Initial Web Page - Example 1**



**Figure 4.2: Second Web Page - Example 1**

Both these Web pages contain dynamic information (obtained from the database) and therefore can easily be produced by PHP scripts. The PHP script (`p1post.php`) that produces the first Web page is shown below.

```

<html>
<head>
<title>Simple Database Access</title>
</head>
<body>

<?
$username="user";
$password="password";
$databse="company";
mysql_connect(localhost,$username,$password);
@mysql_select_db($databse) or die( "Unable to select database");
$query="SELECT ssn FROM employee";
$result=mysql_query($query);
$num=mysql_numrows($result);
mysql_close();
?>

<h4>Employee Details for:</h4>
<form method="post" action="p1.php">
<select name="ssn">

<?
$i=0;
while ($i < $num) {
    $ssn=mysql_result($result,$i,"ssn");
    echo "<option>", $ssn, "\n";
    $i++;
}
?>

</select>
<input type="submit" value="Get Employee Details">
</form>
</body>
</html>

```

A PHP script typically consists of HTML code to display “static” parts of the Web page interspersed with procedural PHP statements that produce “dynamic” parts of the Web page. The dynamic content may come from a database such as MySQL and hence most of the PHP procedural code involves connecting to the database, running queries and using the query result to produce parts of the Web page.

In the above example script, a simple HTML page is produced which has:

- Some static content such as text headers and a HTML form with a submit button. The HTML form when submitted invokes the second PHP script called “p1.php”.

- A dynamic “select” GUI element within the HTML form which contains a list of employee social security numbers for the users to choose from.

PHP statements are enclosed within `<? And >?`. HTML code can be produced within PHP code using the “echo” command as is seen in several places in the code. As can be seen, there are two blocks of PHP code in the example: one to connect to the database and execute a query and the second to use the results of the query to produce the HTML “select” list options.

The PHP script (p1.php) to produce the second Web page is shown next.

```
<html>
<head>
<title>Simple Database Access</title>
</head>
<body>
<h4>Employee Information</h4>

<?
$username="user";
$password="password";
$databse="company";
$ssn=$_POST['ssn'];
mysql_connect(localhost,$username,$password);
@mysql_select_db($databse) or die( "Unable to select database");
$query="SELECT * FROM employee where ssn=$ssn";
$result=mysql_query($query);
$num=mysql_numrows($result);
mysql_close();
if ($num == 1) {
    $fname=mysql_result($result,$i,"fname");
    $minit=mysql_result($result,$i,"minit");
    $lname=mysql_result($result,$i,"lname");
    echo "<b>$fname $minit, $lname</b>";
}?>

</body>
</html>
```

In this script, the employee social security number posted in the first Web page is retrieved using the PHP statement

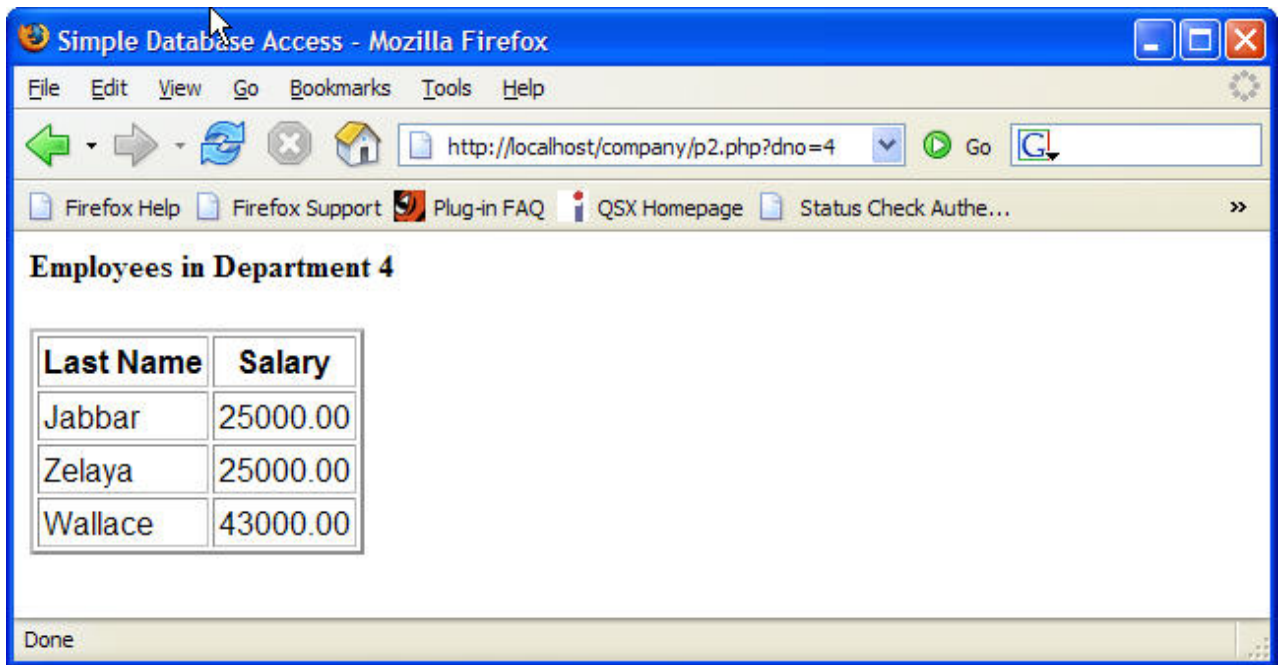
```
$ssn=$_POST['ssn'];
```

This social security number is then used in an SQL query to retrieve employee name. The script contains one block of PHP code surrounded by some HTML code.

**Example 2:** In this example, a PHP script that lists all employees in a given department is shown. The script takes as input the department number as a “GET” parameter in the URL itself as follows:

`http://localhost/company/p2.php?dno=4`

The Web page produced by the script is shown in Figure 4.3.



**Figure 4.3: Web Page for Example 2**

The PHP code (p2.php) is shown below:

```
<html>
<head>
<title>Simple Database Access</title>
</head>
<body>

<?
$username="user";
$password="password";
$databse="company";
$dno=$_GET['dno'];
mysql_connect(localhost,$username,$password);
@mysql_select_db($databse) or die( "Unable to select database");
$query="SELECT lname,salary FROM employee where dno=$dno";
$result=mysql_query($query);
```



```

$num=mysql_numrows($result);
mysql_close();
?>

<table border="2" cellspacing="2" cellpadding="2">
<tr>
<th><font face="Arial,Helvetica,sans-serif">Last Name</font></th>
<th><font face="Arial,Helvetica,sans-serif">Salary</font></th>
</tr>

<?
echo "<h4>Employees in Department $dno</h4>";
$i=0;
while ($i < $num) {
    $lname=mysql_result($result,$i,"lname");
    $salary=mysql_result($result,$i,"salary");
    ?>

<tr>
<td><font face="Arial, Helvetica, sans-serif">
<? echo $lname; ?>
</font></td>
<td><font face="Arial, Helvetica, sans-serif">
<? echo $salary; ?>
</font></td>
</tr>

<?
    $i++;
}
?>

</table>
</body>
</html>

```

The PHP script performs an SQL query to retrieve employee names and salaries who work for the given department. This information is then formatted neatly into an HTML table for display. This example illustrates more intricate embedding of PHP code within HTML code as is seen in the “while” loop towards the end of the script.

**Example 3:** A COMPANY database browser application is shown in this example. The initial Web page in this application lists all the departments in the company. By following hyperlinks, the user may see more details of departments, employees, and projects in three separate Web pages. The browser program is implemented using four PHP scripts:

- (a) `companyBrowse.php`: This script lists all the departments in the company in a tabular form as shown in Figure

- (b) deptView.php:
- (c) empView.php:
- (d) projectView.php:

The code for companyBrowse script is shown below:

```
<html>
<head>
<title>All Departments</title>
</head>
<body>

<?
$username="user";
$password="password";
$database="company";

mysql_connect("host.domain.edu",$username,$password);
@mysql_select_db($database) or die( "Unable to select database");
$query="SELECT dnumber,dname FROM department order by dnumber";
$result=mysql_query($query);
$num=mysql_numrows($result);
mysql_close();
?>

<h4>Departments of Company</h4>
<table border="2" cellspacing="2" cellpadding="2">
<tr>
<th><font face="Arial, Helvetica, sans-serif">
    Department Number</font></th>
<th><font face="Arial, Helvetica, sans-serif">
    Department Name</font></th>
</tr>

<?
$i=0;
while ($i < $num) {
    $dno=mysql_result($result,$i,"dnumber");
    $dname=mysql_result($result,$i,"dname");
    ?>

<tr>
<td><font face="Arial, Helvetica, sans-serif">
    <a href="deptView?dno=<? echo $dno; ?>">
    <? echo $dno; ?></a></font></td>
<td><font face="Arial, Helvetica, sans-serif">
    <? echo $dname; ?></font></td>
```

```

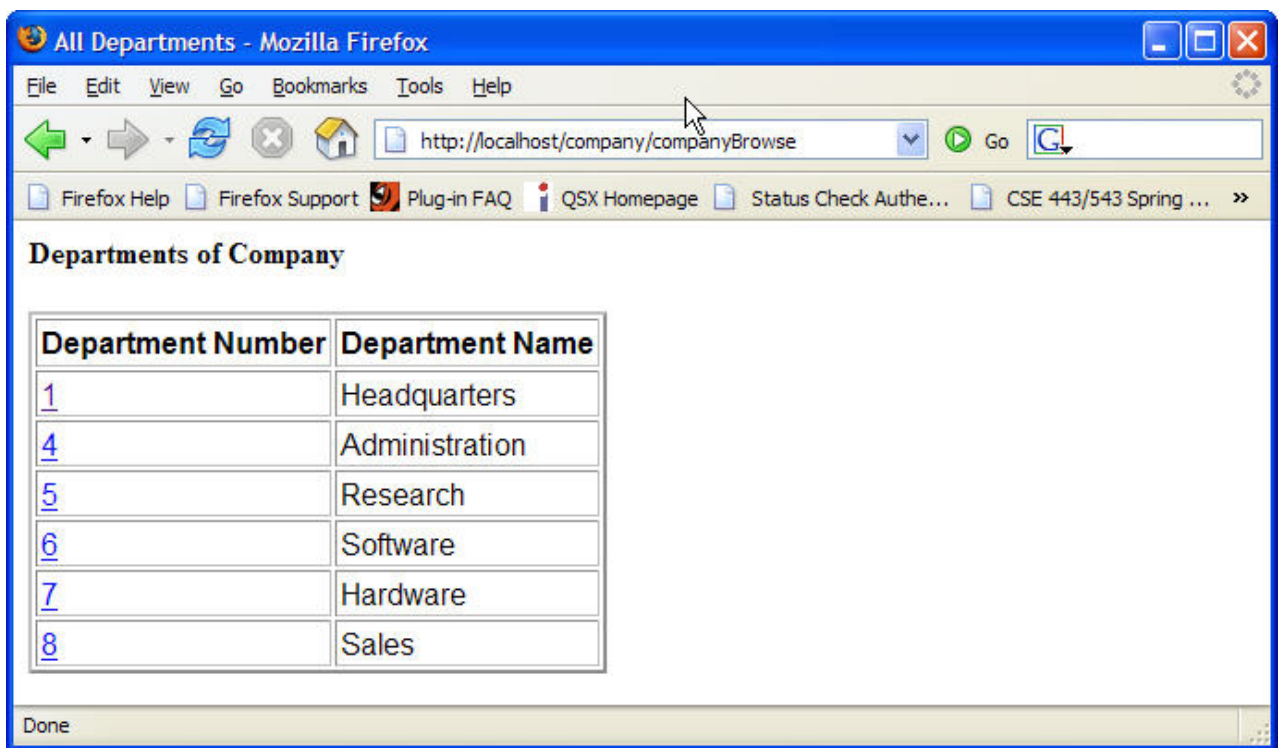
</tr>

<?
    $i++;
}
?>

</table>
</body>
</html>

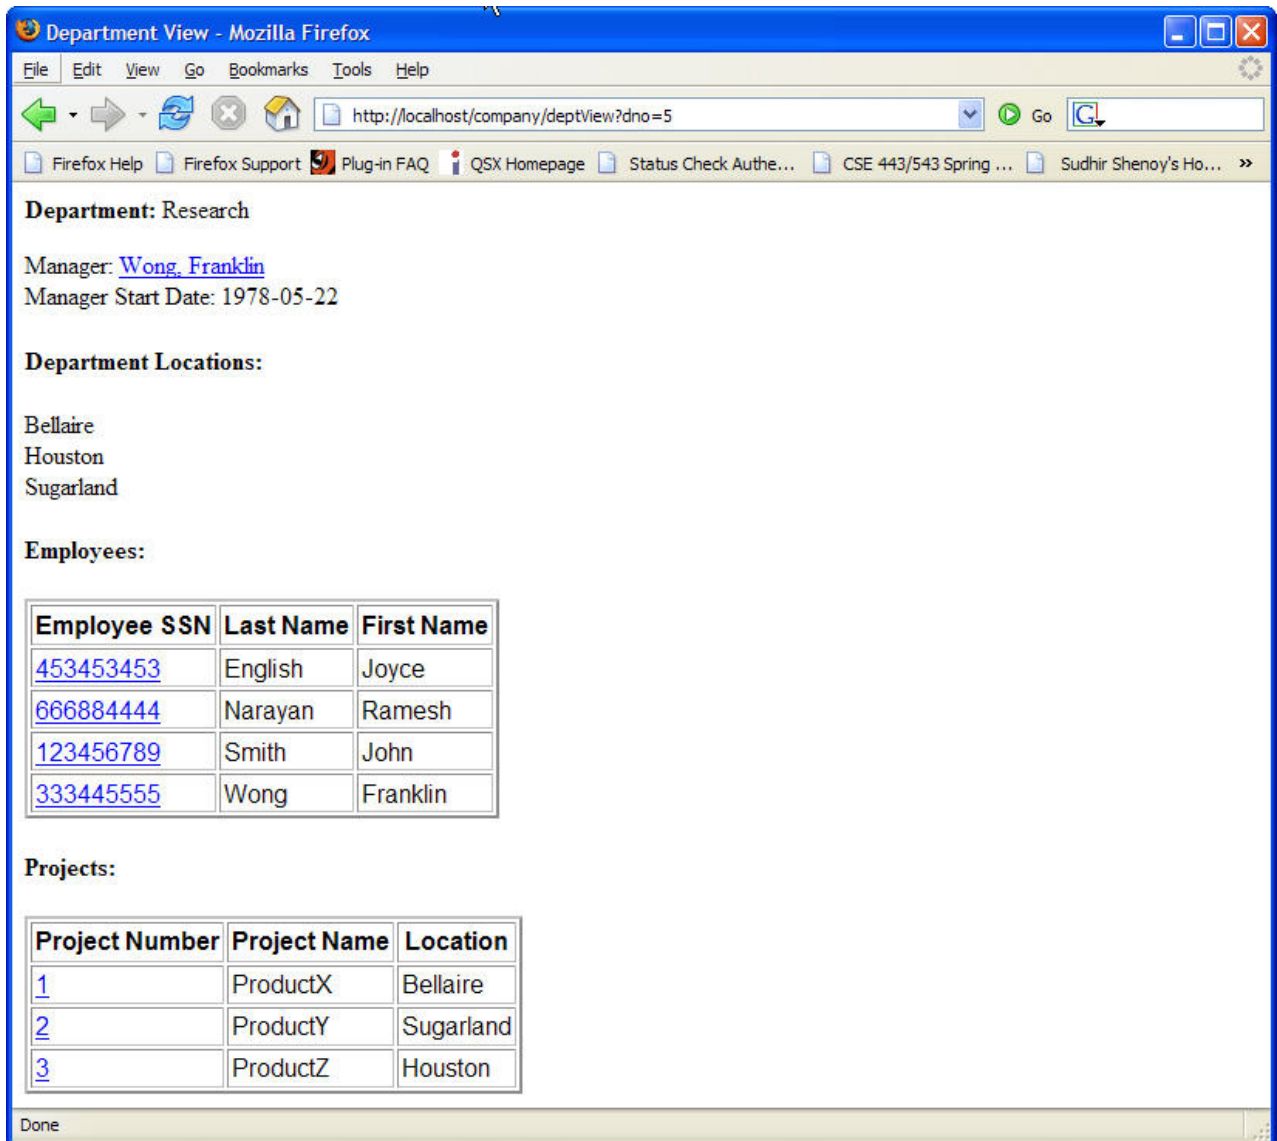
```

The script performs a simple query on the DEPARTMENT table and outputs the list of department numbers and names formatted as an HTML table as shown in Figure 4.4.



**Figure 4.4: All Departments Web Page – COMPANY database browser**

The department numbers in this list are formatted as HTML hyperlinks that, when traversed by the user, will produce a Web page containing more details of the chosen department. The detailed department view Web page is shown in Figure 4.5.



**Figure 4.5: Department Detail Web Page – COMPANY database browser**

The PHP script (deptView.php) that produces the detailed department view is shown below.

```
<html>
<head>
<title>Department View</title>
</head>
<body>

<?
$username="user";
$password="password";
$databse="company";
```

```

$dno=$_GET['dno'];
mysql_connect("host.domain.edu",$username,$password);
@mysql_select_db($database) or die( "Unable to select database");

$query="SELECT dname,mgrssn,mgrstartdate,lname,fname FROM
department,employee where dnumber=$dno and mgrssn=ssn";
$result=mysql_query($query);
$num=mysql_numrows($result);

$dname=mysql_result($result,0,"dname");
$mssn=mysql_result($result,0,"mgrssn");
$mstart=mysql_result($result,0,"mgrstartdate");
$mlname=mysql_result($result,0,"lname");
$mfname=mysql_result($result,0,"fname");

echo "<b>Department: </b>", $dname;
echo "<P>Manager: <a href=\"empView?\", $mssn, \">\", $mlname, ",
", $mfname, "</a></BR>";
echo "Manager Start Date: ", $mstart;

echo "<h4>Department Locations:</h4>";
$query="SELECT dlocation FROM dept_locations where dnumber=$dno";
$result=mysql_query($query);
$num=mysql_numrows($result);
$i=0;
while ($i < $num) {
    $dloc=mysql_result($result,$i,"dlocation");
    echo $dloc, "<BR>\n";
    $i++;
}

echo "<h4>Employees:</h4>";
$query="SELECT ssname,lname,fname FROM employee where dno=$dno";
$result=mysql_query($query);
$num=mysql_numrows($result);
?>

<table border="2" cellspacing="2" cellpadding="2">
<tr>
<th><font face="Arial, Helvetica, sans-serif">
Employee SSN</font></th>
<th><font face="Arial, Helvetica, sans-serif">
Last Name</font></th>
<th><font face="Arial, Helvetica, sans-serif">
First Name</font></th>
</tr>

```

```

<?
$i=0;
while ($i < $num) {
    $ssn=mysql_result($result,$i,"ssn");
    $lname=mysql_result($result,$i,"lname");
    $efname=mysql_result($result,$i,"fname");
?>
<tr>
    <td><font face="Arial, Helvetica, sans-serif">
        <a href="empView?ssn=<? echo $ssn; ?>">
            <? echo $ssn; ?></a></font></td>
    <td><font face="Arial, Helvetica, sans-serif">
        <? echo $lname; ?></font></td>
    <td><font face="Arial, Helvetica, sans-serif">
        <? echo $efname; ?></font></td>
</tr>
<?
    $i++;
}
?>

</table>

<?
echo "<h4>Projects:</h4>";
$query="SELECT pnumber,pname,plocation FROM project where
dnum=$dno";
$result=mysql_query($query);
$num=mysql_numrows($result);
?>

<table border="2" cellspacing="2" cellpadding="2">
<tr>
<th><font face="Arial, Helvetica, sans-serif">Project
Number</font></th>
<th><font face="Arial, Helvetica, sans-serif">Project
Name</font></th>
<th><font face="Arial, Helvetica, sans-serif">Location</font></th>
</tr>

<?
$i=0;
while ($i < $num) {
    $pnum=mysql_result($result,$i,"pnumber");
    $pname=mysql_result($result,$i,"pname");
    $ploc=mysql_result($result,$i,"plocation");

```

```

?>
<tr>
  <td><font face="Arial, Helvetica, sans-serif">
    <a href="projView?ssn=<? echo $pnum; ?>"><? echo $pnum;
?></a></font></td>
  <td><font face="Arial, Helvetica, sans-serif"><? echo $pname;
?></font></td>
  <td><font face="Arial, Helvetica, sans-serif"><? echo $ploc;
?></font></td>
</tr>
<?
  $i++;
}

mysql_close();
?>

</body>
</html>

```

The deptView script executes the following four queries and formats the results of the queries as shown in the Web page.

```

SELECT  dname,mgrssn,mgrstartdate,lname,fname
FROM    department,employee
WHERE   dnumber=$dno and mgrssn=ssn

```

```

SELECT  dlocation
FROM    dept_locations
WHERE   dnumber=$dno

```

```

SELECT  ssn,lname,fname
FROM    employee
WHERE   dno=$dno

```

```

SELECT  pnumber,pname,plocation
FROM    project
WHERE   dnum=$dno

```

Each of these queries uses the PHP variable \$dno containing the department number for which the view is generated.

The department view also contains hyperlinks for employees and projects which when traversed by the user produces detailed employee and project Web pages. The code for PHP scripts that produce these Web pages is omitted as they are similar to the deptView script.

## Exercises

NOTE: All references to Exercises and Laboratory Exercises in the following problems refer to the numbering in the Elmasri/Navathe text.

7. Consider the SQL schema generated for the `CONFERENCE_REVIEW` database in Laboratory Exercise 7.12.
  - a. Create the database tables in MySQL.
  - b. Create comma separated data files containing data for at least 10 papers, 15 authors, and 8 reviewers. You must assign authors to papers, assign a contact author, and assign reviewers to papers. It is possible for some authors to be reviewers. In such a case, the reviewer should not be assigned to a paper in which he or she is an author. You should also create data for the individual reviews for each paper. The values for various rankings must be between 1 and 10.
  - c. Using the MySQL loader command, load the data created in (b) into the database.
  - d. Write SQL queries for the following and execute them in a MySQL interactive session.
    - i. Retrieve the email and names of authors who have submitted more than one paper for review.
    - ii. Retrieve the email id of the contact author, the title of paper submitted by that author, and the average rating received on its reviews such that the average rating was above 8.0.
    - iii. Retrieve the paper id, title, and average ratings received for all papers sorted in decreasing order of average rating. Do not show any papers below an average rating of 3.0.
    - iv. Retrieve the email and names of reviewers along with the average ratings they gave to papers assigned to them. Sort the result in increasing order of average rating.
    - v. Retrieve the paper id and title of papers along with the author name and email id such that the author is also a reviewer.
8. Consider the SQL schema generated for the `GRADE_BOOK` database in Laboratory Exercise 7.13.
  - a. Create the database tables in MySQL.
  - b. Create comma separated data files containing data for at least 20 students and 3 courses with an average of 8 students in each course. Create data for 3 to 4 grading components for each course. Then, create data to assign points to each student in every course they are enrolled in for each of the grading components.
  - c. Using the MySQL loader command, load the data created in (b) into the database.
  - d. Write SQL queries for the following and execute them in a MySQL interactive session.
    - i. Retrieve the course numbers, titles, and term of courses in which student with id=1234 has enrolled.
    - ii. Retrieve the term, section number, and course numbers along with their total enrollments.



- iii. Retrieve the term, section number, and course numbers of courses in which there is a grading component that has a weight of more than 50 and has fewer than 10 students.
  - iv. Retrieve the names of students who have enrolled in all courses taught by the instructor in the Fall 2005 term.
- 9. Consider the SQL schema generated for the `ONLINE_AUCTION` database in Laboratory Exercise 7.14.
  - a. Create the database tables in MySQL.
  - b. Create comma separated data files containing data for at least 20 items with at least 10 buyers and at least 10 sellers. Observe the queries below and make sure that your data will be appropriate for those queries.
  - c. Using the MySQL loader command, load the data created in (b) into the database.
  - d. Write SQL queries for the following and execute them in a MySQL interactive session.
    - i. Retrieve item numbers and descriptions of items that are still active (i.e. bidding has not closed) along with their current bidding price.
    - ii. Retrieve the item numbers and descriptions of all items in which the member with member id `abc24` is the winner.
    - iii. Retrieve the member ids and names of members who have an average rating of 8.0 or above.
    - iv. Retrieve the bidding history for the item with item number `i246`. The bidding history should contain a chronological listing of all bids till the current time including the member id of the bidder and the bidding price.
    - v. Retrieve the item number and titles of items that fall under the `/COMPUTER/HARDWARE/PRINTER` category and have “HP” in their titles and on which bidding has not yet closed sorted by time of bid closing.
- 10. Consider the `CONFERENCE_REVIEW` database of Laboratory Exercise 3.34 for which the relational database was created and populated in Laboratory Exercise 8.28 in MySQL. There are three types of users of the system: *contact author*, *reviewer*, and *administrator*. The contact author should be able to sign in and submit a paper to the conference; the reviewer should be able to submit reviews of papers assigned to him or her; and the administrator should be able to assign reviewers to papers (at least three reviewers to each paper) as well as print a report of all papers sorted by average rating assigned by reviewers. The reviewers and administrator have pre-assigned user ids and passwords, however, contact authors must register themselves first before they can start using the system. Write a PHP-based application that implements the following functions:
  - a. A user registration page that allows a contact author to register their email and other information with the system. They choose a password in the registration page.
  - b. A user login page with GUI textbox elements to accept user id and password and radio buttons to choose between contact author, reviewer, and administrator. A separate menu for each user should be displayed after authentication.
  - c. The contact author should be able to enter details of his paper submission and upload a file.

- d. The reviewer should be able to choose one of his assigned papers and submit a review for the paper.
  - e. The administrator should be able to assign at least three reviewers to each paper. The administrator should also be able to display a report of all papers sorted by average rating (high to low).
11. Consider the `GRADE_BOOK` database of Laboratory Exercise 4.28 for which the relational database was created and populated in Laboratory Exercise 8.29 in MySQL. Write and test PHP scripts that access the MySQL database that implements the following two functions for a given course section (the course number, term, and section information will be provided as GET parameters in the URL itself):
- i. Allows the instructor to enter the points/scores for a particular grading component. The instructor should be able to choose a grading component using a pull down list.
  - ii. Allows the instructor to view the scores for all components for all students in a tabular form with one row per student and one column per grading component.
12. Consider the `ONLINE_AUCTION` database of Laboratory Exercise 4.29 for which the relational database was created and populated in Laboratory Exercise 8.30 in MySQL. Write PHP scripts to implement an online auction Web site. A member may register himself/herself and choose his/her password. After that, they may sign into the system. The following functions should be implemented:
- a. A buyer should be able to search items by entering a keyword. The result of the search should be a listing of items, each with a hyper-link. Upon clicking on the hyper-link, a detailed page describing the item should be displayed. On this page, the user may place a bid on the item.
  - b. A seller should be able to place an item for sale by providing all details of the item along with auction deadlines etc.
  - c. Buyers and sellers should be able to look at a list of all items on which bidding has ended.
  - d. Buyers and sellers should be able to place a rating as well as view ratings on transactions they are involved in.

## CHAPTER 5

### Database Design (DBD) Toolkit

This chapter introduces a toolkit written in SWI-Prolog (a freely available Prolog system for the Unix as well as the Windows operating systems) that allows the student to learn and test out the various concepts, definitions, and algorithms associated with relational database design using functional dependency theory.

The DBD system is available as SWI-Prolog source code in a single file, `dbd.pl`. This file needs to be consulted (read into the Prolog interpreter's memory) before invoking any of the DBD predicates.

#### 5.1 Coding Relational Schemas and Functional Dependencies

The DBD system uses lists to code the relational schemas and functional dependencies. Variables in Prolog begin with an upper-case letter and hence attributes of relational schemas are coded as Prolog constants – terms that begin with a lower-case letter. Consider the following relational schema and associated set of functional dependencies:

```
R(A,B,C,D,E)
F = { A → B, AC → D, E → AB, AD → E, ABC → D }
```

These will be coded as the following Prolog predicates in DBD:

```
schema([a,b,c,d,e]).
fds([ [[a],[b]], [[a,c],[d]], [[e],[a,b]],
      [[a,d],[e]], [[a,b,c],[d]] ]).
```

#### 5.2 Invoking the SWI-Prolog Interpreter

Once the SWI-Prolog system is installed in your server or local computer, the interpreter is invoked using the following terminal command:

```
$ p1
```

Here \$ is the command prompt. In the rest of the chapter, all user inputs are shown in **bold** font and system displays in regular font.

After printing a welcome message, the interpreter responds with the following prompt:

```
?-
```

At this prompt, the user may enter a Prolog query to invoke a predicate. In order to invoke one of the predicates of the DBD system, the user must first load Prolog interpreter's memory with the DBD program using the following command/query:

```
?- ['dbd.pl'].
```

Note that all commands/queries to Prolog end with a period (".") symbol. To exit the Prolog system, the halt command/query is entered at the Prolog prompt as follows:

```
?- halt.
```

To solve a particular problem using the DBD system, the student will usually create a file containing the code to represent the relational schema and the functional dependencies. In addition, one or more Prolog rules may be included to execute predicates to solve the problem at hand. For example, consider the problem of finding one or more candidate keys for the relational schema and functional dependencies mentioned earlier in the section. The student will create a text file with the following code:

```
schema([a,b,c,d,e]).
fds([ [a],[b]], [[a,c],[d]], [[e],[a,b]],
      [[a,d],[e]], [[a,b,c],[d]] ).
answer(K) :- schema(R), fds(F), candkey(R,F,K).
```

The first two predicates (schema and fds) code the relational schema and functional dependencies. The third rule invokes the candkey predicate of the DBD system to find the candidate keys for R and F. The answer is captured in the answer predicate with one argument. Let us assume that the above code is present in a file called example.pl. The following interaction with the SWI-Prolog system shows the steps necessary to find the answer(s) to the problem:

```
$ pl
Welcome to SWI-Prolog (Multi-threaded, Version 5.4.7)
Copyright (c) 1990-2003 University of Amsterdam.
SWI-Prolog comes with ABSOLUTELY NO WARRANTY. This is free software,
and you are welcome to redistribute it under certain conditions.
Please visit http://www.swi-prolog.org for details.
```

```
For help, use ?- help(Topic). or ?- apropos(Word).
```

```
?- ['dbd.pl'].
% dbd.pl compiled 0.02 sec, 30,740 bytes

Yes
?- ['example.pl'].
% example.pl compiled 0.00 sec, 1,116 bytes

Yes
?- answer(K).

K = [c, e] ;
```

```

K = [a, c] ;

No
?- answer(K).

K = [c, e]

Yes
?- halt.
$

```

The user inputs are shown in **bold** font. After invoking the interpreter, the user first loads the DBD system code (`dbd.pl`). The user then loads the code to solve the problem at hand (`example.pl`). Finally, the user submits the query, **`answer(K)`**, to determine the candidate keys. By default, every Prolog system responds with the first answer to the query and waits for user input. At this point the user has two choices: press the ENTER key to stop looking at further answers to the query or press the semi-colon key to request Prolog to display the next answer to the query. In the above screen capture, both options are shown.

### 5.3 DBD System Predicates

The DBD system predicates are discussed in this section. Almost all of the predicates take as input the relation schema  $R$  and a set of functional dependencies  $F$ .  $R$  and  $F$  must be coded properly as indicated in earlier sections. In particular, the attribute names must begin with a lower-case letter. The predicate names themselves also must begin with a lower-case letter.

#### 5.3.1 `xplus(R,F,X,Xplus)`

The `xplus` predicate takes as input a relation schema  $R$ , a set of functional dependencies  $F$ , and a subset  $X$  of  $R$  and produces as output the attribute-closure of  $X$  under  $F$ , i.e. the set of attributes that are functionally dependent on  $X$ . The following shows a sample interaction with the Prolog system to compute the attribute closure of  $\{A, C\}$  for the relational schema and functional dependencies mentioned earlier in this chapter (the file `e2.pl` is assumed to contain the `schema` and `fds` predicates):

```

?- ['dbd.pl'].
% dbd.pl compiled 0.02 sec, 30,740 bytes

Yes
?- ['e2.pl'].
% e2.pl compiled 0.00 sec, 1,096 bytes

Yes
?- schema(R),fds(F),xplus(R,F,[a,c],Xplus).

R = [a, b, c, d, e]
F = [[a], [b]], [[a, c], [d]], [[e], [a, b]], [[a, d], [e]], [[a, b, c], [d]]]
Xplus = [a, b, c, d, e]

```

```
Yes
?- halt.
$
```

### 5.3.2 finfplus(R,F,[X,Y])

The `finfplus` predicate takes as input a relation schema  $R$ , a set of functional dependencies  $F$ , and a functional dependency  $X \rightarrow Y$  (in list coded form  $[X, Y]$ ) and returns “yes” if  $X \rightarrow Y$  is in the closure of  $F$  and “no” otherwise. An example using the same  $R$  and  $F$  as before follows:

```
?- ['dbd.pl'].
% dbd.pl compiled 0.02 sec, 30,740 bytes

Yes
?- ['e2.pl'].
% e2.pl compiled 0.00 sec, 1,096 bytes

Yes
?- schema(R), fds(F), finfplus(R,F,[[a,b],[c]]).

No
?- schema(R), fds(F), finfplus(R,F,[[a,d],[b]]).

R = [a, b, c, d, e]
F = [[[a], [b]], [[a, c], [d]], [[e], [a, b]], [[a, d], [e]], [[a, b, c], [d]]]

Yes
```

In the above screen capture we observe that the functional dependency  $AB \rightarrow C$  is not in the closure of  $F$  whereas the functional dependency  $AD \rightarrow B$  is in the closure of  $F$ .

### 5.3.3 fplus(R,F,Fplus)

The `fplus` predicate takes as input a relational schema  $R$  and a set of functional dependencies  $F$  and returns as output the closure of  $F$  in the third parameter. Typically the number of functional dependencies in the closure of  $F$  is large and in the worst case can be exponential in the number of attributes of  $R$ .

Consider the following Prolog code in file `e3.pl`:

```
schema([a,b,c,d,e]).
fds([ [[a],[b]], [[a,c],[d]], [[e],[a,b]],
      [[a,d],[e]], [[a,b,c],[d]] ]).
displayFDs([]) :- nl.
displayFDs([X,Y|Rest]) :-
    write(X),write('-->'), write(Y),nl,displayFDs(Rest).
go :- schema(R), fds(F), fplus(R,F,Fplus), displayFDs(Fplus).
```

This program adds the display rules to the `e2.pl` file and invokes the `displayFDs` predicate on the closure of `F`. All functional dependencies that are implied by `F` are displayed by the following invocation:

```
?- ['dbd.pl'].
% dbd.pl compiled 0.02 sec, 30,740 bytes

Yes
?- ['e3.pl'].
% e3.pl compiled 0.01 sec, 1,700 bytes

Yes
?- go.
[a]-->[a, b]
[a]-->[b]
...
...
```

**Note:** The `fplus` predicate does not generate any trivial functional dependencies (dependencies whose RHS attributes are also present in the LHS attributes).

### 5.3.4 `implies(R,F1,F2)` and `equiv(R,F1,F2)`

The `implies` predicate takes as input a relational schema `R` and two sets of functional dependencies `F1` and `F2`. It returns “yes” if `F1` implies `F2`, i.e. if each functional dependency in `F2` is in the closure of `F1`; returns “no” otherwise.

The `equiv` predicate takes as input a relational schema `R` and two sets of functional dependencies `F1` and `F2`. It returns “yes” if `F1` is equivalent to `F2`, i.e. if each functional dependency in `F2` is in the closure of `F1` and if each functional dependency in `F1` is in the closure of `F2`; returns “no” otherwise.

Consider the relational schema `R(A, B, C)` and the two sets of functional dependencies:

```
F1 = {A → BC, B → A}
F2 = {A → B, A → C}
```

The above schema and functional dependencies can be coded into a file `e4.pl` as follows:

```
schema([a,b,c]).
fds1([[[a],[b,c]], [[b],[a]]]).
fds2([[[a],[b]], [[a],[c]]]).
```

The following screen capture illustrates the `implies` and `equiv` predicates:

```
?- ['dbd.pl'].
% dbd.pl compiled 0.02 sec, 30,740 bytes
```

```

Yes
?- ['e4.pl'].
% e4.pl compiled 0.00 sec, 1,120 bytes

Yes
?- schema(R), fds1(F1), fds2(F2), implies(R,F1,F2).

R = [a, b, c]
F1 = [[[a], [b, c]], [[b], [a]]]
F2 = [[[a], [b]], [[a], [c]]]

Yes
?- schema(R), fds1(F1), fds2(F2), implies(R,F2,F1).

No
?- schema(R), fds1(F1), fds2(F2), equiv(R,F1,F2).

No
?-

```

### 5.3.5 superkey(R,F,K) and candkey(R,F,K)

The superkey and candkey predicates both take a relation schema R and a set of functional dependencies as input in the first two parameters. The third parameter may be a constant (i.e. an instantiated list of attributes of R) or a variable.

In case the third parameter is instantiated to a list of attributes the superkey (candkey) predicate with return “yes” if the instantiated third parameter is a super key (candidate key) of R; “no” otherwise.

In case the third parameter is a variable, the predicate superkey (candkey) will generate all super keys (candidate keys) one at a time.

Consider the following e5.pl file containing a relation schema and associated functional dependencies:

```

schema([a,b,c]).
fds([[[a],[b,c]], [[b],[a]]]).

```

The following screen capture illustrates usage of the superkey and candkey predicates:

```

?- ['dbd.pl'].
% dbd.pl compiled 0.01 sec, 30,428 bytes

Yes
?- ['e5.pl'].
% e5.pl compiled 0.00 sec, 1,124 bytes

Yes
?- schema(R), fds(F), candkey(R,F,K).

```



```
R = [a, b, c]
F = [[[a], [b, c]], [[b], [a]]]
K = [b] ;
```

```
R = [a, b, c]
F = [[[a], [b, c]], [[b], [a]]]
K = [a] ;
```

```
No
?- schema(R), fds(F), superkey(R,F,K).
```

```
R = [a, b, c]
F = [[[a], [b, c]], [[b], [a]]]
K = [a, b, c] ;
```

```
R = [a, b, c]
F = [[[a], [b, c]], [[b], [a]]]
K = [b, c] ;
```

```
R = [a, b, c]
F = [[[a], [b, c]], [[b], [a]]]
K = [b] ;
```

```
R = [a, b, c]
F = [[[a], [b, c]], [[b], [a]]]
K = [a, c] ;
```

```
R = [a, b, c]
F = [[[a], [b, c]], [[b], [a]]]
K = [a] ;
```

```
R = [a, b, c]
F = [[[a], [b, c]], [[b], [a]]]
K = [a, b] ;
```

```
No
?- schema(R), fds(F), candkey(R,F,[a,b]).
```

```
No
?- schema(R), fds(F), superkey(R,F,[a,b]).
```

```
R = [a, b, c]
F = [[[a], [b, c]], [[b], [a]]]
```

```
Yes
?-
```

### 5.3.6 mincover(R,F,FC)

The `mincover` predicate takes as input a relation schema `R` and a set of functional dependencies `F` and returns as output in the third parameter the minimal cover for `F`.

Consider the following relation schema and functional dependencies in file `e6.pl`:

```
schema([a,b,c,d,e]).
fds([[[a],[b,c]], [[b],[a]], [[a,b],[d]], [[a],[d]], [[b,c],[a]]]).
```

The functional dependencies coded in the file are:

$$F = \{A \rightarrow BC, B \rightarrow A, AB \rightarrow D, A \rightarrow D, BC \rightarrow A\}$$

The following screen capture illustrates the mincover predicate usage:

```
?- ['dbd.pl'].
% dbd.pl compiled 0.02 sec, 30,740 bytes

Yes
?- ['e6.pl'].
% e6.pl compiled 0.00 sec, 1,072 bytes

Yes
?- schema(R),fds(F),mincover(R,F,MinF).

R = [a, b, c, d, e]
F = [[[a], [b, c]], [[b], [a]], [[a, b], [d]], [[a], [d]], [[b, c], [a]]]
MinF = [[[a], [b, c, d]], [[b], [a]]]

Yes
?-
```

The minimum cover for the set of functional dependencies is

$$\text{MinF} = \{A \rightarrow BCD, B \rightarrow A\}.$$

### 5.3.7 $\text{lj d}(R, F, R1, R2)$ , $\text{lj d}(R, F, D)$ , and $\text{fp d}(R, F, D)$

The predicates discussed in this sub-section are related to decompositions of relation schemas. We represent decompositions using lists of relational schemas. For example, the decomposition  $D = \{ABC, ABD, BDE\}$  of  $R = ABCDE$  is represented as the following Prolog term:

$$D = [[a,b,c], [a,b,d], [b,d,e]]$$

#### Testing for lossless join decomposition

There are two versions of the  $\text{lj d}$  predicate.

The first version of the  $\text{lj d}$  predicate takes as input a relation scheme  $R$ , a set of functional dependencies  $F$ , and two relation schemas  $R1$  and  $R2$  that are decompositions of  $R$ . The predicate returns “yes” if  $(R1, R2)$  is a lossless join decomposition of  $R$ ; “no” otherwise.

As an example, consider the  $R = ABC$  and  $F = \{A \rightarrow B\}$ . Consider the decomposition of  $R$  into two sub-schemas  $R_1 = AB$  and  $R_2 = AC$ . The following screen capture illustrates the use of the first version of `ljd` predicate:

```
?- ['dbd.pl'].
% dbd.pl compiled 0.02 sec, 30,740 bytes

Yes
?- ljd([a,b,c],[[a],[b]], [a,b],[a,c]).

Yes
?-
```

The second version of the `ljd` predicate takes as input a relation scheme  $R$ , a set of functional dependencies  $F$ , and a decomposition  $D$  into two or more sub-schemas.  $D$  is represented as a list of relation schemas where each schema is itself a list of attributes. The predicate returns “yes” if  $D$  is a lossless join decomposition of  $R$ ; “no” otherwise. This version implements the “matrix” method (Algorithm 11.1 of the Elmasri/Navathe text) to test for lossless join and prints the final matrix.

As an example, consider the following file (named `e7a.pl`) containing a relational schema, a set of functional dependencies and a decomposition into more than two sub-schemas:

```
schema([a,b,c,d,e]).
fds([[a],[c]], [[b],[c]], [[c],[d]], [[d,e],[c]], [[c,e],[a]]).
decomp([[a,d],[a,b],[b,e],[c,d,e],[a,e]]).
```

The following screen capture illustrates the usage of the more general version of the `ljd` predicate:

```
?- ['dbd.pl'].
% dbd.pl compiled 0.02 sec, 30,740 bytes

Yes
?- ['e7a.pl'].
% e7a.pl compiled 0.01 sec, 1,440 bytes

Yes
?- schema(R), fds(F), decomp(D), ljd(R,F,D).
[[a, 1], [b, 1, 2], [a, 3], [a, 4], [b, 1, 5]]
[[a, 1], [a, 2], [a, 3], [a, 4], [b, 2, 5]]
[[a, 1], [a, 2], [a, 3], [a, 4], [a, 5]]
[[a, 1], [b, 4, 2], [a, 3], [a, 4], [a, 5]]
[[a, 1], [b, 5, 2], [a, 3], [a, 4], [a, 5]]

R = [a, b, c, d, e]
F = [[[a],[c]], [[b],[c]], [[c],[d]], [[d,e],[c]], [[c,e],[a]]]
D = [[a,d],[a,b],[b,e],[c,d,e],[a,e]]

Yes
?-
```

Note that the `ljd` predicate execution prints the final matrix to the screen; In this case it can be observed that the third row is turned into an “all-a” row.

### Testing for FD-preserving decomposition

The `fpd` predicate takes as input a relation scheme  $R$ , a set of functional dependencies  $F$ , and a decomposition  $D$  into two or more sub-schemas.  $D$  is represented as a list of relation schemas where each schema is itself a list of attributes. The predicate returns “yes” if  $D$  is a functional dependency preserving decomposition of  $R$ ; “no” otherwise.

As an example, consider the following file (named `e7b.pl`) containing a relational schema, a set of functional dependencies and a decomposition into sub-schemas:

```
schema([a,b,c,d]).
fds([[a],[b]], [[b],[c]], [[c],[d]], [[d],[a]]).
decomp([[a,b],[b,c],[c,d]]).
```

The following screen capture illustrates the usage of the `fpd` predicate:

```
?- ['dbd.pl'].
% dbd.pl compiled 0.03 sec, 30,740 bytes

Yes
?- ['e7b.pl'].
% e7b.pl compiled 0.00 sec, 1,260 bytes

Yes
?- schema(R), fds(F), decomp(D), fpd(R,F,D).
Considering [a]-->[b]
Xplus=[a, b, c, d]
Considering [b]-->[c]
Xplus=[a, b, c, d]
Considering [c]-->[d]
Xplus=[a, b, c, d]
Considering [d]-->[a]
Xplus=[a, b, c, d]

R = [a, b, c, d]
F = [[a],[b]], [[b],[c]], [[c],[d]], [[d],[a]]
D = [[a,b],[b,c],[c,d]]

Yes
?-
```

Note that the `fpd` predicate execution prints each functional dependency in  $F$  and displays the attribute closure of the LHS of the functional dependency with respect to the “projected” functional dependencies. The algorithm implemented in this predicate is from “*Jeffrey D. Ullman, Principles of Database Systems, Second Edition, Computer Science Press, 1982*”.

### 5.3.8 **is3NF(R,F)** and **threenf(R,F,D)**

The **is3NF** predicate takes as input a relation schema  $R$  and a set of associated functional dependencies and returns “yes” if the relation schema is in 3NF; “no” otherwise.

The **threenf** predicate takes as input a relation schema  $R$  and a set of associated functional dependencies and returns in the third parameter a decomposition of  $R$  into 3NF. The **threenf** predicate implements Algorithm 11.4 of the Elmasri/Navathe text.

As an example, consider the following file (named `e8.pl`) containing a relational schema and a set of functional dependencies:

```
schema([a,b,c,d]).
fds([[a,b],[c]], [[b],[d]], [[b,c],[a]]) .
```

The following screen capture illustrates the 3NF test as well as a 3NF decomposition of the relation schema.

```
?- ['dbd.pl'].
% dbd.pl compiled 0.02 sec, 30,740 bytes

Yes
?- ['e8.pl'].
% e8a.pl compiled 0.01 sec, 932 bytes

Yes
?- schema(R), fds(F), is3NF(R,F).

No
?- schema(R), fds(F), threenf(R,F,R3NF).

R = [a, b, c, d]
F = [[a, b], [c]], [[b], [d]], [[b, c], [a]]
R3NF = [[a, b, c], [b, d]]

Yes
?-
```

As can be seen in the above interaction with the DBD system, the schema  $R = ABCD$  is not in 3NF and is decomposed into 3NF sub-schemas  $ABC$  and  $BD$ .

### 5.3.9 **isBCNF(R,F)** and **bcnf(R,F,D)**

The **isBCNF** predicate takes as input a relation schema  $R$  and a set of associated functional dependencies and returns “yes” if the relation schema is in BCNF; “no” otherwise.

The **bcnf** predicate takes as input a relation schema  $R$  and a set of associated functional dependencies and returns in the third parameter a decomposition of  $R$  into BCNF. The **bcnf** predicate implements Algorithm 11.3 of the Elmasri/Navathe text.

As an example, consider the same `e8.pl` file used in the previous sub-section.

The following screen capture illustrates the BCNF test as well as a BCNF decomposition of the relation schema.

```
?- ['dbd.pl'].
% dbd.pl compiled 0.02 sec, 30,740 bytes

Yes
?- ['e8.pl'].
% e8.pl compiled 0.00 sec, 928 bytes

Yes
?- schema(R), fds(F), isBCNF(R,F).

No
?- schema(R), fds(F), bcnf(R,F,D).
Scheme to decompose = [a, b, c, d] Offending FD: [b]-->[d]
Final Result is:
[a, b, c]
[b, d]

R = [a, b, c, d]
F = [[[a, b], [c]], [[b], [d]], [[b, c], [a]]]
D = [[a, b, c], [b, d]];
```

As can be seen in the above interaction with the DBD system, the schema  $R = ABCD$  is not in BCNF and is decomposed into BCNF sub-schemas  $ABC$  and  $BD$ . Note that the `bcnf` predicate implementation also prints out the schema that is being decomposed along with the “offending” functional dependency.

## Exercises

For all the problems in this set of exercises, use the predicates in the DBD system to answer the questions.

### Chapter 10 (Elmasri/Navathe Text) Problems:

1. Consider the following two sets of functional dependencies:

$$F = \{A \rightarrow C, AC \rightarrow D, E \rightarrow AD, E \rightarrow H\}$$

$$G = \{A \rightarrow CD, E \rightarrow AH\}$$

Check whether they are equivalent.

2. Consider the relation schema

`EMP_DEPT(ename,ssn,bdate,address,dnumber,dname,dmgrssn)`

and the following set  $G$  of functional dependencies on EMP\_DEPT:

$ssn \rightarrow ename, bdate, address, dnumber$   
 $dnumber \rightarrow dname, dmgrssn$

Calculate the closures  $\{ssn\}^+$  and  $\{dnumber\}^+$ .

- Is the set of functional dependencies  $G$  in Exercise 2 minimal? If not, find a minimal set of functional dependencies that is equivalent to  $G$ .
- Consider the universal relation  $R = \{A, B, C, D, E, F, G, H, I\}$  and the set of functional dependencies:

$F = \{AB \rightarrow C, A \rightarrow DE, B \rightarrow F, F \rightarrow GH, D \rightarrow IJ\}$

What is the key for  $R$ ? Decompose  $R$  into 3NF relations.

- Repeat Exercise 4 for the same  $R$  but a different set of functional dependencies

$G = \{AB \rightarrow C, BD \rightarrow EF, B \rightarrow F, F \rightarrow GH, D \rightarrow IJ\}$

- Consider a relation  $R(A,B,C,D,E)$  with the following functional dependencies:

$F = \{AB \rightarrow C, CD \rightarrow E, DE \rightarrow B\}$

Is  $AB$  a candidate key of this relation? If not, is  $ABD$ ?

- Consider the relation  $R$ , which has attributes that hold schedules of courses and sections at a university;

$R = \{courseno, secno, offeringdept, credithours, courselevel, instructorssn, semester, year, Dayshours, roomno, noofstudents\}$

Suppose the following functional dependencies hold on  $R$ :

$courseno \rightarrow offeringdept, credithours, courselevel$

$courseno, secno, semester, year \rightarrow$   
 $dayshours, roomno, noofstudents, instructorssn$

$roomno, dayshours, semester, year \rightarrow$   
 $instructorssn, courseno, secno$

Find the candidate keys for  $R$  and decompose the relation into 3NF relations.

**Chapter 11 (Elmasri/Navathe Text) Problems:**

8. Consider the relation schema

LOTS (propertyid, countyname, lotno, area, price, taxrate)

and associated set of functional dependencies:

propertyid  $\rightarrow$  countyname, lotno, area, price, taxrate  
 countyname, lotno  $\rightarrow$  propertyid, area, price, taxrate  
 countyname  $\rightarrow$  taxrate  
 area  $\rightarrow$  price

Determine if the decomposition of LOTS into

LOTS1AX (propertyid, area, lotno)  
 LOTS1AY (area, countyname)  
 LOTS1B (area, price)  
 LOTS2 (countyname, taxrate)

is a lossless join decomposition?

9. Consider the universal relation  $R = \{ A, B, C, D, E, F, G, H, I \}$  and the set of functional dependencies:

$F = \{ AB \rightarrow C, A \rightarrow DE, B \rightarrow F, F \rightarrow GH, D \rightarrow IJ \}$

Determine the candidate keys for R, find a minimal cover for F, and decompose R into 3NF relations.

10. Consider the universal relation  $R = \{ A, B, C, D, E, F, G, H, I \}$  and the set of functional dependencies:

$G = \{ AB \rightarrow C, BD \rightarrow EF, B \rightarrow F, F \rightarrow GH, D \rightarrow IJ \}$

Determine the candidate keys for R, find a minimal cover for F, and decompose R into 3NF relations.

11. Consider the universal relation  $R = \{ A, B, C, D, E, F, G, H, I \}$  and the set of functional dependencies:

$F = \{ AB \rightarrow C, A \rightarrow DE, B \rightarrow F, F \rightarrow GH, D \rightarrow IJ \}$

Decompose R into BCNF relations. Repeat the exercise for

$G = \{ AB \rightarrow C, BD \rightarrow EF, B \rightarrow F, F \rightarrow GH, D \rightarrow IJ \}$



12. Consider a relation  $R(A,B,C,D,E)$  with the following functional dependencies:

$$F = \{AB \rightarrow C, CD \rightarrow E, DE \rightarrow B\}$$

Determine the candidate keys for  $R$  and decompose  $R$  into 3NF relations and into BCNF relations.

Repeat the exercise for

$R = \{\text{courseno}, \text{secno}, \text{offeringdept}, \text{credithours}, \text{courselevel}, \text{instructorssn}, \text{semester}, \text{year}, \text{Dayshours}, \text{roomno}, \text{noofstudents}\}$

and  $F$  containing the following functional dependencies:

$\text{courseno} \rightarrow \text{offeringdept}, \text{credithours}, \text{courselevel},$

$\text{courseno}, \text{secno}, \text{semester}, \text{year} \rightarrow \text{dayshours}, \text{roomno}, \text{noofstudents}, \text{instructorssn}$

$\text{roomno}, \text{dayshours}, \text{semester}, \text{year} \rightarrow \text{instructorssn}, \text{courseno}, \text{secno}$

13. Consider the following decompositions for the relation schema

$R(a,b,c,d,e,f,g,h,i)$

and the set of associated functional dependencies:

$$F = \{ab \rightarrow c, a \rightarrow de, b \rightarrow f, f \rightarrow gh, d \rightarrow ij\}.$$

Determine whether each of the following decompositions has the (a) dependency preservation property, and (b) the lossless join property with respect to  $F$ :

- i.  $D1 = \{R1, R2, R3, R4, R5\};$   
 $R1 = \{a,b,c\}, R2 = \{a,d,e\}, R3 = \{b,f\}, R4 = \{f,g,h\},$   
 $R5 = \{d,i,j\}.$
- ii.  $D2 = \{R1, R2, R3\};$   
 $R1 = \{a,b,c,d,e\}, R2 = \{b,f,g,h\}, R3 = \{d,i,j\}$
- iii.  $D3 = \{R1, R2, R3, R4, R5\};$   
 $R1 = \{a,b,c,d\}, R2 = \{d,e\}, R3 = \{b,f\}, R4 = \{f,g,h\},$   
 $R5 = \{d,i,j\}.$

14. Consider the relation `REFRIG(modelno, year, price, manufplant, color)` which is abbreviated as `REFRIG(m, y, p, mp, c)` and the set of functional dependencies

$$F = \{m \rightarrow mp, \{m, y\} \rightarrow p, mp \rightarrow c\}$$

- iv. Evaluate each of the following as candidate keys:  $\{m\}$ ,  $\{m, y\}$ ,  $\{m, c\}$ .
- v. Test if `REFRIG` is in 3NF? In BCNF?
- vi. Test if the decomposition  $D = \{R1(m, y, p), R2(m, mp, c)\}$  is lossless.

## CHAPTER 6

### Projects

This chapter presents several projects to be completed by the student using Java/JDBC and Oracle database or using PHP and MySQL database. These projects may be assigned as group projects with teams of two or three students. A written documentation as well as a class presentation of the project may be required.

#### 6.1 STUDENT REGISTRATION System (GoLunar)

Consider the following relational database and sample data for the student registration database (written in Oracle SQL):

```
drop table students cascade constraints;
create table students (
    sid      number(4) primary key,
    password number(5),
    fname    varchar2(20),
    lname    varchar2(20),
    sType    varchar2(5) check (sType in ('GRAD','UGRAD')),
    major    char(4) check (major in ('CSC','MATH','POLS','HIST')),
    gradAssistant char(1) check (gradAssistant in ('Y','N')),
    inState  char(1) check (inState in ('Y','N'))
);
insert into students values
    (1111,1111,'John','Davison','UGRAD','CSC','N','Y');
insert into students values
    (2222,2222,'Jacob','Oram','UGRAD','CSC','N','N');
insert into students values
    (3333,3333,'Ashish','Bagai','GRAD','CSC','Y','N');
insert into students values
    (4444,4444,'Joe','Harris','GRAD','CSC','N','Y');
insert into students values
    (5555,5555,'Andy','Blignaut','GRAD','CSC','N','Y');
insert into students values
    (6666,6666,'Pommie','Mbangwa','GRAD','CSC','N','Y');
insert into students values
    (7777,7777,'Ian','Healy','GRAD','CSC','N','Y');
insert into students values
    (8888,8888,'Dougie','Marillier','GRAD','CSC','N','Y');
--
drop table staff cascade constraints;
create table staff (
    tid  number(4) primary key,
    password number(5),
    fname varchar2(20),
    lname varchar2(20),
    staffType varchar2(10) check (staffType in ('REGISTRAR','DEPARTMENT'))
);
insert into staff values
    (1000,1000,'Venette','Rice','DEPARTMENT');
```

```

insert into staff values
  (2000,2000,'Alison','Payne','REGISTRAR');
--
create or replace view lunarUsers as
  (select sid uid, password, 'STUDENT' uType
   from students) union
  (select tid uid, password, staffType uType
   from staff);
--
drop table courses cascade constraints;
create table courses (
  cprefix char(4),
  cno      number(4),
  ctitle   varchar2(50),
  chours   number(2),
  primary key (cprefix,cno)
);
insert into courses values ('CSC',1010,'Computers and Applications',3);
insert into courses values ('CSC',2010,'Introduction to Computer Science',3);
insert into courses values ('CSC',2310,'Intro to Programming in Java',3);
insert into courses values ('CSC',2311,'Introduction to Programming in C++',3);
insert into courses values ('CSC',3410,'Data Structures',3);
insert into courses values ('CSC',3210,'Computer Organization',3);
insert into courses values ('CSC',3320,'Systems Programming in Unix and C',3);
insert into courses values ('MATH',2211,'Calculus I',5);
insert into courses values ('MATH',2212,'Calculus II',5);
insert into courses values ('MATH',2420,'Discrete Mathematics',3);
insert into courses values ('CSC',6220,'Networks',4);
insert into courses values ('CSC',8220,'Advanced Networks',4);
insert into courses values ('CSC',6710,'Database',4);
insert into courses values ('CSC',8710,'Advanced Database',4);
insert into courses values ('CSC',6820,'Graphics',4);
insert into courses values ('CSC',8820,'Advanced Graphics',4);
insert into courses values ('POLs',1200,'Intro Political Sci',3);
--
drop table sections cascade constraints;
create table sections (
  term      char(2) check (term in ('FA','SP','SU')),
  year      number(4),
  crn       number(5),
  cprefix   char(4),
  cno       number(4),
  section   number(2),
  days      char(6),
  startTime char(5), -- example 08.15, 13.30 etc.
  endTime   char(5),
  room      varchar2(10),
  cap       number(3),
  instructor varchar2(30),
  auth      char(1) check (auth in ('Y','N')),
  primary key (term,year,crn),
  foreign key (cprefix,cno) references courses
);
--
insert into sections values
('SU',2002,10101,'CSC',1010,1,'MWF','09.00','09.50','105G',35,'Bhola','N');
insert into sections values

```

```

('SU',2002,10701,'POLS',1200,1,'TR','09.00','09.50','205Sp',25,'Jones','N');
--
insert into sections values
('FA',2002,10101,'CSC',2010,1,'MWF','09.00','09.50','105G',35,'Bhola','N');
insert into sections values
('FA',2002,10102,'CSC',2010,2,'MWF','10.00','10.50','105CS',40,'Henry','N');
insert into sections values
('FA',2002,10103,'CSC',2310,1,'MWF','12.00','12.50','106G',30,'Henry','N');
insert into sections values
('FA',2002,10104,'CSC',2311,1,'MWF','15.00','15.50','205G',35,'Liu','N');
insert into sections values
('FA',2002,10201,'CSC',6220,1,'TR','19.00','20.40','405G',25,'Hundewale','N');
insert into sections values
('FA',2002,10202,'CSC',6710,1,'TR','16.00','17.15','115CS',25,'Madiraju','N');
insert into sections values
('FA',2002,10203,'CSC',8820,1,'MWF','09.00','09.50','605G',25,'Owen','N');
insert into sections values
('FA',2002,10301,'MATH',2211,1,'TR','11.00','12.50','305G',35,'Li','N');
insert into sections values
('FA',2002,10302,'MATH',2211,2,'MWF','09.00','10.50','106GB',35,'Davis','N');
--
--This data will be loaded into the database in your application program
--insert into sections values
--('SP',2003,10101,'CSC',2010,1,'MWF','09.00','09.50','105G',35,'Bhola','N');
--insert into sections values
--('SP',2003,10102,'CSC',2010,2,'MWF','10.00','10.50','105CS',40,'Henry','N');
--insert into sections values
--('SP',2003,10103,'CSC',2310,1,'MWF','12.00','12.50','106G',30,'Henry','N');
--insert into sections values
--('SP',2003,10104,'CSC',2311,1,'MWF','15.00','15.50','205G',35,'Liu','N');
--insert into sections values
--('SP',2003,10201,'CSC',6220,1,'TR','19.00','20.40','405G',25,'Hundewale','N');
--insert into sections values
--('SP',2003,10202,'CSC',6710,1,'TR','16.00','17.15','115CS',25,'Madiraju','N');
--insert into sections values
--('SP',2003,10203,'CSC',8220,1,'MWF','09.00','09.50','605G',25,'Bourgeois','Y');
--insert into sections values
--('SP',2003,10301,'MATH',2211,1,'TR','11.00','12.50','305G',35,'Li','N');
--insert into sections values
--('SP',2003,10302,'MATH',2211,2,'MWF','09.00','10.50','606GB',35,'Miller','N');
--insert into sections values
--('SP',2003,10303,'MATH',2212,1,'MWF','09.00','10.50','706GB',35,'Davis','N');
--insert into sections values
--('SP',2003,10304,'MATH',2420,1,'TR','14.00','14.50','106GB',35,'Domke','N');
--insert into sections values
--('SP',2003,10405,'CSC',8710,1,'MW','17.30','18.45','206GB',35,'Dogdu','N');
--insert into sections values
--('SP',2003,10406,'CSC',8820,1,'TR','19.15','20.55','306GB',3,'Owen','N');
--
drop table enrolls cascade constraints;
create table enrolls (
    sid      number(4),
    term     char(2) check (term in ('FA','SP','SU')),
    year     number(4),
    crn      number(5),
    grade    char(2) check (grade in ('A','B','C','D','F','I','IP','S','U')),
    primary key (sid,term,year,crn),
    foreign key (sid) references students,
    foreign key (term,year,crn) references sections
);

```

```

--
insert into enrolls values (1111,'SU',2002,10101,'A');
insert into enrolls values (1111,'SU',2002,10701,'C');
--
insert into enrolls values (1111,'FA',2002,10101,null);
insert into enrolls values (1111,'FA',2002,10103,null);
insert into enrolls values (1111,'FA',2002,10301,null);
insert into enrolls values (3333,'FA',2002,10201,null);
insert into enrolls values (3333,'FA',2002,10202,null);
insert into enrolls values (3333,'FA',2002,10203,null);
--
drop table authorizations cascade constraints;
create table authorizations (
    term    char(2) check (term in ('FA','SP','SU')),
    year    number(4),
    crn     number(5),
    sid     number(4),
    authType char(4) check (authType in ('OVFL','AUTH')),
    primary key (term,year,crn,sid,authType),
    foreign key (sid) references students,
    foreign key (term,year,crn) references sections
);
--
drop table fixedFee cascade constraints;
create table fixedFee (
    feeName varchar2(30) primary key,
    fee number(5,2)
);
--
insert into fixedFee values ('Technology Fee',75.00);
insert into fixedFee values ('Health Fee',30.00);
insert into fixedFee values ('Activity Fee',65.00);
insert into fixedFee values ('Transportation Fee',25.00);
--
drop table variableFeeRate cascade constraints;
create table variableFeeRate (
    sType varchar2(6)
        check (sType in ('GRAD','UGRAD')),
    inOrOutOfState varchar2(10)
        check (inOrOutOfState in ('INSTATE','OUTOFSTATE')),
    fee number(6,2),
    primary key (sType,inOrOutOfState)
);
--
insert into variableFeeRate values ('GRAD','INSTATE',125.00);
insert into variableFeeRate values ('GRAD','OUTOFSTATE',500.00);
insert into variableFeeRate values ('UGRAD','INSTATE',100.00);
insert into variableFeeRate values ('UGRAD','OUTOFSTATE',400.00);

```

The database consists of the following tables:

1. Students: This table records information about students. The gradAssistant attribute records whether the student is a graduate assistant or not. The graduate assistants automatically qualify for a full tuition waiver (they still have to pay the fixed fee). The

`instate` attribute records whether the student is an in-state student or not. Again, this has an impact on the fees the student would pay.

2. **Staff:** This table records information about staff users of the system. There are two categories of staff: “Registrar” and “Department”. These users would have different capabilities and functions in the application to be developed.

**Note:** A view called `lunarUsers` is created to provide a simple way to authenticate users of the system.

3. **Courses:** This table records information about courses in the university catalog which includes course number, title and credit hours. The credit hours value will be used in calculating the GPA in the student’s transcript.
4. **Sections:** This table records the course offerings for each term and includes the term, year, and course record number (`crn`), a unique number assigned to course offerings for a specific term and year. The table also includes start time and end time and meeting days as well as the name of the instructor. Finally, this table records a boolean value (yes or no) called `auth` to indicate if the registration for this course is open to all or is done only by authorization.
5. **Enrolls:** This table records information about which student has registered for which course offering.
6. **Authorizations:** This table records authorizations given to students for specific course offerings. Two types of authorizations are given: `OVFL` for overflow, i.e. allows students to register in a course offering that does not have any open seats, and `AUTH` for authorization to register in a course offering that is designated as a authorization only course offering.
7. **FixedFee:** This table records information about all fixed fees a student is required to pay each term they register.
8. **VariableFeeRate:** This table records per credit hour fee rate for different categories of students (graduate vs undergraduate students and in-state vs out-of-state students).

You will implement a University Registration System in Java using JDBC.

There are 3 kinds of database users:

1. **Registrar Staff:** These users will have the ability to load the database tables, make changes to courses, sections, fee details etc.
2. **Department Staff:** These users will have the ability to authorize students into sections, overflow students into sections, add assistantship information to the system, generate class lists etc.
3. **Student:** These users will be able to register for classes, see their schedules, see fee detail, see transcripts etc.

The following real-world constraints need to be enforced by your Java program:

1. Undergraduate students are not allowed to register for graduate courses numbered 6000 and above.
2. Students should not be allowed to register for a class which is FULL unless they have an overflow.
3. Students should not be allowed to register for a class which is listed as AUTHORIZATION ONLY unless they have an authorization.
4. Undergraduate students are not allowed to register for more than 20 hours in a semester and the limit for graduate students is 15.
5. Students cannot register for two classes that overlap in meeting time.

The Java application will be a terminal-based program that has the following interactions with the users. Based on the username, the program should determine the type of user and provide the appropriate menu.

### Department Staff Menu:

```
$ java GoLunar OracleId
Oracle Password:xxxxxx
Semester (e.g. FA2003,SP2003,SU2003): SP2003
Username: 1000
Password:

*****
***                                     ***
***      Welcome to the GoLunar - Online Registration System      ***
***      Venette Rice - Department Staff                          ***
***                                     ***
*****

1. Authorize Student into Section
2. Overflow Student into Section
3. Add Assistantship on System
4. Generate Class List
q. Quit
```

Type in your option:

### Option 1 Interface:

```
CRN:10101
SID:1111
Student John Davison authorized into CRN 10101, CSC 2010.
OR
No need to authorize - This section does not need authorization.
```



**Option 2 Interface:**

```

CRN:10101
SID:1111
No need to overflow - Space still available in this section.
OR
Student John Davison overflowed into CRN 10101, CSC 2010.

```

**Option 3 Interface:**

```

Student Id: 3333
Ashish Bagai (3333) has been added to the Assistantship List.

```

**Option 4 Interface:**

```

CRN: 10101

CSC 2010, Introduction to Computer Science
SP 2003
Instructor: Bhola

```

SID	LNAME	FNAME
1111	Davison	John
2222	Oram	Jacob
...		
...		

**Student Menu:**

```

$ java GoLunar OracleId
Oracle Password:xxxxxx
Semester (e.g. FA2003,SP2003,SU2003): SP2003
Username: 1111
Password:
*****
***                                     ***
***   Welcome to the GoLunar - Online Registration System   ***
***           John Davison - Student                        ***
***                                                         ***
*****

```

1. Add a Section
2. Drop a Section
3. See Schedule for a Term
4. See Fee detail
5. See Transcript
- q. Quit

Type in your option: 1

**Option 1 Interface:**

CRN: 10101  
 CSC2010, Introduction to Computer Science ADDED.  
 OR  
 Appropriate Error Message.

**Option 2 Interface:**

CRN: 10101  
 CSC2010, Introduction to Computer Science DROPPED.  
 OR  
 Appropriate Error Message.

**Option 3 Interface:**

Term: FA2002

CRN	Course	Title	Days	Time	Room	Instructor
10101	CSC2010	Introduction to Computer Science	MWF	09.00-09.50	105G	Bhola
...						
...						

**Option 4 Interface:**

Term: sp2003

Spring 2003

Tuition - InState	
(12 hours)	1,500.00
Technology Fee	75.00
Health Fee	30.00
Activity Fee	65.00
Transportation Fee	25.00
	-----
	1,695.00
	-----

**Option 5 Interface:**

Summer 2002

CSC	1010	10101	Computers and Applications	3	A	12.00
POLS	1200	10701	Intro Political Sci	3	C	6.00
Semester GPA: 3.00			GPA: 3.00			

Fall 2002

CSC	2010	10101	Introduction to Computer Science	3	A	12.00
Csc	2310		Introduction to Programming in Java	3	A	12.00
Math	2211		Calculus I	5	B	15.00
Semester GPA: 3.54			GPA: 3.35			

...  
 ...



...

**Option 5** is similar to student option except here the system should accept student id as input and display that student's transcript.

**Option 6** is similar to student options except here the system should accept student id (in addition to term) as input and display that student's term schedule and fee detail for the particular term.

Sample files for loading data in the Registrar's options are available in

#### **sections.dat**

```
SP
2003
10101,CSC,2010,1,MWF,09.00,09.50,105G,35,Bhola,N
10102,CSC,2010,2,MWF,10.00,10.50,105CS,40,Henry,N
10103,CSC,2310,1,MWF,12.00,12.50,106G,30,Henry,N
10104,CSC,2311,1,MWF,15.00,15.50,205G,35,Liu,N
10201,CSC,6220,1,TR,19.00,20.40,405G,25,Hundewale,N
10202,CSC,6710,1,TR,16.00,17.15,115CS,25,Madiraju,N
10203,CSC,8220,1,MWF,09.00,09.50,605G,25,Bourgeois,Y
10301,MATH,2211,1,TR,11.00,12.50,305G,35,Li,N
10302,MATH,2211,2,MWF,09.00,10.50,606GB,35,Miller,N
10303,MATH,2212,1,MWF,09.00,10.50,706GB,35,Davis,N
10304,MATH,2420,1,TR,14.00,14.50,106GB,35,Domke,N
10405,CSC,8710,1,MW,17.30,18.45,206GB,35,Dogdu,N
10406,CSC,8820,1,TR,19.15,20.55,306GB,3,Owen,N
```

#### **grades.dat**

```
FA
2002
1111,10101,A
1111,10103,A
1111,10301,B
3333,10201,B
3333,10202,B
3333,10203,A
```

## **6.2 Online Book Store Database System**

Consider the following relational database schema written in Oracle SQL for an online book store application:

```
drop table books cascade constraints;
create table books (
  isbn char(10),
  author varchar2(100) not null,
  title varchar2(128) not null,
  price number(7,2) not null,
  subject varchar2(30) not null,
  primary key (isbn)
);
```

```

drop table members cascade constraints;
create table members (
    fname varchar2(20) not null,
    lname varchar2(20) not null,
    address varchar2(50) not null,
    city varchar2(30) not null,
    state varchar2(20) not null,
    zip number(5) not null,
    phone varchar2(12),
    email varchar2(40),
    userid varchar2(20),
    password varchar2(20),
    creditcardtype varchar2(10)
        check(creditcardtype in ('amex','discover','mc','visa')),
    creditcardnumber char(16),
    primary key (userid)
);
drop table orders cascade constraints;
create table orders (
    userid varchar2(20) not null,
    ono number(5),
    received date not null,
    shipped date,
    shipAddress varchar2(50),
    shipCity varchar2(30),
    shipState varchar2(20),
    shipZip number(5),
    primary key (ono),
    foreign key (userid) references members
);
drop table odetails cascade constraints;
create table odetails (
    ono number(5),
    isbn char(10),
    qty number(5) not null,
    price number(7,2) not null,
    primary key (ono,isbn),
    foreign key (ono) references orders,
    foreign key (isbn) references books
);
drop table cart cascade constraints;
create table cart (
    userid varchar2(20),
    isbn char(10),
    qty number(5) not null,
    primary key (userid,isbn),
    foreign key (userid) references members,
    foreign key (isbn) references books
);

```

The database consists of five tables:

1. Books: This table records information about the books on sale in the book store. Each book is classified under a “subject” to enable subject searches.

2. **Members:** This table records information about members of the application. Each member chooses their own user id and password at the time of registration.
3. **Orders:** This table records information about orders placed by members place orders. The orders may contain one or more books and the details of the order are kept in a separate table. A unique order number is generated by the system.
4. **OrderDetails:** This table records information about each order including the isbn and quantity of books in the order.
5. **Cart:** This table contains isbn and quantity of each book placed in the shopping cart of a member. Once a member checks out, the shopping cart is emptied and an order is created.

The book store application should be developed as a terminal application in Java and should be implemented in three phases:

**Phase I** of the project requires:

- (a) Each student to create data for approximately 10 books for two different subjects (the subjects may be assigned by the instructor of the class to each student). The instructor may then consolidate the data into a large data set and give it out to the entire class. This is an easy way to create a large data set of books.
- (b) Each student to build program the following interface implementing only the member registration and member login functions:

```
$ java OnlineBookStore
```

```
*****
***                                     ***
***           Welcome to the Online Book Store           ***
***                                     ***
*****
          1. Member Login
          2. New Member Registration
          q. Quit
```

Type in your option: 2

```
Welcome to the Online Book Store
  New Member Registration
```

```
Enter first name: Raj
Enter last name: Sunderraman
Enter street address: 123 Main Street
Enter city: Atlanta
Enter state: GA
Enter zip: 30303
Enter phone: 555-1212
Enter email address: raj@cs.gsu.edu
Enter userID: raj
Enter password: raj
```

Do you wish to store credit card information(y/n): y  
 Enter type of Credit Card(amex/visa): amex  
 Enter Credit Card Number: 121212121212121  
 Invalid Entry  
 Enter Credit Card Number: 1212121212121212  
 Invalid Entry  
 Enter Credit Card Number: 12121212121212

You have registered successfully.

Name: Raj Sunderraman  
 Address: 123 Main Street  
 City: Atlanta GA 30303  
 Phone: 555-1212  
 Email: raj@cs.gsu.edu  
 UserID: raj  
 Password: raj  
 CreditCard Type: amex  
 CreditCard Number: 12121212121212  
 Press Enter to go back to Menu

```

*****
***                                     ***
***           Welcome to the Online Book Store           ***
***                                     ***
*****
1. Member Login
2. New Member Registration
q. Quit
  
```

Type in your option: 1

Enter userID: raj  
 Enter password: raj

```

*****
***                                     ***
***           Welcome to Online Book Store           ***
***           Member Menu                             ***
***                                     ***
*****
1. Browse by Subject
2. Search by Author/Title/Subject
3. View/Edit Shopping Cart
4. Check Order Status
5. Check Out
6. One Click Check Out
7. View/Edit Personal Information
8. Logout
  
```

Type in your option: 8  
 You have successfully logged out.

**Phase II** of the project requires the student to implement the following member functions:

**1. Browse by Subject:** This option should first list all subjects alphabetically; It then allows user to choose one subject; Upon choosing a subject, the program displays book details (2 books at a time on a screen); The option allows user to

- (a) enter isbn to put in cart;
- (b) press ENTER to return to main menu
- (c) press n ENTER to continue browsing

User Interface follows:

Type in your option: 1

- 1. Cooking
- 2. Jokes
- 3. Sports

Enter your choice: 3

5 books available on this Subject

Author: Dom Parker  
 Title: 1,001 Baseball Questions Your Friends Can't Answer  
 ISBN: 0451191323  
 Price: 22.46  
 Subject Sports

Author: Timothy Jacobs  
 Title: 100 Athletes Who Shaped Sports History  
 ISBN: 0912517131  
 Price: 32.56  
 Subject Sports

Enter ISBN to add to Cart or  
 n Enter to browse or  
 ENTER to go back to menu:  
 0451191323  
 Enter quantity: 2

Author: Michael Dregni  
 Title: 100 Years of Fishing  
 ISBN: 0896584305  
 Price: 15.95  
 Subject Sports

Author: David Claerbaut  
 Title: The 1999 NBA Analyst: The Science of Hoops Magic  
 ISBN: 0878332103  
 Price: 20.95  
 Subject Sports



Enter ISBN to add to Cart or  
 n Enter to browse or  
 ENTER to go back to menu: n

Author: Sports Collectors Digest  
 Title: 1999 Sports Collectors Almanac (Serial)  
 ISBN: 0987654234  
 Price: 17.56  
 Subject Sports

Enter ISBN to add to Cart or  
 n Enter to browse or  
 ENTER to go back to menu:  
 0987654234  
 Enter quantity: 1

## 6. One Click Check Out

This option should move items in the cart to the `order` and `odetails` tables. Cart is emptied in the process and an invoice is printed. Shipping address is same as member address in this option.

User Interface follows:

Invoice for Order no.117

Shipping Address	Billing address
Name: Raj Sunderraman	Name: Raj Sunderraman
Address: 123 Main Street	Address: 123 Main Street
Atlanta	Atlanta
GA 33333	GA 33333

ISBN	Title	\$	Qty	Total
0451191323	1,001 Baseball Questions Your Friends Can't Answer	22.45	1	22.45
0987654234	1999 Sports Collectors Almanac (Serial)	17.55	1	17.55
Total =				\$40.01

Press enter to go back to Menu

## 4. Check Order Status

This option should list all orders for member and should allow user to choose one order to see details. User Interface follows:

Orders placed by Raj Sunderraman

ORDER NO	RECEIVED DATE	SHIPPED DATE
117	3-1-2001	3-3-2001

Enter the Order No to display its details or (q) to quit: 117

## Details for Order no.117

## Shipping Address

Name: Raj Sunderraman  
 Address: 123 Main Street  
 Atlanta  
 GA 33333

## Billing address

Name: Raj Sunderraman  
 Address: 123 Main Street  
 Atlanta  
 GA 33333

ISBN	Title	\$	Qty	Total
0451191323	1,001 Baseball Questions Your Friends Can't Answer	22.45	1	22.45
0987654234	1999 Sports Collectors Almanac(Serial)	17.55	1	17.55
Total =				\$40.01

Press Enter to go back to Menu

**Phase III** of the project requires the student to implement the following member functions:

## 2. Search by Author/Title

This option should provide 3 sub-options:

1. Author Search
2. Title Search
3. Go Back to Member Menu

In the Author or Title search sub-option, the user may enter a substring and the system should respond with all books which contain the substring in the title/author. The display should be done 2 books at a time on a screen.

The system should also allow user to enter isbn to put in cart;  
 to press ENTER to return to main menu  
 to press n ENTER to continue browsing

User Interface follows:

1. Author Search
2. Title Search
3. Go Back to Member Menu

Type in your option: 2

Enter title or part of the title: cook

2 books found

Author: Irma S. Rambauer  
 Title: Joy of Cooking  
 ISBN: 0452279232  
 Price: 15.25  
 Subject Cooking

Author: Jennifer E. Darling  
 Title: Better Homes and Gardens New Cook Book  
 ISBN: 0696201887

Price: 21.96  
Subject Cooking

Enter ISBN to add to Cart or  
Enter to browse or  
n ENTER to return to menu: 0696201887  
Enter quantity: 1

1. Author Search
2. Title Search
3. Go Back to Member Menu

Type in your option: 2

Enter title or part of the title: Computer  
0 books found

Enter ISBN to add to Cart or  
Enter to browse or  
n ENTER to return to menu: n

1. Author Search
2. Title Search
3. Go Back to Member Menu

Type in your option: 1

Enter name or part of the name: am  
1 books found

Author: Irma S. Rambauer  
Title: Joy of Cooking  
ISBN: 0452279232  
Price: 15.25  
Subject Cooking

Enter ISBN to add to Cart or  
Enter to browse or  
n ENTER to return to menu: 0452279232  
Enter quantity: 2

1. Author Search
2. Title Search
3. Go Back to Member Menu

Type in your option: 3

### **3. View/Edit Shopping Cart**

This option should show the contents of the cart; It should then provide options to delete items or edit (change quantity) items. User Interface (for delete and update cart) follows:

Current Cart Contents:

ISBN	Title	\$	Qty	Total
0696201887	Better Homes and Gardens New Cook Book	21.95	1	21.95
0452279232	Joy of Cooking	15.25	2	30.50
Total =				\$52.45

Enter d to delete item  
 e to edit cart or  
 q to go back to Menu: d  
 Enter isbn of item: 0452279232  
 Delete Item Completed  
 Press enter to go back to Menu

Type in your option: 3

Current Cart Contents:

ISBN	Title	\$	Qty	Total
0696201887	Better Homes and Gardens New Cook Book	21.95	1	21.95
Total =				\$21.95

Enter d to delete item  
 e to edit cart or  
 q to go back to Menu: e  
 Enter isbn of item: 0696201887  
 Enter new Quantity: 2  
 Edit Item Completed  
 Press enter to go back to Menu

## 5. Check Out

This option should display invoice; request user of they want to provide shipping address (if no use current address in file for shipping); Also this option should ask if a new credit card should be used. Finally, an invoice should be printed. User Interface follows:

Current Cart Contents:

ISBN	Title	\$	Qty	Total
0696201887	Better Homes and Gardens New Cook Book	21.95	2	43.91
Total				\$43.91

Proceed to check out(Y/N): y  
 Do you want to enter new shipping address(y/n): y  
 Enter first name: John  
 Enter last name: Smith  
 Enter street: 123 Elm Street  
 Enter city: Atlanta

Enter state: GA  
 Enter zip: 11111  
 Do you want to enter new CreditCard Number(y/n): n

Invoice for Order no.118

Shipping Address	Billing address
Name: John Smith	Name: Raj Sunderraman
Address: 123 Elm Street	Address: 123 Main Street
Atlanta	Atlanta
GA 11111	GA 33333

ISBN	Title	\$	Qty	Total
0696201887	Better Homes and Gardens New Cook Book	21.95	2	43.91
Total =				\$43.91

Press enter to go back to Menu

### 6.3 Online Shopping System

Using PHP and MySQL, implement a Web-based application for an online video-store. The online video-store maintains an inventory of DVDs. Customers become member of this online store. They are able to search for DVDs of their interest and add DVDs to their shopping cart. At any time, they are able to edit the shopping cart and also are able to check out. The initial login screen shown in the Figure 6.1 allows an existing customer to sign in or a new customer to register. The new customer registration screen is shown in Figure 6.2. Upon successful sign-in, the 3-frame page shown in Figure 6.3 is displayed. As one can see, there are six different options for the customer:

1. **Search by Keyword:** This option allows the customer to perform a keyword search of the DVD titles (substring; case insensitive comparison). Successful matches are shown on the right frame (Figure 6.4). The customer may then choose certain quantities of the DVDs and add them to the shopping cart. Upon successful addition to the shopping cart, a message should be shown on the right frame.
2. **View/Edit Cart:** Upon clicking this option, the system should display the shopping cart on the right side frame (Figure 6.5). Here, the customer may edit the shopping cart by changing quantities including replacing a value with a zero. Upon submission, the cart should be updated and a message should be displayed.
3. **Update Profile:** This option brings up the customers profile (Figure 6.6) on the right frame. The user may modify any field and submit. Upon successful update, a message should be displayed.
4. **Check Order Status:** This option allows the customer to see all their orders (Figure 6.7). Upon clicking the order number link the details for that order should be displayed in a tabular format.
5. **Check Out:** Upon clicking this link, the system should empty the cart and move the items into the `orders` and `odetails` tables. An invoice (Figure 6.8) should be printed on the right frame.

6. **Logout:** Upon logout, the system should display 3 Options (Figure 6.9) to the user. The user may check out, save cart and logout or empty cart and logout. Upon checkout a similar action should take place as earlier. Upon the other two options, appropriate action should take place and a message should be displayed. If the cart was empty to begin with these 3 options should not be shown and the customer should be logged out.

The database schema for the online shopping cart example is shown below:

```

create table parts(
    pno         integer(5) not null,
    pname       varchar(30),
    qoh         integer,
    price       decimal(6,2),
    olevel      integer,
    primary key (pno));

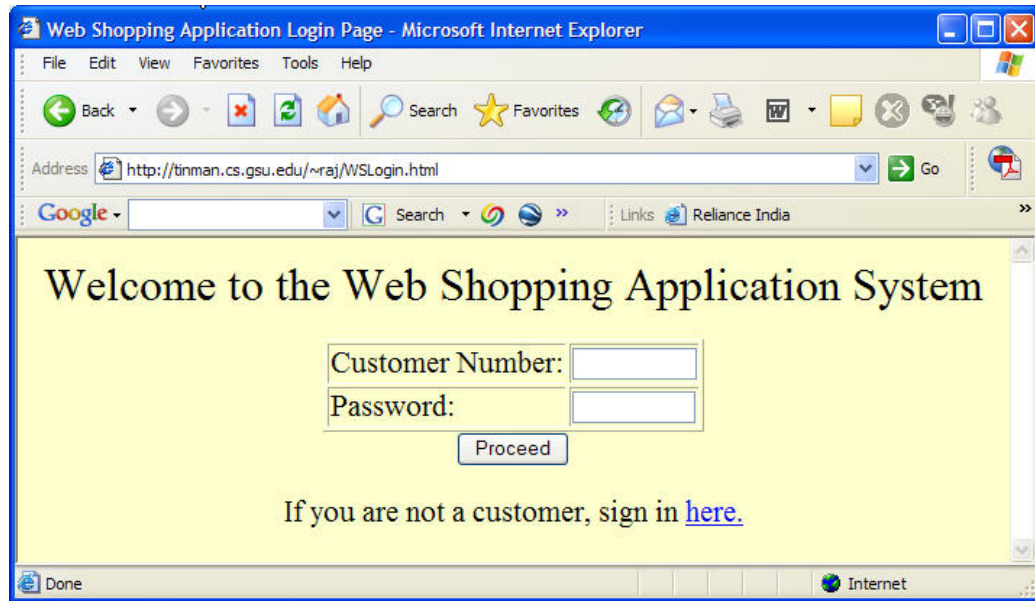
create table customers (
    cno         integer(10) not null auto_increment=100,
    cname       varchar(30),
    street      varchar(50),
    city        varchar(30),
    state       varchar(30),
    zip         integer(5),
    phone       char(12),
    email       varchar(50),
    password    varchar(15),
    primary key (cno));

create table cart(
    cartno      integer(10) not null auto_increment,
    cno         integer(10) not null,
    pno         integer(5) not null,
    qty         integer,
    primary key (cartno, pno),
    foreign key (cno) references customers,
    foreign key (pno) references parts);

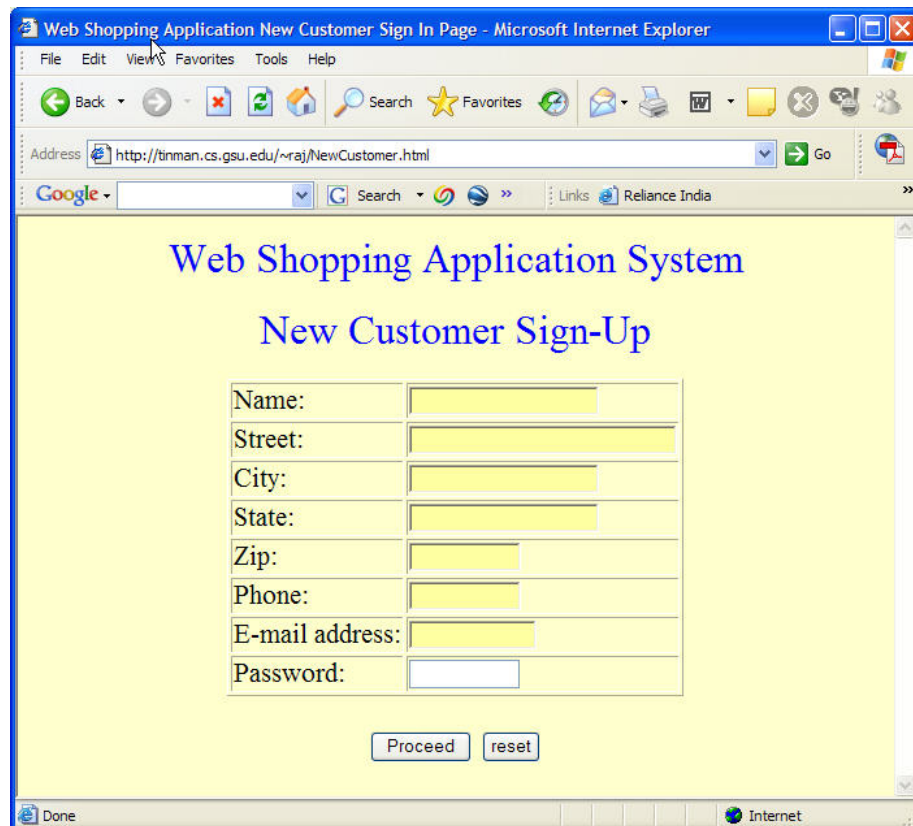
create table orders (
    ono         integer(5) not null auto_increment=1000,
    cno         integer(10),
    received    date,
    shipped     date,
    primary key (ono),
    foreign key (cno) references customers);

create table odetails (
    ono         integer(5) not null,
    pno         integer(5) not null,
    qty         integer,
    primary key (ono,pno),
    foreign key (ono) references orders,
    foreign key (pno) references parts);

```



**Figure 6.1 Web Shopping – Initial Screen**



**Figure 6.2 Web Shopping – New Customer Registration**

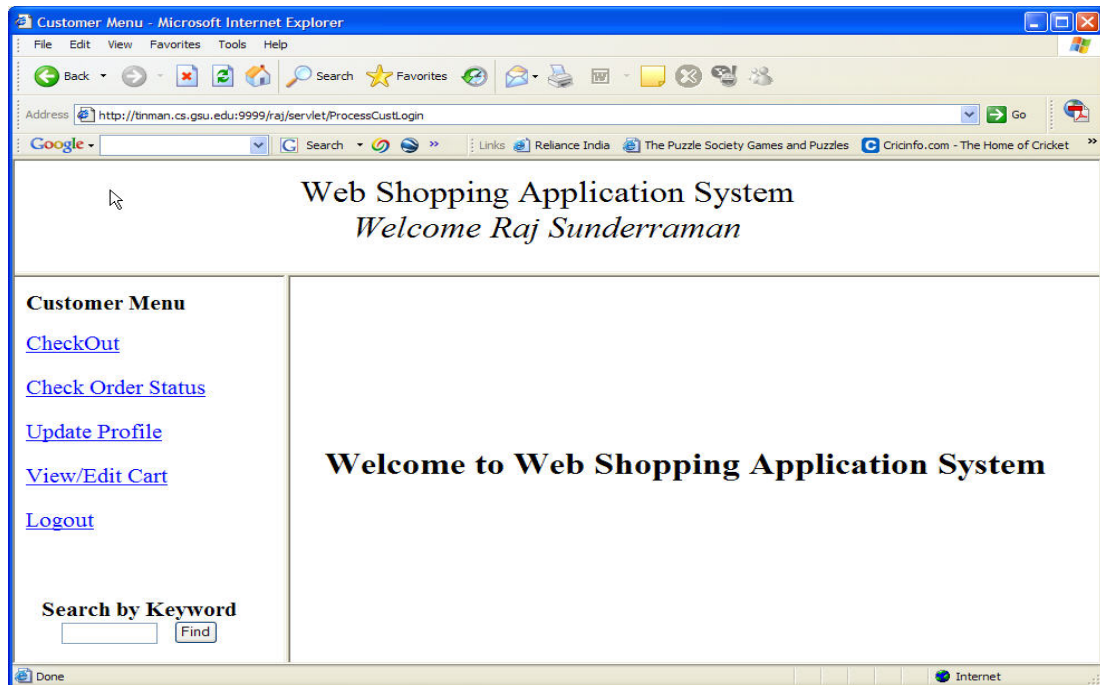


Figure 6.3: Web Shopping – Successful Sign-In 3-Frame Page

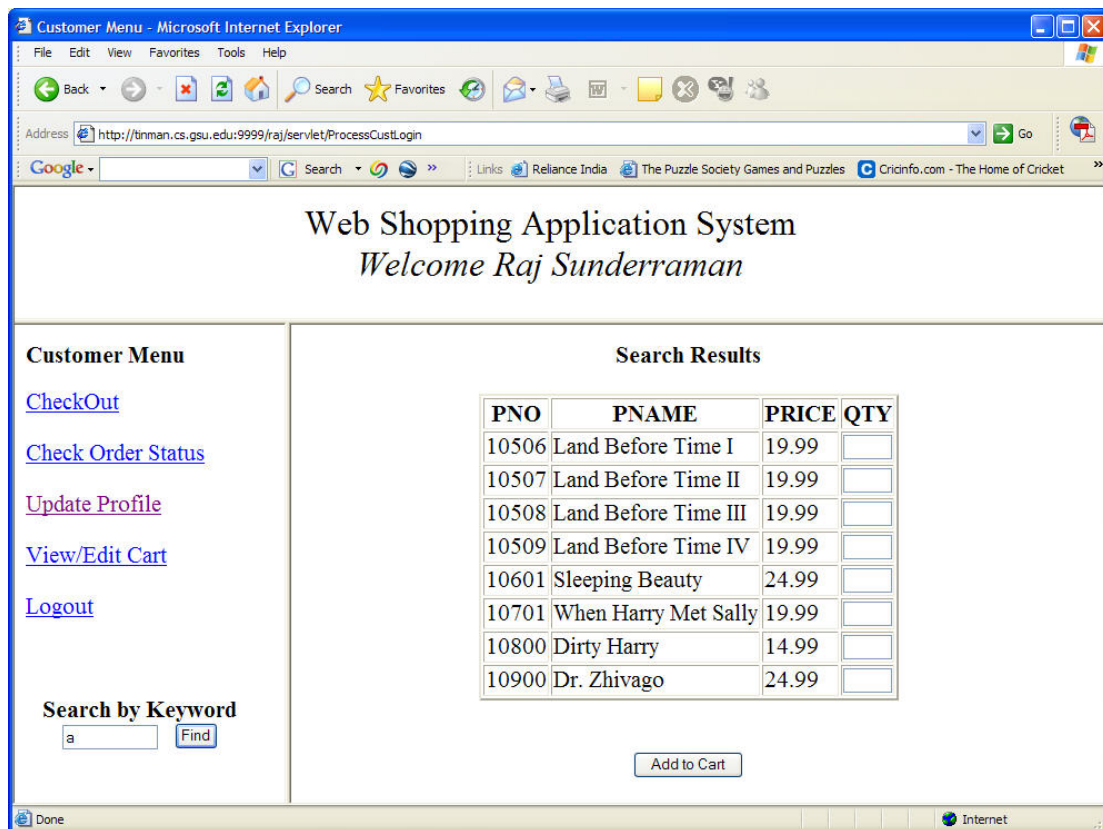


Figure 6.4: Web Shopping – Search Result Page



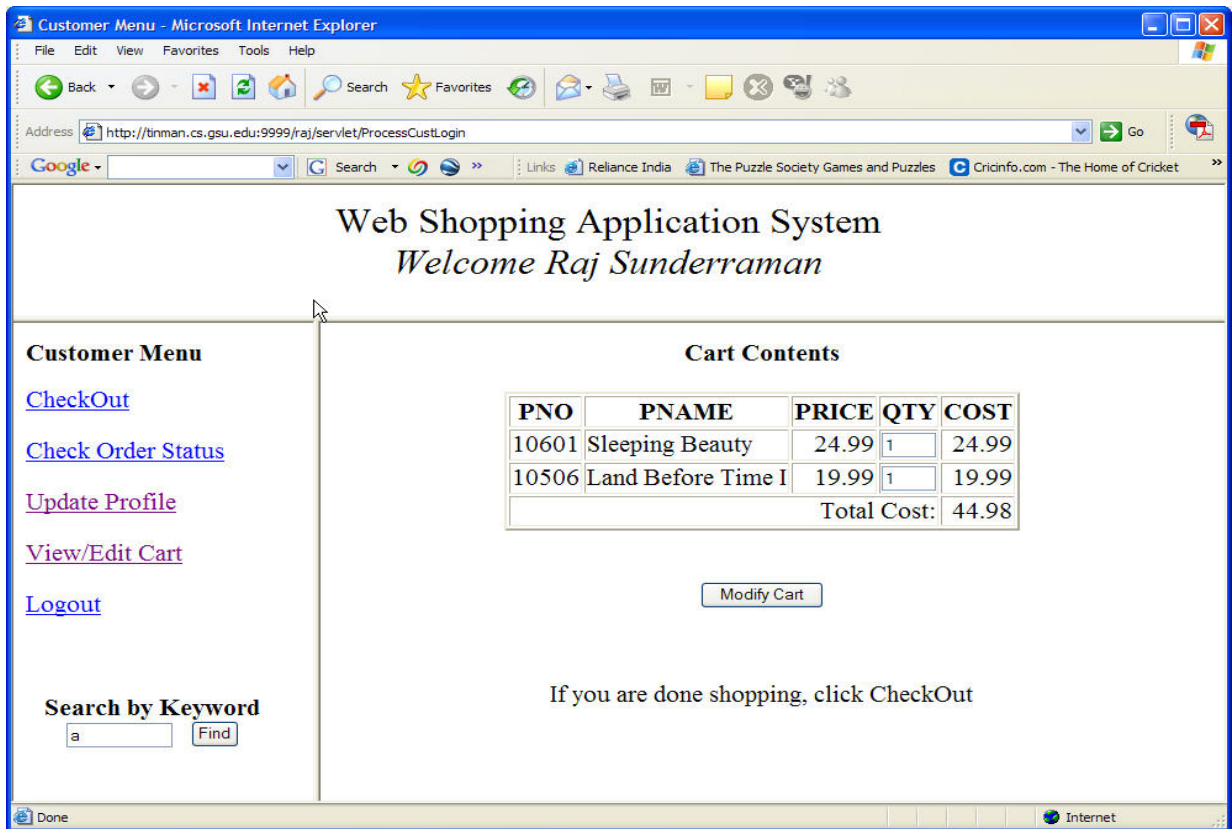


Figure 6.5: Web Shopping – View/Edit Cart

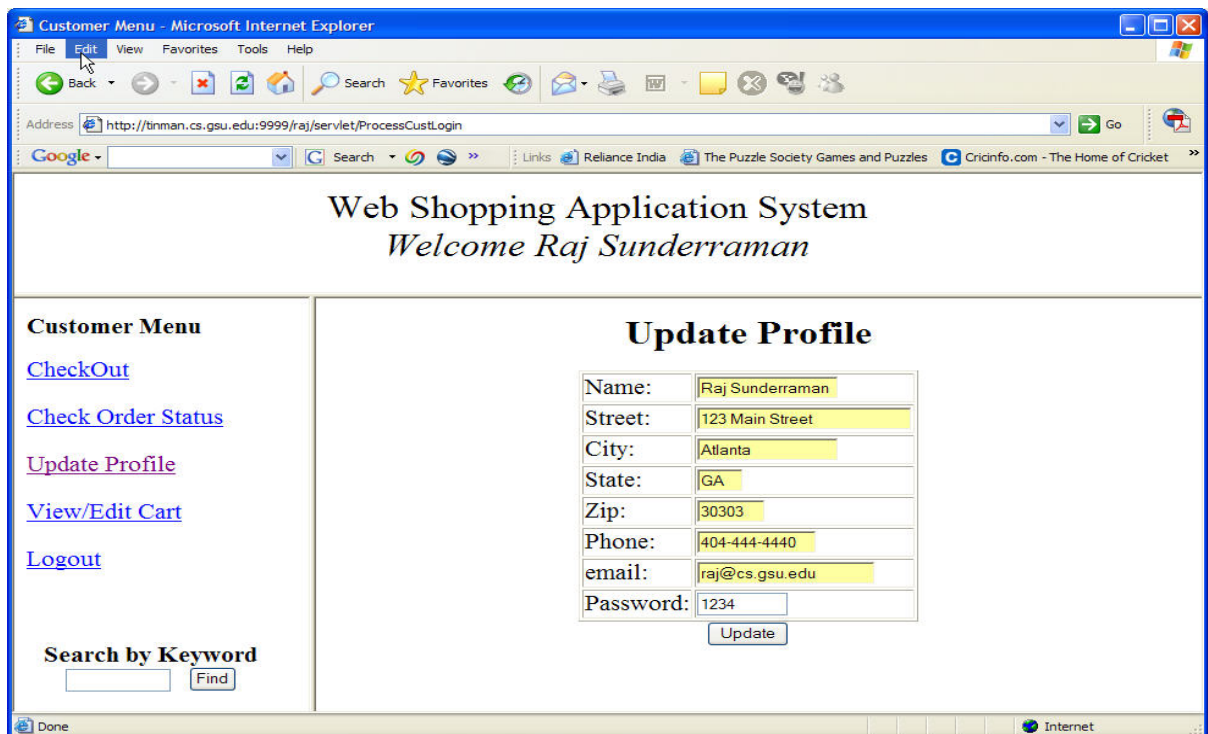


Figure 6.6: Web Shopping – Update Profile

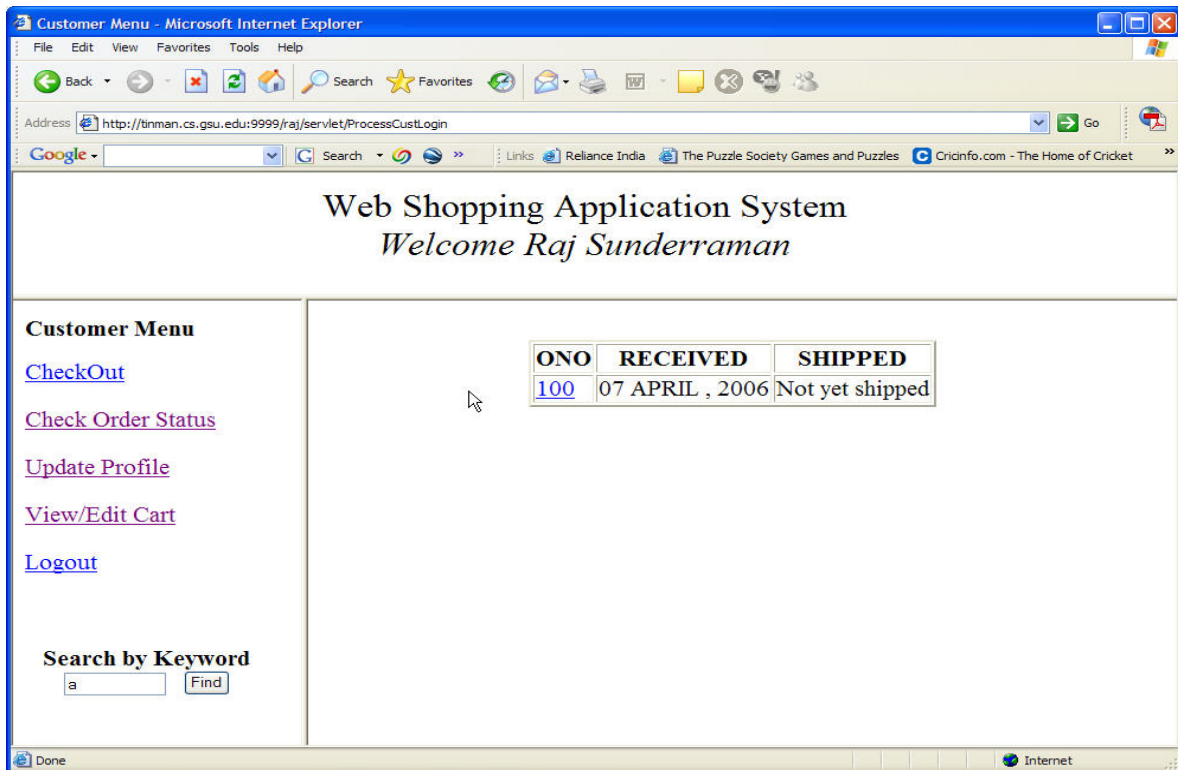


Figure 6.7: Web Shopping – Check Order Status

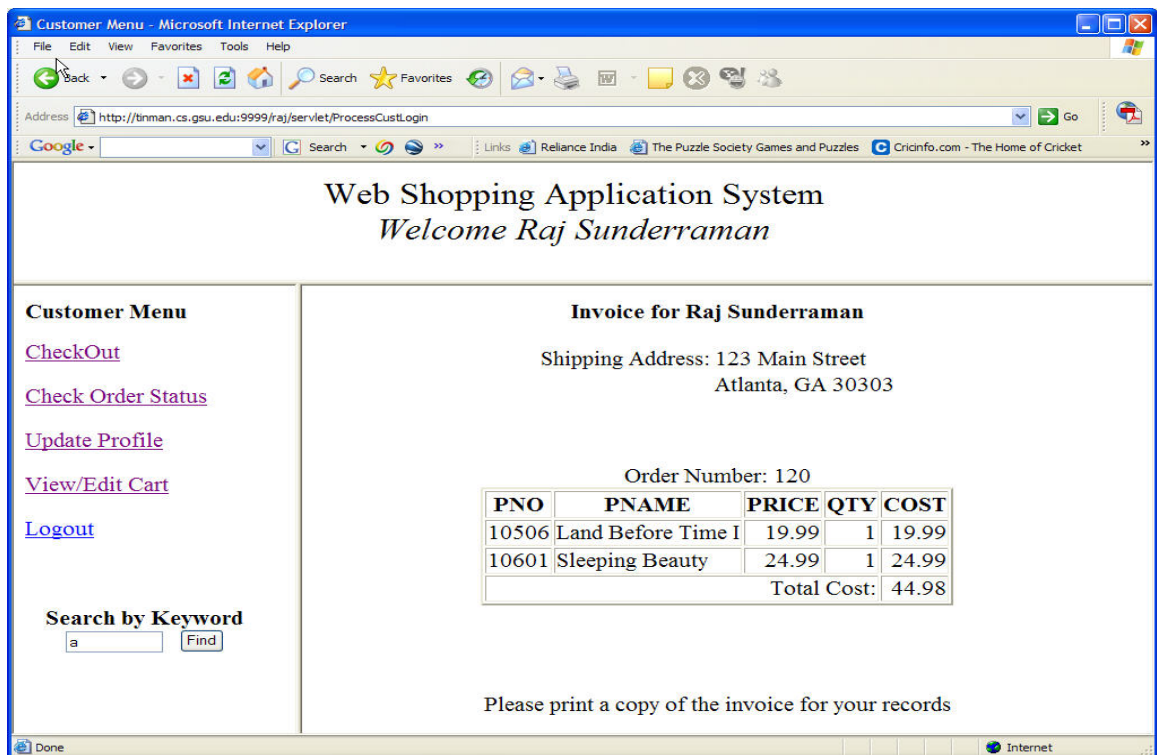
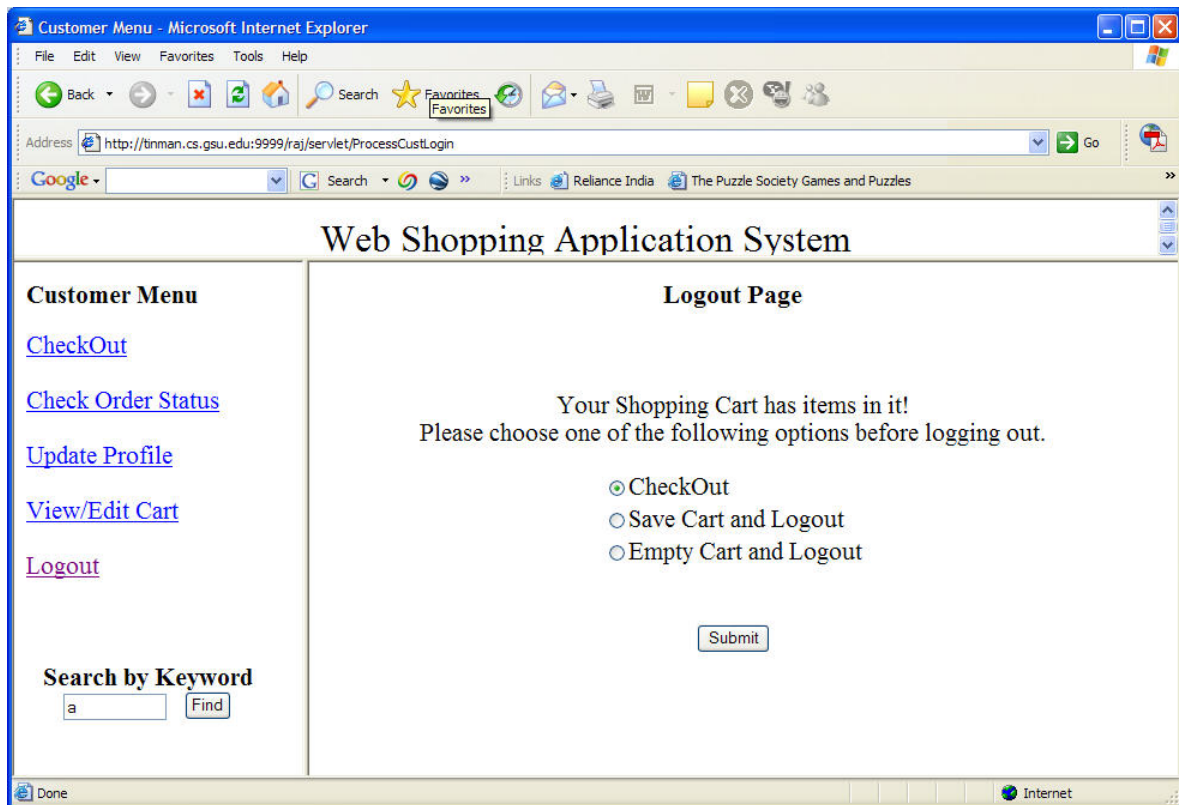


Figure 6.8: Web Shopping – Check Out



**Figure 6.9: Web Shopping – Log Out**

## 6.4 Online Bulletin Board System

Using PHP and MySQL implement an online bulletin board system that allows a set of authorized users to participate in an online discussion forum. The data for the bulletin board system should be stored in a MySQL database with the following schema:

```
create table bbusers (
  email      varchar(50),
  name       varchar(30),
  password   varchar(10),
  nickname   varchar(30),
  primary key (email)
);

create table postings (
  postId      integer(5) auto_increment,
  postDate    datetime,
  postedBy    varchar(50),
  postSubject varchar(100),
  content     varchar(512),
  ancestorPath varchar(100),
  primary key (postId),
  foreign key (postedBy) references bbusers
);
```

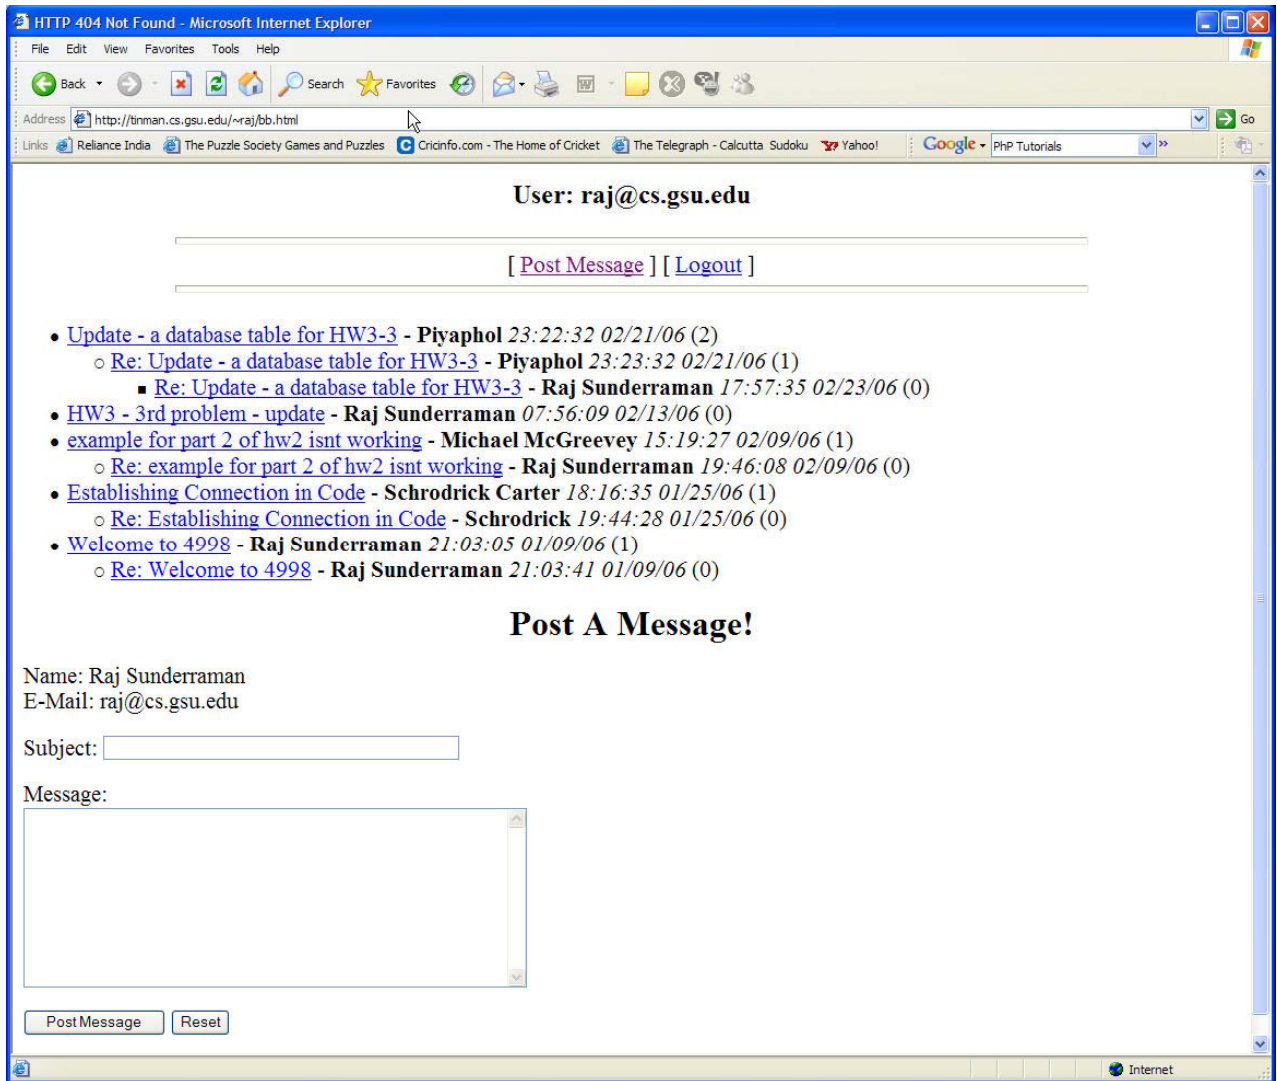
The database has two tables:

1. `bbusers`: This table records information about users of the bulletin board. The `email` and `password` fields are used for signing into the system.
2. `postings`: This table records information about all postings as well as follow-up postings of the bulletin board. Each posting is assigned a unique `postId`. To keep track of the “tree-structure” generated by follow-up postings, the system keeps track of the path from root message to the posting in the `ancestorPath` attribute. The path is recorded as a colon separated list of posting Ids; for example the ancestor path `1:5:6:12` would indicate that the current posting has a parent posting with `postId=12`, a grand-parent posting with `postId=6`, a great-grand-parent with `postId=5`, and a great-great-grandparent with `postId=1`. With this structure, the entire bulletin board messages can be viewed as a collection (forest) of trees.

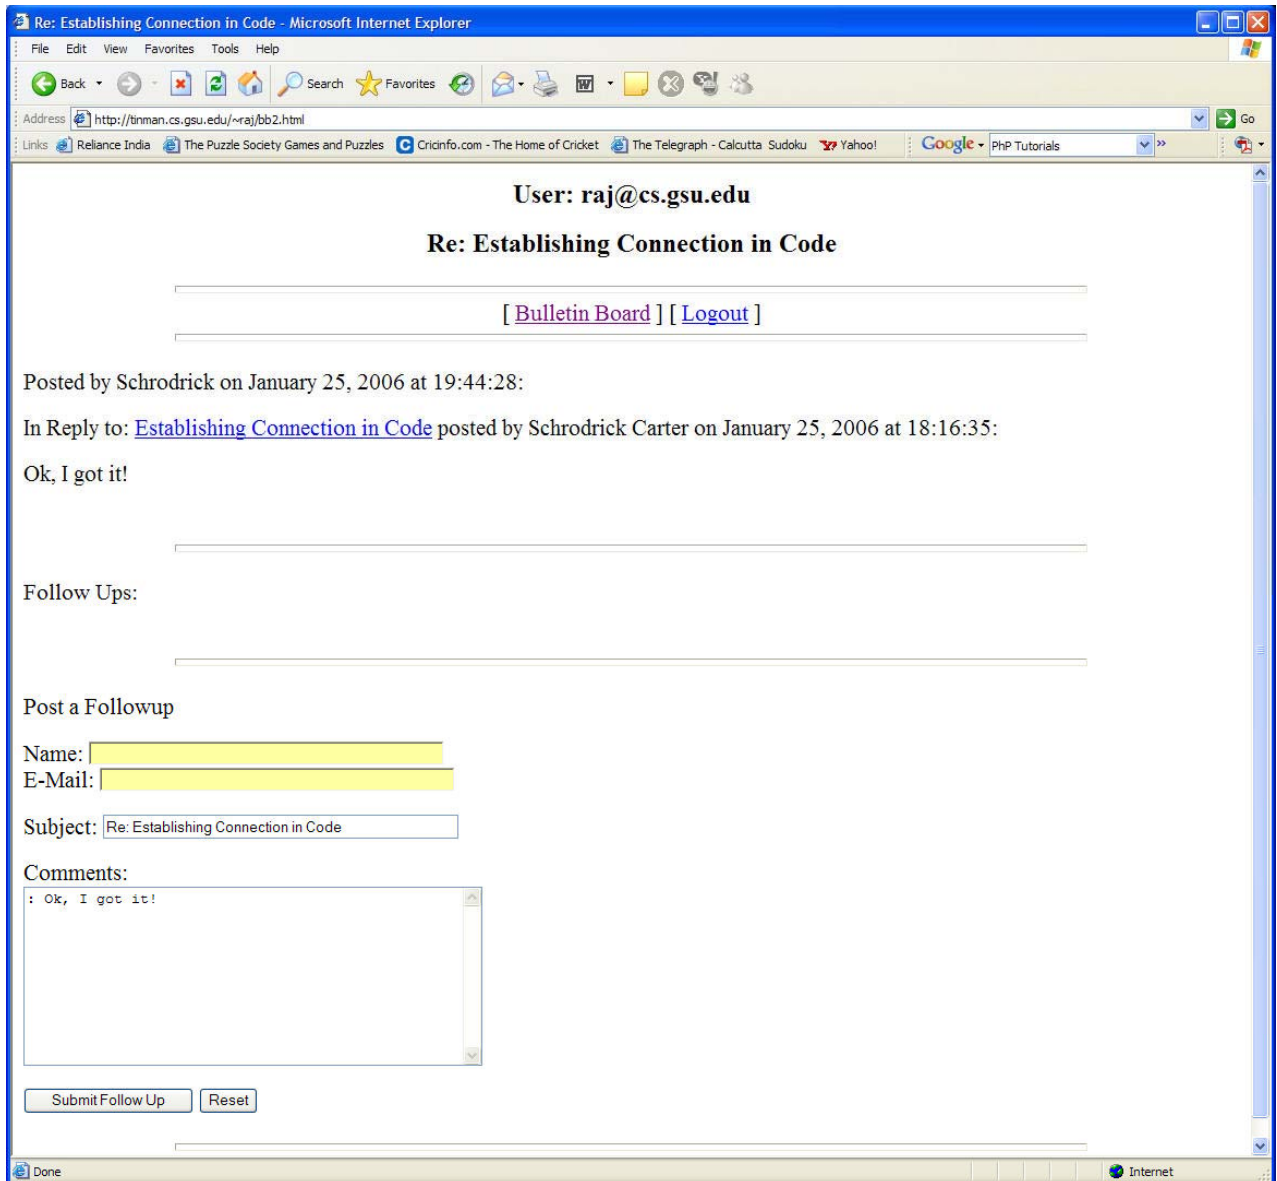
The Web application should implement the following basic functions:

1. User sign-in and sign-out.
2. Default display of messages in reverse chronological order and properly indented follow-up messages.
3. Post message and post follow-up message by user.

Figure 6.10 and 6.11 show possible user interfaces for the main display page and the follow-up display page.



**Figure 6.10: Bulletin Board – Main Display Page**



**Figure 6.11: Bulletin Board – Follow-up Listing Page**

## BIBLIOGRAPHY

1. R. Elmasri and S. Navathe, Fundamentals of Database Systems, 5<sup>th</sup> Edition, Addison-Wesley, 2007.
2. M. Fisher, J. Ellis , and J. Bruce, JDBC API Tutorial and Reference, 3<sup>rd</sup> Edition, Addison-Wesley Professional, 2003.
3. A. Silberschatz, H.F. Korth, and S. Sudarshan, Database Systems Concepts, 4<sup>th</sup> Edition, McGraw Hill, 2002.
4. R. Sunderraman, Oracle 9i Programming: A Primer, Addison-Wesley, 2004.
5. J.D. Ullman, Principles of Database and Knowledge-Base Systems, Volume 1, Computer Science Press, 1988.
6. J.D. Ullman and J. Widom, A First Course in Database Systems, 2<sup>nd</sup> Edition, Prentice Hall 2002.
7. H. Williams and D. Lane, Web Database Applications with PHP and MySQL, O'Reilley, 2004.

Useful URLs:

1. Computer Associates: <http://www.ca.com/>
2. Java: <http://java.sun.com>
3. JFlex: <http://www.jflex.de/>
4. JCup: <http://www.cs.princeton.edu/~appel/modern/java/CUP/>
5. SWI Prolog: <http://www.swi-prolog.org/>
6. MySQL: <http://www.mysql.com/>
7. Oracle: <http://www.oracle.com/index.html>
8. PHP: <http://www.php.net/>
9. SWI Prolog: <http://www.swi-prolog.org/>